
Discovering Cooperative Pipelines: Autoresearch for Sequential Social Dilemmas

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 We study *two-level autoresearch* for cooperation: an outer-loop AI agent au-
2 tonomously redesigns the inner-loop pipeline of an LLM policy-synthesis sys-
3 tem for multi-agent Sequential Social Dilemmas (SSDs). A *researcher agent* $\mathcal{R}$
4 (Claude Opus 4.6, run as a coding agent) reads the inner-loop source code, edits
5 system prompts, feedback functions, helper libraries, and iteration logic, runs eval-
6 uations, and decides what to keep, following the *autoresearch* paradigm. Across
7 two games (Cleanup and Gathering), two policy-synthesizer LLMs (Claude Son-
8 net 4.6, Gemini 3.1 Pro), and two welfare objectives (utilitarian efficiency and
9 Rawlsian maximin), the researcher reliably exceeds hand-designed baselines,
10 sharply tightens run-to-run variance, and outperforms prompt-only optimization.
11 The discovered pipelines are objective-dependent: under maximin the researcher
12 independently rediscovers *duty rotation* as a fairness mechanism, supporting an
13 *information-design* reading in which the researcher chooses what to reveal to
14 the boundedly rational synthesizer. Code at [https://github.com/anon590/](https://github.com/anon590/llm-policies-social-dilemmas)
15 [llm-policies-social-dilemmas](https://github.com/anon590/llm-policies-social-dilemmas).

16 1 Introduction

17 Sequential Social Dilemmas (SSDs) [1] are the multi-agent analogue of the prisoner’s dilemma in
18 temporally rich Markov games: individually rational play leads to collectively suboptimal outcomes
19 through pollution, over-harvesting, or open conflict. Standard multi-agent reinforcement learning
20 (MARL) struggles in this regime due to credit assignment, non-stationarity, and large joint action
21 spaces [2]. A complementary approach, recently introduced by Gallego [3], replaces parameter-
22 space optimization with *algorithm-space synthesis*: a frozen LLM writes a Python policy function,
23 evaluates it in self-play, and iteratively refines it from performance feedback. A single generation
24 step can produce coordination logic (territory partitioning, role assignment, conditional cooperation)
25 at a sample efficiency several orders of magnitude beyond what gradient-based MARL achieves on
26 the same environments.

27 This shifts where the design problem lives, rather than removing it. The inner-loop pipeline that
28 drives the synthesizer has many free parameters: which system prompt, which feedback variables,
29 which helper functions, how many refinement steps. Each materially affects the resulting policies,
30 and prior work tuned them by hand. A natural question follows: *can an AI agent design the pipeline?*

31 We answer affirmatively with a two-level autoresearch framework. An *outer-loop researcher agent*
32 (Claude Opus 4.6, run as a coding agent) edits the source files of an *inner-loop policy synthesizer*
33 (another LLM), runs evaluations on held-out seeds, and keeps modifications that improve a fixed
34 welfare objective Φ . The outer agent operates on an ordinary git repository (reading code, writing
35 diffs, running shell commands, etc) without task-specific scaffolding, mirroring the autoresearch
36 paradigm of Karpathy [4] for single-GPU LLM pretraining. Although the inner-loop SSDs are

gridworld benchmarks rather than physical systems, the outer-loop discovery process itself runs under conditions a deployed discovery agent faces: noisy multi-seed evaluations, stochastic code generation, an LLM-evaluation budget that bounds how often Φ can be queried, and a heterogeneous code repository the agent must navigate end-to-end on its own.

Our contributions are: i) a general two-level framework that delegates the design of an LLM synthesis pipeline to a coding agent operating on a real software repository (Section 3); ii) the first instantiation of the autoresearch paradigm in a multi-agent decision-making domain, with experiments across two SSDs (Cleanup and Gathering), two policy LLMs, and two welfare objectives (utilitarian efficiency U and Rawlsian maximin $\min_i R_i$) (Section 4); and iii) a mechanism-design interpretation supported by the qualitatively different pipelines the agent produces under different welfare objectives, including the autonomous discovery of *time-based duty rotation* as a fairness mechanism, never found under efficiency optimization.

2 Background

We build on the iterative LLM policy synthesis framework of Gallego [3], which serves as the (frozen) inner loop of our two-level system. This section recalls the SSD formalism, the social outcome metrics, and the base synthesis loop.

2.1 Sequential Social Dilemmas

A Sequential Social Dilemma is a partially observable Markov game $\mathcal{G} = \langle N, \mathcal{S}, \{\mathcal{A}_i\}_{i=1}^N, T, \{R_i\}_{i=1}^N, H \rangle$ with N agents, state space $\mathcal{S}$ (the gridworld configuration), per-agent action spaces $\mathcal{A}_i$, transition function T , reward functions R_i , and episode horizon H [1]. Beyond the dilemma’s matrix-game structure, SSDs add temporal richness: agents must learn *when* and *where* to cooperate, not just *whether* to.

We study two canonical SSDs that capture complementary dilemma types.

Cleanup [5] is a *public goods provision* game ($N=10$). The map has two regions: a river that accumulates waste, and an orchard where apples grow. Apples regrow only when river pollution is below threshold. Each agent can fire a cleaning beam (cost -1) that removes waste, collect apples ($+1$ each), or fire a tagging beam (cost -1 , inflicting -50 on the target and removing it for 25 steps). The dilemma: cleaning is costly but its benefits are public, so purely self-interested agents free-ride.

Gathering [1, 6] is a *common pool resource* game ($N=4$). Agents collect shared apples on a fixed respawn timer and may fire tagging beams to temporarily remove rivals. The dilemma: agents can coexist and share resources, or aggress to monopolize them; aggression wastes time and reduces total welfare.

Both games use 8–9 discrete actions (movement, rotation, beam, stand, optionally clean) and episodes of $H=1000$ steps. The two dilemmas differ in cost structure: asymmetric provision (cleaners pay, all benefit) vs. symmetric restraint (every agent faces the same temptation), a distinction that drives our experimental findings.

Social metrics. Following Perolat et al. [6], let $R_i = \sum_{t=0}^{H-1} r_i^t$ denote agent i ’s episode return. We evaluate four social outcomes:

$$U = \frac{1}{H} \sum_{i=1}^N R_i \quad (\text{efficiency}) \quad (1)$$

$$E = 1 - \frac{\sum_{i,j} |R_i - R_j|}{2N \sum_i R_i} \quad (\text{equality}) \quad (2)$$

$$S = \frac{1}{N} \sum_{i=1}^N \bar{t}_i \quad (\text{sustainability}) \quad (3)$$

$$P = \frac{1}{H} \sum_{t=0}^{H-1} |\{i : \text{active}_i^t\}| \quad (\text{peace}) \quad (4)$$

75 where $\bar{t}_i$ is the mean timestep at which agent i collects positive reward (higher $\Rightarrow$ resources preserved
 76 later) and active_i^t indicates that agent i is not tagged out at step t . We additionally consider the
 77 *maximin* (Rawlsian) welfare criterion $\min_i R_i$, which optimizes the worst-off agent’s return and
 78 serves as the second objective in our experiments.

79 2.2 Iterative LLM Policy Synthesis

80 Let Π denote the space of *code-based policies*: deterministic functions $\pi : \mathcal{S} \times [N] \rightarrow \mathcal{A}$ expressed
 81 as executable Python code. Each policy has access to the full environment state and a library of
 82 helpers (BFS pathfinding, beam targeting, coordinate transforms). This state access is a deliberate
 83 design choice: programmatic policies operate in algorithm space rather than the reactive observation-
 84 to-action space of neural policies, which lets a single LLM generation step encode rich coordination
 85 logic.

86 A frozen LLM $\mathcal{M}$ acts as the *policy synthesizer*. Given a system prompt p describing the environ-
 87 ment API and a feedback prompt q_k , it produces a new policy

$$\pi_{k+1} = \mathcal{M}(p, q(\pi_k, \mathcal{F}_k)),$$

88 where π_k is the previous policy (its source code) and $\mathcal{F}_k$ is the evaluation feedback. All N agents ex-
 89 ecute the same policy in homogeneous self-play, and evaluation over a set of random seeds S yields
 90 the mean per-agent return $\bar{r}_k$ and the social metrics vector $\mathbf{m}_k = (U_k, E_k, S_k, P_k)$. Each generated
 91 policy passes an AST-based safety check (blocking eval, file I/O, network access) followed by a
 92 short smoke test; failures trigger regeneration (up to R attempts) with the error message appended
 93 to the prompt.

94 **Feedback.** The hand-designed feedback used by Gallego [3] as our starting point packages the
 95 previous policy’s source code together with all available evaluation signals:

$$\mathcal{F}_k = (\text{code}(\pi_k), \bar{r}_k, \mathbf{m}_k, \mathbf{d}), \quad (5)$$

96 where $\mathbf{d}$ contains natural-language definitions of each social metric. Exposing the full social-metrics
 97 vector (rather than the scalar reward alone) disambiguates distinct failure modes that share the same
 98 average return (e.g., under- vs. over-cleaning in Cleanup) and consistently improves cooperative
 99 outcomes. On Cleanup, this baseline feedback reaches $U=1.93/2.70$ (mean/max) with Gemini and
 100 $U=0.86/1.56$ with Sonnet across 4 runs each. The choice of feedback content is a single design
 101 decision within a much larger pipeline configuration space; our two-level framework (Section 3)
 102 opens the full space to automated search.

103 3 Two-Level Framework

104 We introduce a two-level system where a *researcher agent* $\mathcal{R}$ autonomously discovers configurations
 105 that optimize the output of an inner-loop system. While we instantiate this for multi-agent policy
 106 synthesis, the architecture is general: any pipeline where an LLM generates artifacts, evaluates them,
 107 and iterates can serve as the inner loop. The key insight is that the entire inner-loop codebase is a
 108 designable artifact that an AI agent can search over. Figure 1 illustrates the architecture.

109 3.1 Configuration Space

110 Let $\mathcal{C}$ denote the space of *pipeline configurations*. Each configuration $c \in \mathcal{C}$ specifies the full inner-
 111 loop setup:

$$c = (p, \phi, \mathcal{H}, \iota, v) \quad (6)$$

112 where p is the system prompt, ϕ is the feedback construction function (which metrics and diagnostics
 113 to include, how to frame it, whether to inject adaptive hints), $\mathcal{H}$ is the helper function library, ι
 114 specifies the iteration logic (number of iterations K , sampling strategy, temperature schedule), and
 115 v is the validation pipeline. Table 1 provides concrete examples.

116 The hand-designed feedback of [3] corresponds to a single fixed instantiation of ϕ . Our framework
 117 opens the full configuration space to automated search.




Figure 1: Two-level automated research framework (Algorithm 1).

Algorithm 1 Two-Level Automated Research

Require: Game $\mathcal{G}$, policy synthesizer $\mathcal{M}$, researcher $\mathcal{R}$, system prompt $p_{\mathcal{R}}$, initial config c_0 , outer iterations J , ground-truth objective Φ , held-out seeds S_{ho}

Ensure: Best configuration c^* , best policy π^*

```

1:  $\pi_0^* \leftarrow \text{INNERLOOP}(\mathcal{M}, \mathcal{G}, c_0)$ 
2:  $\Phi_0 \leftarrow \text{EVAL}(\pi_0^*; \mathcal{G}, S_{\text{ho}})$ 
3:  $\text{history} \leftarrow \{(c_0, \Phi_0, \mathbf{m}_0, \emptyset)\}$ 
4: for  $j = 1, \dots, J$  do
5:    $c_j \leftarrow \mathcal{R}(p_{\mathcal{R}}, \text{CODE}(c_{j-1}), \text{history})$ 
6:   if  $\neg \text{VALIDATECONFIG}(c_j)$  then  $\text{retry} \leq R$ 
7:      $\pi_j^* \leftarrow \text{INNERLOOP}(\mathcal{M}, \mathcal{G}, c_j)$ 
8:      $\Phi_j \leftarrow \text{EVAL}(\pi_j^*; \mathcal{G}, S_{\text{ho}})$ 
9:      $\Delta_j \leftarrow \text{DIFF}(c_{j-1}, c_j)$  // code diff
10:     $\text{history} \leftarrow \text{history} \cup \{(c_j, \Phi_j, \mathbf{m}_j, \Delta_j)\}$ 
11:  end for
12:  $j^* \leftarrow \arg \max_j \Phi_j$ 
13: return  $c_{j^*}, \pi_{j^*}^*$ 

```

Table 1: Configuration components modifiable by the researcher, with concrete edits that $\mathcal{R}$ made in our runs. The environment simulator, ground-truth evaluation, and policy LLM weights are frozen.

Scope	Concrete edits by $\mathcal{R}$ (see appendix)
p System prompt	Multi-section strategic briefing with explicit duty-rotation template (<code>agent_id + step//T</code>) % n (App. D.2, Listing 3)
ϕ Feedback fn	Thresholded diagnostics: “FAIRNESS ALERT” fires when $\min_i R_i < 0$; “DO NOT REGRESS” guard fires when $U \geq 2.5$ (App. D.3, Listing 4)
$\mathcal{H}$ Helper library	BFS-Voronoi territories with respawn-timer awareness; band-based apple zoning by <code>agent_id</code> (App. D.1, Listings 2, 1)
l Iteration logic	Per-condition $K \in \{2, 3\}$; $ S $ walked $5 \rightarrow 8 \rightarrow 12$ for variance control; thinking budget 10–32k (App. D.4, Table 6)
v Validation	AST blocks <code>eval/file</code> I/O/network; 50-step smoke test; $\text{retry} \leq R$ on failure

3.2 Inner Loop (Policy Synthesis)

Given a configuration c , the inner loop executes K iterations of LLM policy synthesis:

$$\pi_c^* = \text{INNERLOOP}(\mathcal{M}, \mathcal{G}, c) \quad (7)$$

Each iteration k proceeds in four stages, following [3]:

- Synthesize.** The policy LLM $\mathcal{M}$ receives the system prompt p , the previous policy’s source code π_{k-1} , and feedback $\mathcal{F}_{k-1}$ constructed by ϕ . It generates a new Python policy function π_k that has access to full environment state and the helper library $\mathcal{H}$.
- Validate.** The generated code undergoes AST-based safety checks (blocking dangerous operations such as file I/O and network access) followed by a short smoke test. Failures trigger re-generation (up to R retries), with the error message appended to the prompt.
- Evaluate.** All N agents execute the same policy π_k in self-play over $|S|$ random seeds. The evaluation yields the mean per-agent reward $\bar{r}_k$ and the social metrics vector $\mathbf{m}_k = (U_k, E_k, S_k, P_k)$.
- Feedback.** The feedback function ϕ constructs the prompt for the next iteration from $(\pi_k, \bar{r}_k, \mathbf{m}_k)$, packaging the previous policy’s source code together with the scalar reward, the social metrics vector, their natural-language definitions, and any adaptive diagnostics that ϕ injects.

The inner loop output is evaluated on held-out seeds via a fixed ground-truth objective:

$$\Phi(c) = \text{EVAL}(\pi_c^*; \mathcal{G}, S_{\text{held-out}}) \quad (8)$$

134 where Φ maps from social outcomes to a scalar. We consider two objectives:

$$\Phi_U = U \quad (\text{utilitarian efficiency}) \quad (9)$$

$$\Phi_{\min} = \min_i R_i \quad (\text{Rawlsian maximin}) \quad (10)$$

135 3.3 Outer Loop (Automated Research)

136 The researcher agent $\mathcal{R}$ iteratively modifies the pipeline configuration. Following the autoresearch
137 paradigm [4], $\mathcal{R}$ operates on the inner-loop codebase as a modifiable artifact, proposing changes,
138 observing outcomes, and refining. The procedure is formalized in Algorithm 1.

139 At each outer iteration j , the researcher $\mathcal{R}$ receives: i) the **full source code** of the current config-
140 uration c_{j-1} (prompts, feedback construction, helpers, iteration logic); ii) the **experiment history**:
141 for each prior iteration, the code diff Δ_i , ground-truth score Φ_i , and social metrics vector $\mathbf{m}_i$; iii)
142 the **environment source code** (read-only), enabling the researcher to reason about game mechanics.
143 The researcher proposes a new configuration c_j by generating code modifications. Concretely, $\mathcal{R}$
144 is a coding agent (Claude Code CLI) that operates on a real software repository: it reads and edits
145 Python source files, runs shell commands, inspects evaluation outputs, and commits changes to a
146 dedicated git branch, following the same workflow a human researcher would follow. This mirrors
147 the original autoresearch setup [4], where an AI coding agent modifies `train.py` and observes
148 validation loss; here, the researcher modifies the inner-loop codebase and observes social welfare.

149 3.4 Connection to Automated Mechanism Design

150 The two-level structure admits a mechanism design interpretation. The researcher $\mathcal{R}$ acts as a *mech-*
151 *anism designer*: it controls the information structure (what metrics to reveal, how to frame them),
152 the action space (what helper functions are available), and the incentive structure (how feedback is
153 presented) under which the policy synthesizer $\mathcal{M}$ operates. The synthesizer acts as the *agent* within
154 the designed mechanism.

155 This connects to the automated mechanism design literature [7], where a principal designs rules to
156 induce desired behavior from self-interested agents. In our setting:

- 157 • The **principal** is the researcher $\mathcal{R}$, optimizing the ground-truth objective Φ .
- 158 • The **agent** is the synthesizer $\mathcal{M}$, optimizing per-agent reward as instructed.
- 159 • The **mechanism** is the configuration c : prompts, feedback, helpers, iteration logic.
- 160 • The **outcome** is the social welfare of the resulting multi-agent policy π_c^* .

161 A key distinction from classical mechanism design: $\mathcal{M}$ is not *strategically* self-interested—it fol-
162 lows instructions—but it has *bounded rationality* in the sense that its ability to synthesize effective
163 policies depends on the information and tools provided. The researcher’s task is thus closer to *infor-*
164 *mation design* [8]: choosing what to reveal to help $\mathcal{M}$ navigate the cooperation–defection tension.
165 Our experiments (Section 4) show that the researcher designs qualitatively different information
166 structures depending on the welfare objective Φ , supporting this interpretation empirically.

167 4 Experiments

168 4.1 Setup

169 We conduct 12 autonomous researcher runs across a factorial design. The researcher agent $\mathcal{R}$ is
170 Claude Opus 4.6, invoked via the Claude Code CLI as a coding agent. Each run operates on a dedi-
171 cated git branch of a real Python codebase: $\mathcal{R}$ edits source files in `pipeline/`, executes evaluation
172 scripts, reads metric outputs, and iterates—without human intervention.

173 **Design.** For Cleanup: 2 policy LLMs $\times$ 2 objectives $\times$ 2 replications = 8 runs. For Gathering: 2
174 policy LLMs $\times$ 1 objective $\times$ 2 replications = 4 runs. Maximin runs are unnecessary for Gather-
175 ing because efficiency optimization alone achieves equality $E > 0.94$ (the game’s symmetric cost
176 structure makes fairness a free byproduct of coordination).

177 **Models.** Policy synthesizer $\mathcal{M}$: Gemini 3.1 Pro (Google) or Claude Sonnet 4.6 (Anthropic). Both
178 use extended thinking. These are the same models evaluated in [3]. We additionally test with

Table 2: Results. Mean / max across 2 runs per condition (4 for Cleanup baselines). Baseline: unmodified pipeline from [3]. GEPA: automated prompt optimization [9] with matched compute budget.

Policy LLM	Target	U	E	$\min_i R_i$
<i>Cleanup (N=10) — Baseline</i>				
Gemini 3.1 Pro	—	1.93/2.70	0.17/0.62	−159/−84
Sonnet 4.6	—	0.86/1.56	−0.02/0.25	−151/−59
<i>Cleanup (N=10) — GEPA [9]</i>				
Gemini 3.1 Pro	Φ_U	1.34/1.37	−0.10/−0.05	−149/−126
	$\Phi_{\min}$	1.76/2.76	0.90/0.96	143/245
Sonnet 4.6	Φ_U	1.04/1.20	−0.59/−0.21	−164/−126
	$\Phi_{\min}$	−0.63/1.60	0.77/1.00	−147/6
<i>Cleanup (N=10) — Autoresearch</i>				
Gemini 3.1 Pro	Φ_U	3.20 /3.25	0.55/0.61	−211/−182
	$\Phi_{\min}$	3.16/3.19	0.98 /0.98	290 /296
Sonnet 4.6	Φ_U	3.12 /3.14	0.66/0.70	−196/−146
	$\Phi_{\min}$	2.57/2.93	0.91 /0.97	179 /200
<i>Gathering (N=4) — Baseline</i>				
Gemini 3.1 Pro	—	2.04/2.42	0.90/0.98	412/571
Sonnet 4.6	—	0.03/0.03	0.54/0.54	0/0
<i>Gathering (N=4) — GEPA [9]</i>				
Gemini 3.1 Pro	Φ_U	2.08/2.35	0.94/0.96	436/518
Sonnet 4.6	Φ_U	1.20/1.23	0.63/0.71	86/123
<i>Gathering (N=4) — Autoresearch</i>				
Gemini 3.1 Pro	Φ_U	2.49/2.51	0.98/0.98	582/582
Sonnet 4.6	Φ_U	2.52/2.52	0.96/0.98	576/597

179 Gemma 4 26B-A4B-IT (Google), a smaller open-weight model, to probe the framework’s behavior
 180 when the policy synthesizer has substantially lower capability (Appendix C).

181 **Baselines.** The hand-designed feedback configuration from [3] serves as the initial pipeline c_0
 182 for all runs. On Cleanup ($N=10$), this baseline achieves $U=1.93/2.70$ (mean/max) with Gemini
 183 and $U=0.86/1.56$ with Sonnet across 4 runs each. We additionally compare against GEPA [9], an
 184 automated prompt optimization method that iteratively refines the system prompt via LLM reflection.
 185 GEPA optimizes only the system prompt p , whereas our researcher modifies the full pipeline
 186 $(p, \phi, \mathcal{H}, \iota, v)$. We give GEPA a matched compute budget: the same number of optimization steps
 187 as outer iterations used by our method (each GEPA step costs ~ 3 policy evaluations, comparable to
 188 our inner loop’s $K+1=4$).

189 **Inner loop parameters.** $K=3$ synthesis iterations (or $K=2$ if the researcher elects), evaluated
 190 over $|S|=5-12$ random seeds (researcher-configurable). The researcher ran $J=3-17$ outer iterations
 191 per run (mean: 8.3), with a keep rate of 18%–100% (mean: 63%).

192 4.2 Results

193 Table 2 presents the main results, comparing autoresearch against the hand-designed baseline of [3]
 194 and GEPA prompt optimization [9].

195 **Finding 1: The researcher reliably improves over hand-designed baselines and outperforms**
 196 **prompt-only optimization.** All 12 runs show substantial improvement regardless of starting point
 197 (Figure 2). In Cleanup, the researcher exceeds the unmodified pipeline baseline by 66% with Gemini
 198 ($U: 1.93 \rightarrow 3.20$) and 263% with Sonnet ($U: 0.86 \rightarrow 3.12$). Strikingly, the Sonnet improvement
 199 nearly closes the gap with Gemini ($U=3.12$ vs. 3.20), despite a $2\times$ baseline disadvantage (0.86 vs.
 200 1.93). Final efficiency values cluster tightly around $U \approx 3.1-3.2$ despite baselines varying from
 201 0.29 to 2.70 across individual runs, suggesting the researcher reliably discovers the performance
 202 ceiling of each policy LLM. Beyond improving the mean, the researcher yields remarkably consistent
 203 results: under efficiency optimization, Gemini’s runs cluster within $U=3.20-3.25$ (gap 0.05)




Figure 2: Efficiency (U) across researcher iterations for all 12 runs. *Left*: Cleanup ($N=10$) with 8 runs across 2 LLMs $\times$ 2 objectives (solid lines = kept experiments, open circles = discarded). Dashed horizontal line: hand-designed baseline from [3]. All runs converge to $U \approx 3.1$ – 3.2 despite diverse starting points. *Right*: Gathering ($N=4$) with 4 runs (2 per LLM). All converge to $U \approx 2.5$ within 2–4 iterations despite baselines ranging from 0.03 to 2.42.

despite baselines spanning 1.15–2.70, and Sonnet’s within 3.12–3.14 (gap 0.02) from baselines of 0.38–1.56. In Gathering, all four runs converge to $U=2.47$ – 2.52 from baselines ranging 0.03–2.42. The researcher discovers pipeline modifications—structured helpers (App. D.1), strategic hints (App. D.2), tailored feedback (App. D.3)—that constrain the policy LLM toward consistently good solutions.

Compared to GEPA (Table 2), autoresearch achieves $2.4\times$ higher efficiency under Φ_U in Cleanup ($U=3.20$ vs. 1.34), $2\times$ higher maximin under $\Phi_{\min}$ (290 vs. 143), and 20% higher efficiency in Gathering ($U=2.49$ vs. 2.08). Critically, GEPA shows far wider spread across runs on the harder fairness objective ($\Phi_{\min}$ maximin in Cleanup: 245/41 vs. autoresearch’s 296/284). The gap widens further with Sonnet: on Cleanup Φ_U , GEPA stalls at $U=1.04$ ($3\times$ below autoresearch’s 3.12), and on $\Phi_{\min}$ its two runs split between a modestly positive solution ($\min_i R_i=6$, $U=1.60$) and a pathological “everyone cleans, nobody eats” regime (perfect equality $E=1.00$ but negative efficiency $U=-2.87$ and $\min_i R_i=-299$), while autoresearch with the same model reaches $\min_i R_i=179/200$ reliably. This highlights the advantage of modifying the full pipeline—helpers, feedback, and iteration logic provide structured scaffolding that reduces run-to-run variance, whereas prompt-only optimization cannot address the underlying stochasticity of code generation.

Finding 2: No efficiency–fairness tradeoff in Cleanup (Gemini). Maximin-optimized Gemini pipelines sacrifice only 1% efficiency (U : 3.16 vs. 3.20) while achieving near-perfect equality (E : 0.98 vs. 0.55) and transforming maximin from deeply negative baselines (-99 to -84) to $\min_i R_i = 290$. The researcher discovers *fair duty rotation* (Listing 6; primed by the prompt rewrite of Listing 3)—time-based cycling using `agent_id` and `env._step_count`—simultaneously improving worst-off welfare *and* collective output. Because cleaning is a public good, distributing the cleaning cost fairly ensures enough cleaners to sustain apple production.

For Sonnet, there is a moderate tradeoff: efficiency drops from 3.12 to 2.57 under maximin optimization. The gap reflects Sonnet’s harder time implementing complex coordination mechanisms (role rotation, zone assignment) from strategic hints alone.

Finding 3: Game structure determines whether fairness requires explicit optimization. In Cleanup, where cleaning costs are borne asymmetrically (cleaners pay -1 , free-riders collect apples), baseline equality ranges from $E=0.04$ to 0.62, and maximin optimization is required to reach $E > 0.9$. In Gathering, where all agents face a symmetric landscape, efficiency optimization alone achieves $E > 0.94$ across all 4 runs—no separate maximin runs are needed.

235 This generalizes: in *provision* dilemmas with asymmetric costs, fairness requires designed mecha-
 236 nisms (role rotation, duty sharing); in *restraint* dilemmas with symmetric costs, fairness emerges as
 237 a free byproduct of efficient coordination. The researcher independently discovers this—it creates
 238 role differentiation pipelines only for Cleanup, and pure spatial-coordination pipelines for Gather-
 239 ing.

240 **Finding 4: Convergent discovery of qualitatively different strategies per objective.** Despite
 241 fully independent runs, the researcher converges on the same core strategies within each condition
 242 (Table 3, Appendix A). In Cleanup, waste-counting helpers and spatial zone partitioning appear
 243 across all 8 runs. The key qualitative difference is *role rotation*, appearing in 4/4 maximin runs
 244 (Listings 6, 7) but 0/4 efficiency runs (Listing 5). Under efficiency optimization, the researcher
 245 assigns static cleaning roles (some agents always clean), achieving high collective output at the cost
 246 of equality. Under maximin optimization, it discovers rotation as a fairness mechanism.

247 In Gathering, the researcher discovers BFS-Voronoi territory partitioning and respawn-timer aware-
 248 ness (Listings 2, 8), with no role differentiation—optimization is purely spatial (who collects which
 249 apples) and temporal (respawn-aware positioning).

250 **Common failure modes.** Several patterns led to discarded experiments across all runs: (1) *over-*
 251 *prescription*—adding too many strategic hints confuses the policy LLM, causing generation failures;
 252 (2) *iteration regression*— $K=4$ or $K=5$ inner iterations often cause catastrophic regression as the
 253 LLM “over-refines” a working policy; (3) *feedback overload*—verbose per-agent diagnostics cause
 254 over-correction rather than gradual improvement; (4) *variance exposure*—identical configurations
 255 can yield $U=3.2$ then $U=0.09$ on different seeds, requiring multi-seed evaluation for stable signals
 256 (per-run K , $|S|$ in Table 6).

257 5 Related Work

258 **Sequential social dilemmas.** SSDs were introduced by Leibo et al. [1] (Gathering) and extended
 259 to public-goods settings by Hughes et al. [5] (Cleanup); Perolat et al. [6] formalized the social
 260 outcome metrics (U, E, S, P) used here. Subsequent work studies inequity aversion, intrinsic moti-
 261 vation, and reputational mechanisms for promoting cooperation in MARL agents. We complement
 262 this line by automating the search for cooperative *programs* rather than evolving neural policies,
 263 and by treating the welfare objective Φ as a designable parameter that the outer agent optimizes
 264 for—obtaining qualitatively different cooperative behaviors (static role assignment vs. duty rotation)
 265 under different objectives without changing the environment.

266 **LLMs for policy and program synthesis.** FunSearch [10] evolves programs for combinato-
 267 rial discovery; Eureka [11] synthesizes reward functions from environment source code; Voy-
 268 ager [12] and Code as Policies [13] generate executable skill code for single-agent embodied control;
 269 ReEvo [14] evolves heuristics with reflective feedback. These works target single-agent settings or
 270 non-strategic optimization. Gallego [3], on which our inner loop is based, extended LLM program
 271 synthesis to the multi-agent SSD setting, where one program must coordinate N self-play copies.
 272 Our contribution is one level above: rather than tuning the synthesizer for one task, we let an outer
 273 agent rewrite the synthesis pipeline itself.

274 **LLM reflection and prompt optimization.** Reflexion [15], Self-Refine [16], OPRO [17], and
 275 GEPA [9] demonstrate that structured verbal feedback loops improve LLM outputs; ERL [18] in-
 276 ternalizes such reflection via self-distillation. These methods optimize the prompt or the model’s
 277 reasoning trajectory in isolation. Our framework instead modifies the entire surrounding pipeline—
 278 system prompt, feedback construction, helper library, iteration logic, validation—making prompt
 279 optimization (cf. GEPA) one component of a strictly larger search space. The empirical gap is
 280 large: Section 4 shows full-pipeline autoresearch achieving $3\times$ higher efficiency in Cleanup and
 281 37% higher in Gathering than GEPA at matched compute, with run-to-run spread an order of mag-
 282 nitude tighter on the harder fairness objective.

283 **Automated AI research.** A small but growing body of work delegates the design of ML pipelines
 284 to LLM coding agents. Karpathy’s autoresearch [4] runs a coding agent that modifies `train.py`

for nanoGPT pretraining and is rewarded for validation loss. Our system shares this autoresearch architecture—coding agent + frozen evaluation harness + diff-based history—but targets a different inner loop (multi-agent policy synthesis instead of single-model pretraining) and a different objective (multi-agent social welfare instead of validation bits-per-byte). To our knowledge this is the first instantiation of the autoresearch paradigm in a multi-agent decision-making domain.

Automated mechanism and information design. Classical automated mechanism design [7] computes optimal allocation rules for self-interested agents under explicit incentive constraints, while information design [8] studies how a principal commits to a signaling policy that shapes a receiver’s behavior. Our outer agent solves a related but distinct problem: $\mathcal{M}$ is not strategically deceptive (it follows instructions) but is *boundedly rational*, so the researcher must decide what information, helpers, and structure to expose so that $\mathcal{M}$ writes policies achieving the principal’s welfare goal. Section 4 shows that the pipelines $\mathcal{R}$ produces are genuinely a function of the welfare objective, supporting this information-design framing empirically.

6 Discussion and Conclusion

We have presented a two-level framework where an AI coding agent autonomously discovers pipeline configurations that improve the output of an inner-loop LLM system. Applied to multi-agent policy synthesis in social dilemmas, the researcher agent reliably exceeds hand-designed baselines across 12 independent runs, and converges on qualitatively similar strategies within each game-objective condition. The framework requires no domain-specific scaffolding: the agent operates on a standard software repository using the same tools (file editing, shell commands, git) available to a human researcher.

Mechanism design in action. Our results empirically support the mechanism design interpretation of Section 3.4. The researcher acts as an information designer: under Φ_U , it reveals efficiency-oriented information (waste counts, zone assignments) that guides $\mathcal{M}$ toward productive but unequal role allocation (Listing 5). Under $\Phi_{\min}$, it additionally reveals fairness-oriented structure—rotation schedules and per-agent equity feedback (Listings 3, 4)—guiding $\mathcal{M}$ toward egalitarian coordination. The striking absence of a fairness–efficiency tradeoff with Gemini suggests that in public goods provision, the mechanism design problem has a cooperative optimum where both objectives align—the researcher discovers this independently.

Two types of dilemma. The Cleanup–Gathering contrast reveals a structural insight about social dilemmas: *provision* dilemmas (asymmetric costs, one agent’s contribution benefits all) require explicit fairness mechanisms, while *restraint* dilemmas (symmetric costs, each agent faces the same temptation) achieve fairness naturally when coordination improves. This distinction, well-known in the common-pool resource literature, is rediscovered here by an autonomous AI agent—suggesting it is a robust structural feature of these environments.

Human oversight and audit. Our system is fully autonomous by design, but its architecture offers natural affordances for human-in-the-loop oversight. Because the researcher $\mathcal{R}$ operates on a standard git repository—committing each pipeline modification to a dedicated branch—a domain expert can audit the full history of code diffs Δ_j alongside their metric outcomes $(\Phi_j, \mathbf{m}_j)$, exactly as they would review a colleague’s pull requests. The keep/discard decision (Algorithm 1, line 11) is a lightweight intervention point: a human could override the accept threshold, veto modifications that introduce undesirable mechanisms (e.g., exploitative role assignments), or steer the researcher toward under-explored regions of the configuration space by editing the experiment history. More broadly, the separation between the researcher $\mathcal{R}$ and the frozen evaluation Φ creates a *delegation boundary*: the human defines *what* to optimize (the welfare objective), while the agent decides *how*. This mirrors the expert-in-the-loop design patterns identified in deployment-oriented work on AI agents for discovery, where trust calibration depends on the legibility of the agent’s decision trail rather than on real-time supervision.

Limitations. Our experiments use only 2 replications per condition, yielding high uncertainty on some estimates (e.g., Sonnet $\Phi_{\min}$: $U=2.57/2.93$). The two SSDs are small gridworld environments;

scaling to larger, more complex games remains open. The researcher runs are not fully controlled—the number of outer iterations varies from 3 to 17 (mean 8.3)—making direct comparison across conditions approximate. The GEPA comparison uses a single run per condition; additional replications would strengthen confidence in the comparison.

Future work. Five directions emerge: (1) *new domains*—apply the two-level framework to other LLM-driven pipelines (code optimization, scientific experiment design, infrastructure tuning), where the same architecture applies with a different inner loop and objective Φ ; (2) *ablation*—freeze all configuration components except one to isolate the relative importance of prompt, feedback, helpers, and iteration logic; (3) *adversarial objectives*—set $\Phi = R_0$ (agent 0’s reward only) to test whether the researcher discovers exploitative pipeline configurations, connecting to the reward hacking analysis of [3]; (4) *cross-game transfer*—run the best Cleanup configuration on Gathering and vice versa, testing game-generalizability; (5) *heterogeneous policies*—extend to settings where agents run different code, enabling richer role specialization.

References

- [1] Joel Z Leibo, Vinicius Zambaldi, Marc Lanctot, Janusz Janssen, and Thore Graepel. Multi-agent reinforcement learning in sequential social dilemmas. In *Proc. 16th Conference on Autonomous Agents and MultiAgent Systems*, pages 464–473, 2017.
- [2] Lucian Buşoniu, Robert Babuška, and Bart De Schutter. A comprehensive survey of multi-agent reinforcement learning. *IEEE Transactions on Systems, Man, and Cybernetics, Part C (Applications and Reviews)*, 38(2):156–172, 2008.
- [3] Víctor Gallego. Cooperation and exploitation in LLM policy synthesis for sequential social dilemmas. In *Proc. Adaptive and Learning Agents Workshop (ALA 2026)*, 2026.
- [4] Andrej Karpathy. autoresearch: AI agents running research on single-GPU nanochat training automatically. <https://github.com/karpathy/autoresearch>, 2026.
- [5] Edward Hughes, Joel Z Leibo, Matthew Phillips, Karl Tuyls, Edgar A Dueñez-Guzmán, Antonio García Castañeda, Iain Dunning, Tina Zhu, Kevin R McKee, R. Koster, Heather Roff, and Thore Graepel. Inequity aversion improves cooperation in intertemporal social dilemmas. In *Neural Information Processing Systems*, 2018.
- [6] Julien Perolat, Joel Z Leibo, Vinicius Zambaldi, Charles Beattie, Karl Tuyls, and Thore Graepel. A multi-agent reinforcement learning model of common-pool resource appropriation. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- [7] Vincent Conitzer and Tuomas Sandholm. Complexity of mechanism design. In *Proc. 18th Conference on Uncertainty in Artificial Intelligence*, pages 103–110, 2002.
- [8] Emir Kamenica and Matthew Gentzkow. Bayesian persuasion. *American Economic Review*, 101(6):2590–2615, 2011.
- [9] Lakshya A Agrawal et al. GEPA: Reflective prompt evolution can outperform reinforcement learning. *arXiv preprint arXiv:2507.19457*, 2026.
- [10] Bernardino Romera-Paredes, Mohammadamin Barekatain, Alexander Novikov, Matej Balog, M Pawan Kumar, Emilien Dupont, Francisco JR Ruiz, Jordan S Ellenberg, Pengming Wang, Omar Fawzi, et al. Mathematical discoveries from program search with large language models. *Nature*, 625(7995):468–475, 2024.
- [11] Yecheng Jason Ma, William Liang, Guanzhi Wang, De-An Huang, Osbert Bastani, Dinesh Jayaraman, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Eureka: Human-level reward design via coding large language models. In *The Twelfth International Conference on Learning Representations*, 2024.
- [12] Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *Transactions on Machine Learning Research*, 2024.

- [13] Jacky Liang, Wenlong Huang, Fei Xia, Peng Xu, Karol Hausman, Brian Ichter, Pete Florence, and Andy Zeng. Code as policies: Language model programs for embodied control. In *2023 IEEE International Conference on Robotics and Automation (ICRA)*, pages 9493–9500, 2023.
- [14] Haoran Ye, Jiarui Wang, Zhiguang Cao, Federico Berto, Chuanbo Hua, Haeyeon Kim, Jinkyoo Park, and Guojie Song. ReEvo: Large language models as hyper-heuristics with reflective evolution. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [15] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: language agents with verbal reinforcement learning. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.
- [16] Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. Self-refine: Iterative refinement with self-feedback. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.
- [17] Chengrun Yang, Xuezhi Wang, Yifeng Lu, Hanxiao Liu, Quoc V Le, Denny Zhou, and Xinyun Chen. Large language models as optimizers. In *The Twelfth International Conference on Learning Representations*, 2024.
- [18] Taiwei Shi, Sihao Chen, Bowen Jiang, Linxin Song, Longqi Yang, and Jieyu Zhao. Experiential reinforcement learning. *arXiv preprint arXiv:2602.13949*, 2026.

A Additional Results

Table 3: Strategies discovered by the researcher in Cleanup across 4 efficiency-optimized and 4 maximin-optimized independent runs.

Strategy	Φ_U (4 runs)	$\Phi_{\min}$ (4 runs)
Waste-counting helpers	4/4	4/4
Zone/lane partitioning	3/4	4/4
Anti-regression feedback	3/4	2/4
Worked policy examples	2/4	3/4
Cleaning cost economics	2/4	3/4
Time-based role rotation	0/4	4/4

B Compute Requirements

Table 4 reports wall-clock time for all 12 autonomous runs and the 6 GEPA comparison runs. *Inner loop* measures total time spent on policy LLM generation and environment simulation across all evaluations; *total* (estimated from run-directory timestamps) additionally includes the researcher agent’s analysis, code editing, and decision-making between evaluations.

Cost breakdown. Each inner loop evaluation involves $K=2-3$ policy LLM generation calls (each $\sim 2k$ input tokens, $\sim 1.5k$ output tokens) followed by multi-seed simulation (5–12 seeds, $\sim 15-30s$ total). Policy LLM generation dominates inner loop time (86–97%), with Sonnet evaluations taking $\sim 2\times$ longer than Gemini due to extended thinking. The researcher agent (Claude Opus 4.6, running via Claude Code CLI) accounts for 41% of total wall-clock time (25.6h of the 62.2h total across all 12 runs), spent reading results, editing pipeline source files, and planning modifications. In monetary terms, the researcher agent is the dominant cost: each outer iteration consumes $\sim 50-100k$ context tokens in the Opus session, whereas each inner loop evaluation uses only $\sim 5-10k$ policy LLM tokens. The 6 GEPA comparison runs required 5.0h of compute with no researcher agent, operating at lower cost per run but achieving substantially lower performance (Table 2).




Figure 3: Maximin ($\min_i R_i$) across researcher iterations for the 4 Cleanup maximin-optimized runs. All runs transform deeply negative baselines (worst-off agents losing reward) into substantially positive values. Gemini reaches ~ 290 while Sonnet reaches ~ 160 – 200 . The dashed line marks $\min_i R_i = 0$ (no agent loses reward overall).




Figure 4: Final metrics across all conditions. (a) Efficiency: all conditions converge to $U \approx 2.5$ – 3.2 . (b) Equality: maximin-optimized Cleanup runs achieve $E \approx 1.0$, while efficiency-optimized runs show $E \approx 0.5$ – 0.7 ; Gathering achieves high equality regardless. (c) Maximin: the sharpest contrast—efficiency optimization leaves the worst-off agent deeply negative ($\min_i R_i \approx -200$), while maximin optimization transforms it to $+290$ (Gemini) or $+179$ (Sonnet). Gathering achieves high maximin (~ 580) under efficiency optimization alone. Error bars show s.d. across runs.

418 C Smaller Open-Weight Model: Gemma 4 26B

419 We test the framework with Gemma 4 26B-A4B-IT (Google), a 26B-parameter open-weight model
 420 substantially smaller than the frontier models in the main experiments. We run three experiments—
 421 two on Cleanup ($N=10$) optimizing efficiency and maximin respectively, and one on Gathering
 422 ($N=4$) optimizing efficiency—all with Opus 4.6 as the researcher $\mathcal{R}$.

423 **Setup.** In all three experiments, Gemma starts from complete failure: the baseline pipeline pro-
 424 duces $U=-10.0$ on Cleanup (agents CLEAN-spam without collecting apples) and $U=0.0$ on Gath-
 425 ering (broken BFS calls), compared to $U=1.93/0.86$ (Gemini/Sonnet on Cleanup) and $U=2.04/0.03$
 426 (Gemini/Sonnet on Gathering). The researcher ran $J=10$ – 18 outer iterations per condition.

427 **Results.** Table 5 presents the best configurations discovered. Under efficiency optimization, the
 428 researcher achieves $U=0.87$ —a dramatic recovery from the broken baseline but far below fron-

Table 4: Wall-clock time per autonomous run. Evals: total inner loop evaluations including initial baseline. Per eval: mean wall-clock per single evaluation. Total: end-to-end including researcher overhead.

Policy LLM	Φ	Evals	Inner loop (h)	Per eval (min)	Total (h)
<i>Cleanup (N=10)</i>					
Gemini	Φ_U	11	2.6	14	3.7
Gemini	Φ_U	18	4.1	14	6.4
Sonnet	Φ_U	4	2.1	32	6.1
Sonnet	Φ_U	7	5.0	43	6.1
Gemini	$\Phi_{\min}$	17	3.0	11	4.3
Gemini	$\Phi_{\min}$	11	3.6	20	5.0
Sonnet	$\Phi_{\min}$	8	4.5	33	8.8
Sonnet	$\Phi_{\min}$	9	5.3	36	13.2
<i>Gathering (N=4)</i>					
Sonnet	Φ_U	3	1.7	33	2.2
Gemini	Φ_U	5	1.4	17	1.9
Gemini	Φ_U	3	0.9	19	1.1
Sonnet	Φ_U	4	2.3	35	3.4
All 12 runs		100	36.6	22	62.2
GEPA (6 runs)		6	5.0	50	—

429 tier models ($U \approx 3.2$). The researcher compensates for the model’s limited capability by reducing
 430 inner-loop iterations to $K=1$ (avoiding the regression that plagued $K \geq 2$ runs) and providing highly
 431 structured worked examples with explicit cleaning-role assignment.

Table 5: Autoresearch with Gemma 4 26B-A4B-IT. Single run per condition. Baseline: pipeline from [3].

Game	Target	U	E	$\min_i R_i$	Keep rate
Cleanup	Baseline	-10.0	1.00 [†]	-1000	—
	Φ_U	0.87	-1.11	-371	6/19
	$\Phi_{\min}$	1.71	0.94	137	9/17
Gathering	Baseline	0.0	1.00 [†]	0	—
	Φ_U	2.44	0.98	580	4/11

[†]Equality is trivially 1.0 because all agents receive near-zero reward.

432 **Maximin optimization rescues efficiency.** Strikingly, under maximin optimization the researcher
 433 achieves *higher* efficiency ($U=1.71$) than under direct efficiency optimization ($U=0.87$), alongside
 434 strong equality ($E=0.94$) and positive worst-off welfare ($\min_i R_i=137$). This reversal—absent in
 435 frontier models, where both objectives yield $U \approx 3.1$ – 3.2 —occurs because the maximin objective
 436 forces the researcher toward coordination mechanisms (rotating cleaning duties, index-based apple
 437 assignment) that simultaneously improve collective output. Under efficiency-only optimization, the
 438 researcher converges on a local optimum—static role assignment with 25% dedicated cleaners—that
 439 a 26B model can implement but that caps performance well below the frontier.

440 This suggests that *for models below a capability threshold, fairness objectives may serve as better*
 441 *optimization targets for overall social welfare than direct efficiency maximization.* The structured
 442 coordination enforced by the maximin objective provides a scaffold that compensates for the weaker
 443 model’s difficulty implementing complex strategies from strategic hints alone.

444 **Gathering: full recovery on a simpler game.** On Gathering ($N=4$), the researcher brings
 445 Gemma from $U=0.0$ to $U=2.44$ with $E=0.98$ and $\min_i R_i=580$, after 10 outer iterations (4 kept).
 446 These values nearly match frontier models (Gemini: $U=2.49$, Sonnet: $U=2.52$; Table 2). The
 447 researcher discovers the same Voronoi partitioning and respawn-aware camping strategies found
 448 in the main experiments, and increases inner-loop iterations to $K=5$ to allow sufficient refinement.
 449 The contrast with Cleanup—where Gemma peaks at $U=0.87/1.71$ versus frontier $U \approx 3.2$ —suggests

that the researcher can fully compensate for model weakness on simpler coordination tasks (4 agents, symmetric costs) but not on harder provision dilemmas (10 agents, asymmetric costs).

D Researcher-Authored Pipeline Artifacts

The previous appendix shows the policies π_c^* that the inner loop *outputs*; this one shows the configuration $c = (p, \phi, \mathcal{H}, \iota, v)$ (Section 4, Table 1) that the researcher $\mathcal{R}$ *authored* to produce them. Excerpts are taken from the final commit on the dedicated git branch of the corresponding run; we reproduce the prose verbatim, with author commentary in [brackets] and elisions marked We organize by artifact type so each subsection makes a within-type contrast (e.g., Φ_U vs. $\Phi_{\min}$, or Cleanup vs. Gathering).

D.1 Helper library $\mathcal{H}$: coordination primitives

The researcher adds primitives that the policy LLM can call as black boxes, sparing it from re-implementing tricky logic on every iteration. Two patterns dominate: (i) *state inspection* (waste counts, fractions, beam-yield scoring) and (ii) *coordination primitives* (zone assignment, role rotation). We show one of each.

Listing 1: Cleanup, $\Phi_{\min}$ – band-based apple zoning helper added by the researcher in exp5. This is the primitive that lets the policy in Listing 6 keep collectors within their own row band, preventing all 7 gatherers from racing to the same apple.

```

464 def get_my_apples(env, agent_id):
465     """Alive apples in this agent's horizontal row band.
466     Divides the orchard into n_agents bands; falls back to all apples if empty."""
467     aid = int(agent_id)
468     n = int(env.n_agents)
469     band_h = env.height / n
470     band_lo, band_hi = aid * band_h, (aid + 1) * band_h
471
472     my_apples, all_apples = set(), set()
473     for i in range(env.n_apples):
474         if env.apple_alive[i]:
475             ar = int(env._apple_pos[i, 0]); ac = int(env._apple_pos[i, 1])
476             all_apples.add((ar, ac))
477             if band_lo <= ar < band_hi:
478                 my_apples.add((ar, ac))
479     return my_apples if my_apples else all_apples
480

```

Listing 2: Gathering – BFS-Voronoi territory + respawn-aware waiting helpers added by the researcher in gather-exp3. These are the primitives that Listing 8 composes into a one-line policy: “go to my_zone_apples, otherwise nearest_respawning_apple.”

```

482 def voronoi_zones(env):
483     """Multi-source BFS Voronoi over walkable cells.
484     Returns (row, col) -> agent_id; ties broken by lower agent_id."""
485     queue, visited = deque(), {}
486     for a_id in range(env.n_agents):
487         if int(env.agent_timeout[a_id]) == 0:
488             ar = int(env.agent_pos[a_id][0]); ac = int(env.agent_pos[a_id][1])
489             queue.append((ar, ac))
490             visited[(ar, ac)] = a_id
491     while queue:
492         r, c, a_id = queue.popleft()
493         for dr, dc in [(-1,0), (1,0), (0,-1), (0,1)]:
494             nr, nc = r + dr, c + dc
495             if 0 <= nr < env.height and 0 <= nc < env.width \
496                 and not env.walls[nr, nc] and (nr, nc) not in visited:
497                 visited[(nr, nc)] = a_id
498                 queue.append((nr, nc, a_id))
499     return visited
500
501 def nearest_respawning_apple(env, agent_id, zones=None, max_timer=10):
502     """Nearest dead apple in MY zone respawning within max_timer steps.
503     Returns (row, col) of best wait spot, or None."""
504     if zones is None: zones = voronoi_zones(env)
505     best_pos, best_t, best_d = None, max_timer + 1, float('inf')
506     ar = int(env.agent_pos[agent_id][0]); ac = int(env.agent_pos[agent_id][1])
507     for i in range(env.n_apples):
508         if not env.apple_alive[i]:
509             pos = (int(env._apple_pos[i][0]), int(env._apple_pos[i][1]))
510

```

```

511 29         if zones.get(pos) == agent_id:
512 30             t = int(env.apple_timer[i])
513 31             if t <= max_timer:
514 32                 d = abs(pos[0] - ar) + abs(pos[1] - ac)
515 33                 if t < best_t or (t == best_t and d < best_d):
516 34                     best_pos, best_t, best_d = pos, t, d
518 35     return best_pos

```

519 The choice of helper depends on the dilemma’s geometry: row-bands suffice when fairness is the
520 main concern (Cleanup maximin) but full BFS-Voronoi is needed when wall-aware territory owner-
521 ship matters (Gathering). The researcher does not fix this in advance — it picks the simpler primitive
522 whenever it works.

523 D.2 System prompt p : from neutral framing to strategic briefing

524 The unmodified system prompt is a neutral API description: “Write a policy that maximizes per-
525 agent reward.” Under maximin optimization, the researcher rewrites the prompt into a multi-section
526 strategic briefing. The excerpt below shows the pieces it added to the Cleanup prompt in exp5; the
527 full prompt grew from 165 to 325 lines.

Listing 3: Cleanup, $\Phi_{\min}$ – excerpts from the researcher-rewritten system prompt (exp5). The unmodified baseline contained only API documentation; everything below was added by $\mathcal{R}$.

```

528 ## CRITICAL OBJECTIVE: Rawlsian Fairness (Maximin)
529
530 Your goal is to maximize the minimum per-agent total return across all
531 agents. This is the "maximin" or Rawlsian welfare criterion.
532
533 This means: it is NOT enough for the *average* agent to do well. The
534 worst-off agent must do as well as possible. A policy where 8 agents
535 each earn +200 but 2 agents each earn -100 is TERRIBLE (maximin = -100).
536
537 Key implication: cleaning costs (-1 per CLEAN action) must be shared
538 equitably among ALL agents. If you assign fixed "cleaner" roles, those
539 agents will accumulate large negative rewards and destroy your maximin score.
540
541 ## Strategy for Maximin: ROLE ROTATION + APPLE ZONING
542
543 The optimal strategy for maximin involves:
544 1. Shared cleaning duty: ALL agents take turns cleaning using a rotation
545    schedule based on 'agent_id' and 'env._step_count'. For example:
546    'is_my_cleaning_turn = (agent_id + env._step_count // SHIFT) % env.n_agents < NUM_CLEANERS'
547    where SHIFT is ~50 steps and NUM_CLEANERS is 2-3.
548 2. When NOT your turn: collect apples from YOUR ZONE using
549    'get_my_apples(env, agent_id)' -- prevents all gatherers competing
550    for the same nearest apple.
551 3. NEVER use BEAM (action 6): it costs -1 and causes -50 to the target.
552 4. Keep waste density low: aim for ~2-3 active cleaners at any time.
553
554 [... API documentation, helper documentation, working example ...]
555
556 IMPORTANT:
557 - NEVER assign permanent cleaning roles to specific agents -- this kills maximin.
558 - Use env._step_count for time-based role rotation so all agents share cleaning duty.
559 - NEVER use BEAM (action 6) -- it destroys both agents' rewards.
560

```

562 Two things stand out. First, the researcher writes the formula it expects $\mathcal{M}$ to use almost verbatim —
563 `(agent_id + env._step_count // SHIFT) % env.n_agents < NUM_CLEANERS` — and the
564 policy in Listing 6 uses essentially this template. Second, the researcher *teaches by counter-example*:
565 the explicit “8 agents earn +200, 2 earn -100, maximin = -100” makes the failure mode of Φ_U -style
566 static roles (Listing 5) concretely visible to $\mathcal{M}$. This is the information-design move predicted by
567 the mechanism-design framing in Section 3.4.

568 The Gathering prompt evolves along a different axis (no maximin, no rotation). Its researcher-added
569 “Key Strategic Insights” section instead emphasizes (i) *self-play implies never beam* and (ii) *Voronoi*
570 *partition each step*, exactly matching what the policy in Listing 8 implements.

571 D.3 Feedback ϕ : adaptive diagnostics

572 The baseline feedback function from [3] shows reward + four social metrics with definitions and
573 stops there. The researcher turns it into a state machine: it inspects the latest metrics and condi-

tionally injects different hints. Two different runs produced two qualitatively different diagnostics — the same metric ($\bar{r}_k$ or $\min_i R_i$) is repurposed to either *stabilize* a working policy or *redirect* a failing one.

Listing 4: Adaptive feedback diagnostics. *Top*: efficiency-optimized run (exp1) – a stability guard prevents iter-3 regression once U is high. *Bottom*: maximin-optimized run (exp5) – two fairness diagnostics that fire when the worst-off agent loses reward.

```

577 1 # --- Cleanup, Phi_U feedback (exp1, pipeline/feedback.py final) ---
578 2 last_eff = history[-1]["metrics"]["efficiency"]
579 3 if last_eff >= 2.5:
580 4     parts.append(
581 5         """**CRITICAL -- DO NOT REGRESS**: The current policy achieves high "
582 6         "efficiency (>=2.5). You MUST output a policy that is nearly identical "
583 7         "to the current one. Copy the current policy code and make AT MOST "
584 8         "one small targeted improvement (e.g., adjust a single numeric "
585 9         "threshold). If you are not confident a change will help, output the "
586 10        "current policy UNCHANGED. A regression here means the run fails.")
587 11
588 12 # --- Cleanup, Phi_min feedback (exp5, pipeline/feedback.py final) ---
589 13 last_maximin = history[-1]["metrics"].get("maximin", 0)
590 14 last_avg = history[-1]["reward_avg"]
591 15 if last_maximin < 0:
592 16     parts.append(
593 17         f"""**FAIRNESS ALERT**: maximin={last_maximin:.1f} is NEGATIVE. "
594 18         f"The worst-off agent lost reward overall while average was {last_avg:.1f}. "
595 19         "This means cleaning duties are NOT shared equitably. "
596 20         "Use time-based role rotation (env._step_count) so ALL agents share "
597 21         "cleaning costs equally. NEVER assign permanent cleaner roles.")
598 22 elif last_maximin < last_avg * 0.5:
599 23     parts.append(
600 24         f"""**FAIRNESS WARNING**: maximin={last_maximin:.1f} is much lower than "
601 25         f"average={last_avg:.1f}. The gap suggests unequal cleaning burden. "
602 26         "Ensure ALL agents rotate through cleaning duty.")
603

```

The two diagnostics are written by independent researcher runs but converge on a common pattern: *thresholded interventions on a primary metric*. They are also a common cause of the low run-to-run variance reported in Finding 1: when a working policy lands in the high- U basin, the stability guard prevents the LLM from over-refining and tipping out of it (the $K=3$ regressions visible in Section 4’s “Common failure modes” (4)), and when an unfair policy lands in the negative-maximin region, the alert injects exactly the rotation hint that recovers it.

D.4 Iteration logic ι : per-condition hyperparameters

Although ι has the smallest source surface, the researcher’s edits to it explain a meaningful fraction of the variance reduction noted in Finding 1. Table 6 summarizes the values shipped in the final commit of each best-performing run.

Table 6: Iteration-logic (ι) values in the final commit of selected runs. K = inner-loop iterations, $|S|$ = evaluation seeds, “thinking” = extended-thinking token budget passed to $\mathcal{M}$. Baseline ([3]) is $K=3$, $|S|=5$, thinking 16k.

Run	Game / objective	K	$ S $	thinking	Researcher’s stated rationale
exp1 (Gemini)	Cleanup, Φ_U	3	5	16k	default; stability comes from feedback hint
exp4 (Sonnet)	Cleanup, Φ_U	3	5	32k	“Sonnet was over-refining at 16k thinking”
exp5 (Gemini)	Cleanup, $\Phi_{\min}$	2	12	16k	“ $K=2$ avoids iter-3 regression; 12 seeds for variance control”
exp7 (Sonnet)	Cleanup, $\Phi_{\min}$	3	5	10k	“reduced thinking from 16k – Sonnet was generating runaway code at 16k”
gather-exp3 (Gemini)	Gathering, Φ_U	3	5	16k	default; Gathering converges in $K \leq 3$

The maximin Gemini run is the most informative case: the researcher discovered that with $K=3$ and $|S|=5$, identical configurations could yield $\min_i R_i=295$ on one set of seeds and $\min_i R_i=100$ on another (this is the “variance exposure” failure mode of Section 4’s common failures). It then walked $|S|$ from 5 $\rightarrow$ 8 $\rightarrow$ 12 over three outer iterations until the seed-averaged signal was stable enough that the policy LLM stopped chasing noise. The policy in Listing 6 is sampled at $|S|=12$.

620 D.5 Cross-artifact observations

- 621 • *Helpers do the algebra; prompts pick the strategy class.* The researcher uses helpers (1, 2) to put
622 non-trivial primitives at $\mathcal{M}$'s fingertips, and the prompt (3) to tell $\mathcal{M}$ which primitive to compose.
623 Neither is sufficient alone — prompts without supporting helpers produced policies that “mention
624 rotation” but failed to implement it; helpers without prompt updates often went unused.
- 625 • *Feedback is the only adaptive component.* p , $\mathcal{H}$, and ι are static within a run; ϕ is the only place
626 where the researcher gets to change what $\mathcal{M}$ sees during the inner loop, and it uses thresholded
627 diagnostics (Listing 4) to do so. This is what mostly drives the run-to-run variance reduction.
- 628 • *The same artifact has different content under each objective.* The Cleanup prompt under Φ_U omits
629 “rotation” and “maximin” entirely; the same prompt under $\Phi_{\min}$ devotes its entire opening section
630 to it. The configuration c is genuinely a function of the welfare objective Φ , not a fixed pipeline
631 parameterized by it — consistent with the information-design framing of Section 3.4.
- 632 • *The researcher’s edits are short and targeted.* The full diff between the baseline pipeline and the
633 best maximin Gemini configuration totals ~ 300 added lines spread over four files; ~ 80 of them
634 appear in the listings above. The remaining content is API documentation and worked examples,
635 included for completeness but not strategically novel.

636 E Selected Generated Policies

637 The code excerpts below are verbatim outputs of the policy synthesizer $\mathcal{M}$, taken from the final
638 inner-loop iteration of the best-performing run in each condition. To fit the page, we elide repeated
639 boilerplate (BFS scaffolding, fallback branches) with `...` and add commentary in [brackets];
640 nothing else has been edited. The listings illustrate the qualitative differences highlighted in Findings
641 2–4 of Section 4: *static* vs. *rotating* cleaning roles in Cleanup, and the *Voronoi* + *respawn-timer*
642 structure in Gathering. All four policies were independently generated by the LLM under different
643 seeds, models, and outer-loop runs.

644 E.1 Cleanup, Φ_U (Gemini): static interleaved roles + dynamic threshold

645 Best run: $U=3.25$, $E=0.61$, $\min_i R_i = -182$. The synthesizer assigns *permanent* cleaning roles
646 via $\text{agent_id} \bmod n < \tau$, where the threshold τ depends smoothly on the current waste fraction.
647 Higher pollution recruits more cleaners, but the *identity* of each cleaner is fixed — agents $0\text{--}\tau\text{--}1$
648 always pay the cost. This drives U to its ceiling but leaves equality at 0.61 and pushes the worst-off
649 agent to $\min_i R_i = -182$.

Listing 5: Cleanup, Φ_U – Gemini 3.1 Pro. Static role assignment with waste-tiered threshold (ver-
batim, abridged).

```

650 1 def policy(env, agent_id) -> int:
651 2     if int(env.agent_timeout[agent_id]) > 0:
652 3         return 7 # STAND while removed
653 4
654 5     waste_frac = waste_fraction(env)
655 6     n_agents = env.n_agents
656 7
657 8     # [Smooth dynamic threshold: more pollution -> more cleaners]
658 9     if waste_frac > 0.35: threshold_cleaners = n_agents
659 10    elif waste_frac > 0.25: threshold_cleaners = int(n_agents * 0.7)
660 11    elif waste_frac > 0.15: threshold_cleaners = int(n_agents * 0.5)
661 12    elif waste_frac > 0.05: threshold_cleaners = int(n_agents * 0.3)
662 13    else: threshold_cleaners = max(1, int(n_agents * 0.2))
663 14
664 15    # [STATIC role: cleaners are ALWAYS the lowest-id agents -- no rotation]
665 16    is_cleaner = (agent_id % n_agents) < threshold_cleaners
666 17
667 18    if is_cleaner:
668 19        best_orient, n_waste = find_best_clean_orientation(env, agent_id)
669 20        if n_waste >= 1:
670 21            cur_orient = int(env.agent_orient[agent_id])
671 22            if best_orient == cur_orient: return 8 # CLEAN
672 23            if (cur_orient + 1) % 4 == best_orient: return 5 # ROTATE_RIGHT
673 24            elif (cur_orient - 1) % 4 == best_orient: return 4 # ROTATE_LEFT
674 25            else: return 4
675 26        # [Competitor-aware BFS toward waste in own row band]
676 27        c_idx = agent_id % n_agents
677 28        target_r = int(((c_idx + 0.5) / threshold_cleaners) * env.height)

```

```

679 29         # ... pick waste cell minimizing dist + |row - target_r| + 50 *
680         # closer_cleaners
681 30         # ... bfs_toward and direction_to_action
682 31     else:
683 32         # GATHERER: row-banded sweep, competitor-aware
684 33         # assigned_row = ((g_idx + 0.5) / n_gatherers) * env.height
685 34         # ... pick alive apple minimizing dist + |row - assigned_row| + 50 *
686         # closer_gatherers
687 35     ...

```

689 E.2 Cleanup, $\Phi_{\min}$ (Gemini): time-rotated, geographically interleaved roles

690 Best run: $U=3.19$, $E=0.98$, $\min_i R_i=296$. Switching the outer objective from U to $\min_i R_i$ flips
691 the role-assignment scheme entirely: `role_idx` now depends on `agent_id + step/50`, so every
692 agent rotates through cleaner and gatherer duty in a 50-step cycle. The cleaner indices $\{1, 4, 8\}$ are
693 spatially *interleaved* along the river (rather than contiguous) so that whoever is on duty has a short
694 walk to their assigned river slice. This single change accounts for the $E=0.55 \rightarrow 0.98$ jump and the
695 $\min_i R_i=-182 \rightarrow 296$ swing reported in Table 2, with negligible efficiency loss.

Listing 6: Cleanup, $\Phi_{\min}$ – Gemini 3.1 Pro. Time-rotated roles with interleaved geography (verbatim, abridged).

```

696 1 def policy(env, agent_id) -> int:
697 2     if int(env.agent_timeout[agent_id]) > 0:
698 3         return 7 # STAND while removed
699 4
700 5
701 6     aid = int(agent_id)
702 7     step = int(env._step_count)
703 8     n = int(env.n_agents)
704 9
705 10    # [TIME ROTATION: every 50 steps each agent's role index advances]
706 11    if n == 10:
707 12        role_idx = (aid + step // 50) % 10
708 13        # [Cleaner indices {1,4,8} are SPATIALLY INTERLEAVED, not contiguous,
709 14        # so each on-duty cleaner has a short walk to its river slice]
710 15        if role_idx == 1: role_type, zone_idx = 'C', 0
711 16        elif role_idx == 4: role_type, zone_idx = 'C', 1
712 17        elif role_idx == 8: role_type, zone_idx = 'C', 2
713 18        else:
714 19            role_type = 'G'
715 20            g_map = {0:0, 2:1, 3:2, 5:3, 6:4, 7:5, 9:6}
716 21            zone_idx = g_map[role_idx]
717 22            num_cleaners, num_gatherers = 3, 7
718 23
719 24    if role_type == 'C':
720 25        # [Each on-duty cleaner owns one horizontal slice of the river]
721 26        row_min = int( zone_idx * env.height / num_cleaners)
722 27        row_max = int((zone_idx + 1) * env.height / num_cleaners) - 1
723 28
724 29        # [Score every (row, col, orientation) in the slice by beam yield;
725 30        # fire if facing >=2 waste, else move to a >=2-yield position]
726 31        best_positions, good_positions = set(), set()
727 32        for r in range(row_min, row_max + 1):
728 33            for c in range(10):
729 34                if env.walls[r, c] or env.waste[r, c]: continue
730 35                y_max = max(get_clean_yield(env, r, c, o) for o in range(4))
731 36                if y_max >= 2: best_positions.add((r, c))
732 37                if y_max >= 1: good_positions.add((r, c))
733 38            # ... rotate / CLEAN / bfs_to_target_set as appropriate
734 39        else:
735 40            # GATHERER: collect apples ONLY inside the rotating row band
736 41            # row_min/row_max for num_gatherers slices ...
737 42        ...

```

739 E.3 Cleanup, $\Phi_{\min}$ (Sonnet): independent rediscovery of duty rotation

740 Best run: $U=2.93$, $E=0.83$, $\min_i R_i=154$. Run with a different policy LLM (Sonnet 4.6) and on
741 a different outer-loop trajectory, the synthesizer arrives at the *same* structural insight as the Gemini
742 maximin run: a phase counter (`agent_id + step/50`) mod n rotates which 2 of the 10 agents clean
743 at any given time, and a *separate* 200-step zone counter rotates which 5-row band each collector
744 sweeps. The two rotation periods (50 and 200 steps) ensure every agent visits every cleaning slot
745 and every apple zone within one episode — a structural fairness invariant the researcher discovered
746 without any explicit instruction to “rotate.”

Listing 7: Cleanup, $\Phi_{\min}$ – Sonnet 4.6. Phase-rotated cleaning + zone-rotated collection (verbatim, abridged).

```

747 1 def policy(env, agent_id) -> int:
748 2     if int(env.agent_timeout[agent_id]) > 0:
749 3         return 7 # STAND while removed
750 4
751 5     n = env.n_agents
752 6     step = env._step_count
753 7     wf = waste_fraction(env)
754 8
755 9     # --- ROLE ASSIGNMENT: fair 2-cleaner rotation ---
756 10    # [50-step phases x 20 phases / episode -> each agent cleans 4 phases (~20%)]
757 11    phase = step // 50
758 12    cleaner_rank = (agent_id + phase) % n # [cycles 0..9 fairly]
759 13    is_cleaner = cleaner_rank < 2 # [exactly 2 cleaners per phase]
760 14
761 15    # [Emergency override: all agents clean when waste crosses the spawn cliff]
762 16    in_emergency = wf >= 0.40
763 17    if in_emergency: is_cleaner = True
764 18
765 19    # [River split: cleaner_slot 0 -> top half, 1 -> bottom half]
766 20    cleaner_slot = (agent_id % 2) if in_emergency else cleaner_rank
767 21
768 22    if is_cleaner and wf > 0.04:
769 23        # [Scan 4 orientations from current position; if best beam shot
770 24        # covers >=3 waste cells, fire; else step to a strictly better neighbor]
771 25        # ... (best_o, best_cnt) = max over 4 directions
772 26        # ... if best_cnt >= 3: rotate to best_o, then CLEAN
773 27        # ... else: scan 4 adjacent cells x 4 orientations for higher yield
774 28        # ... if no waste in beam range: bfs to waste_set in OWN river half,
775 29        # fall back to whole river
776 30
777 31    # --- COLLECTOR LOGIC: 5-zone rotating apple collection ---
778 32    # [5 zones x 200 steps = 1000 steps -> each agent visits every zone once;
779 33    # 2 agents share each zone at any time, reducing competition 10 -> 2]
780 34    bounds = zone_boundaries(env, 5)
781 35    zone_phase = step // 200
782 36    zone = (agent_id + zone_phase) % 5
783 37    z_start, z_end = bounds[zone], bounds[zone + 1]
784 38
785 39    zone_apples = {(int(env._apple_pos[i][0]), int(env._apple_pos[i][1]))
786 40                  for i in range(env.n_apples) if env.apple_alive[i]
787 41                  and z_start <= int(env._apple_pos[i][0]) < z_end}
788 42    # ... bfs_to_target_set(zone_apples), else bfs_nearest_apple as fallback
789 43    return 7
790

```

E.4 Gathering (Gemini): wall-aware Voronoi + spatiotemporal targeting

Best run: $U=2.47$, $E=0.98$, $\min_i R_i=580$. In Gathering, where costs are symmetric, no role differentiation is needed: the entire optimization is spatial and temporal. The synthesizer (i) partitions cells into BFS-Voronoi territories owned by individual agents, (ii) runs a single centralized BFS from its own position to score every reachable cell with exact wall-aware distances, and (iii) for both alive and respawning apples in its territory, computes the *earliest collection time* $\max(\text{walk_dist}, \text{respawn_timer})$ and targets the minimum. When the agent must wait on a dead apple, it steps onto a *non-spawn* adjacent cell rather than blocking a respawn point. This single sort key — $\max(\text{distance}, \text{timer})$ — subsumes both “go to nearest apple” and “camp respawn” as special cases.

Listing 8: Gathering – Gemini 3.1 Pro. Voronoi territories + spatiotemporal priority targeting (verbatim, abridged).

```

802 1 def policy(env, agent_id) -> int:
803 2     if int(env.agent_timeout[agent_id]) > 0:
804 3         return 7
805 4
806 5     r, c = int(env.agent_pos[agent_id][0]), int(env.agent_pos[agent_id][1])
807 6     orient = int(env.agent_orient[agent_id])
808 7
809 8     # [Wall-aware Voronoi over the gridworld; spawn_points = all apple cells]
810 9     zones = voronoi_zones(env)
811 10    spawn_points = {(int(p[0]), int(p[1])) for p in env._apple_pos}
812 11
813 12    # [Single centralized BFS: exact distance + first move to every reachable cell]
814 13    distances, first_moves = {}, {}
815 14    q = deque([(r, c, 0, None)])

```

```

817 15     visited = {(r, c)}
818 16     while q:
819 17         cr, cc, d, fm = q.popleft()
820 18         distances[(cr, cc)] = d
821 19         first_moves[(cr, cc)] = fm
822 20         for dr, dc in [(-1,0),(1,0),(0,-1),(0,1)]:
823 21             nr, nc = cr + dr, cc + dc
824 22             if 0 <= nr < env.height and 0 <= nc < env.width and not env.walls[nr,
825 23                 nc]:
826 24                 if (nr, nc) not in visited:
827 25                     visited.add((nr, nc))
828 26                     q.append((nr, nc, d + 1, fm if fm is not None else (dr, dc)))
829 27
830 28     # [Both alive AND respawning apples in MY zone, with their respawn timer]
831 29     my_zone_spawns = []
832 30     for i in range(env.n_apples):
833 31         pr, pc = int(env._apple_pos[i][0]), int(env._apple_pos[i][1])
834 32         if zones.get((pr, pc)) == agent_id and (pr, pc) in distances:
835 33             my_zone_spawns.append((pr, pc, int(env.apple_timer[i]), distances[(pr,
836 34                 pc)]))
837 35
838 36     # [SPATIOTEMPORAL PRIORITY: earliest collection time first.
839 37     # score = max(walk_distance, respawn_timer); ties broken by sooner timer]
840 38     my_zone_spawns.sort(key=lambda x: (max(x[3], x[2]), x[2], x[3]))
841 39
842 40     for pr, pc, timer, dist in my_zone_spawns:
843 41         if timer == 0:
844 42             # [Alive apple: walk straight to it]
845 43             if dist == 0: return 7
846 44             fm = first_moves[(pr, pc)]
847 45             return direction_to_action(fm[0], fm[1], orient)
848 46         else:
849 47             # [Dead apple: camp on safe adjacent cell]
850 48             # ... pick nearest neighbor (nr,nc) that is NOT in spawn_points,
851 49             # ... navigate there and STAND while waiting for respawn
852 50             ...
853 51     # [Global poaching fallback if zone is empty: nearest alive apple anywhere]
854 52     ...

```

854 E.5 Cross-condition observations

855 Several patterns are visible across these four listings (and confirmed by inspecting the remaining 8
856 runs not shown):

- 857 • *Role assignment encodes the welfare objective.* Static $\text{agent_id} < \tau$ (Listing 5) maximizes
858 U but harms $\min_i R_i$. Time-rotated $(\text{agent_id} + \text{step} // T) \% n$ (Listings 6, 7) maximizes
859 $\min_i R_i$ at $\leq 1\%$ efficiency cost (Gemini).
- 860 • *Convergent rediscovery across LLMs.* Gemini and Sonnet, on independent runs, both arrive at
861 the $(\text{id} + \text{phase}) \% n$ duty-rotation idiom with similar phase lengths ($T \approx 50$ steps) under
862 maximin — and both omit it under efficiency.
- 863 • *Game structure dictates strategy class.* Cleanup policies are dominated by role-assignment logic
864 (cleaner vs. gatherer); Gathering policies are dominated by territory partitioning and respawn-
865 timer reasoning, with no role differentiation at all. The researcher does not need to be told which
866 class of strategy to write.
- 867 • *Earliest-collection-time as a unifying score.* The Gathering policy’s $\max(\text{walk}, \text{timer})$ key (List-
868 ing 8) is structurally identical across the 4 Gathering runs (both Gemini and Sonnet), differing
869 only in how the Voronoi diagram is computed — Manhattan approximation, fully BFS-based, or
870 cached helper.