
Discovering Cooperative Pipelines: Autoresearch for Sequential Social Dilemmas

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 We study *two-level autoresearch* for cooperation: an outer-loop AI agent au-
2 tonomously redesigns the inner-loop pipeline of an LLM policy-synthesis system
3 for multi-agent Sequential Social Dilemmas (SSDs). A *researcher agent* $\mathcal{R}$ (run as
4 a coding agent) reads the inner-loop source code, edits system prompts, feedback
5 functions, helper libraries, and iteration logic, runs evaluations, and decides what
6 to keep, following the *autoresearch* paradigm. Across two games (Cleanup and
7 Gathering), two policy-synthesizer LLMs, and two welfare objectives (utilitarian
8 efficiency and Rawlsian maximin), the researcher reliably exceeds hand-designed
9 baselines, sharply tightens run-to-run variance, and outperforms prompt-only op-
10 timization. The discovered pipelines are objective-dependent: only under max-
11 imin does the researcher inject an explicit *fairness mechanism* into synthesizer
12 pipelines, a class of mechanism that is absent from its own objective-agnostic
13 system prompt and from every efficiency-optimized pipeline. This supports an
14 *information-design* reading in which the researcher chooses what to reveal to the
15 boundedly rational synthesizer as a function of the welfare objective. Code at
16 <https://github.com/anon590/autoresearch-social-dilemmas>.

17 1 Introduction

18 Sequential Social Dilemmas (SSDs) [1] are the multi-agent analogue of the prisoner’s dilemma in
19 temporally rich Markov games: individually rational play leads to collectively suboptimal outcomes
20 through pollution, over-harvesting, or open conflict. Standard multi-agent reinforcement learning
21 (MRL) struggles in this regime due to credit assignment, non-stationarity, and large joint action
22 spaces [2]. A complementary approach, recently introduced by Gallego [3], replaces parameter-
23 space optimization with *algorithm-space synthesis*: a frozen LLM writes a Python policy function,
24 evaluates it in self-play, and iteratively refines it from performance feedback. A single generation
25 step can produce coordination logic (territory partitioning, role assignment, conditional cooperation)
26 at a sample efficiency several orders of magnitude beyond what gradient-based MRL achieves on
27 the same environments.

28 This shifts where the design problem lives, rather than removing it. The inner-loop pipeline that
29 drives the synthesizer has many free parameters: which system prompt, which feedback variables,
30 which helper functions, how many refinement steps. Each materially affects the resulting policies,
31 and prior work tuned them by hand. A natural question follows: *can an AI agent design the pipeline?*

32 We answer affirmatively with a two-level autoresearch framework. An *outer-loop researcher agent*
33 (Claude Opus 4.6, run as a coding agent) edits the source files of an *inner-loop policy synthesizer*
34 (another LLM), runs evaluations on held-out seeds, and keeps modifications that improve a fixed
35 welfare objective Φ . The outer agent operates on an ordinary git repository (reading code, writing
36 diffs, running shell commands, etc) without task-specific scaffolding beyond a standard CLI and git,

37 mirroring the autoresearch paradigm of Karpathy [4] for single-GPU LLM pretraining. Although the
 38 inner-loop SSDs are gridworld benchmarks rather than physical systems, the outer-loop discovery
 39 process itself runs under conditions a deployed discovery agent faces: noisy multi-seed evaluations,
 40 stochastic code generation, an LLM-evaluation budget that bounds how often Φ can be queried, and
 41 a heterogeneous code repository the agent must navigate end-to-end on its own.

42 Our contributions are: i) a general two-level framework that delegates the design of an LLM syn-
 43 thesis pipeline to a coding agent operating on a real software repository (Section 3); ii) the first
 44 instantiation of the autoresearch paradigm in a multi-agent decision-making domain, with experi-
 45 ments across two SSDs (Cleanup and Gathering), two policy LLMs, and two welfare objectives
 46 (utilitarian efficiency U and Rawlsian maximin $\min_i R_i$) (Section 4); and iii) a mechanism-design
 47 interpretation supported by the qualitatively different pipelines the agent produces under different
 48 welfare objectives, including the autonomous insertion of explicit fairness mechanisms (usually *time-*
 49 *based duty rotation*) into the researcher-authored synthesizer prompts and helpers, in every maximin
 50 run and no efficiency run.

51 2 Background

52 We build on the iterative LLM policy synthesis framework of Gallego [3], which serves as the
 53 (frozen) inner loop of our two-level system. This section recalls the SSD formalism, the social
 54 outcome metrics, and the base synthesis loop.

55 2.1 Sequential Social Dilemmas

56 A Sequential Social Dilemma is a partially observable Markov game $\mathcal{G} =$
 57 $\langle N, \mathcal{S}, \{\mathcal{A}_i\}_{i=1}^N, T, \{R_i\}_{i=1}^N, H \rangle$ with N agents, state space $\mathcal{S}$ (the gridworld configuration),
 58 per-agent action spaces $\mathcal{A}_i$, transition function T , reward functions R_i , and episode horizon H [1].
 59 Beyond the dilemma’s matrix-game structure, SSDs add temporal richness: agents must learn *when*
 60 and *where* to cooperate, not just *whether* to.

61 We study two canonical SSDs that capture complementary dilemma types.

62 **Cleanup** [5] is a *public goods provision* game ($N=10$). The map has two regions: a river that
 63 accumulates waste, and an orchard where apples grow. Apples regrow only when river pollution is
 64 below threshold. Each agent can fire a cleaning beam (cost -1) that removes waste, collect apples
 65 ($+1$ each), or fire a tagging beam (cost -1 , inflicting -50 on the target and removing it for 25 steps).
 66 The dilemma: cleaning is costly but its benefits are public, so purely self-interested agents free-ride.

67 **Gathering** [1, 6] is a *common pool resource* game ($N=4$). Agents collect shared apples on a fixed
 68 respawn timer and may fire tagging beams to temporarily remove rivals. The dilemma: agents can
 69 coexist and share resources, or aggress to monopolize them; aggression wastes time and reduces
 70 total welfare.

71 Both games use 8–9 discrete actions (movement, rotation, beam, stand, optionally clean) and
 72 episodes of $H=1000$ steps. The two dilemmas differ in cost structure: asymmetric provision (clean-
 73 ers pay, all benefit) vs. symmetric restraint (every agent faces the same temptation), a distinction
 74 that drives our experimental findings.

75 **Social metrics.** Following Perolat et al. [6], let $R_i = \sum_{t=0}^{H-1} r_i^t$ denote agent i 's episode return.
 76 We evaluate four social outcomes:

$$U = \frac{1}{H} \sum_{i=1}^N R_i \quad (\text{efficiency}) \quad (1)$$

$$E = 1 - \frac{\sum_{i,j} |R_i - R_j|}{2N \sum_i R_i} \quad (\text{equality}) \quad (2)$$

$$S = \frac{1}{N} \sum_{i=1}^N \bar{t}_i \quad (\text{sustainability}) \quad (3)$$

$$P = \frac{1}{H} \sum_{t=0}^{H-1} |\{i : \text{active}_i^t\}| \quad (\text{peace}) \quad (4)$$

77 where $\bar{t}_i$ is the mean timestep at which agent i collects positive reward (higher $\Rightarrow$ resources preserved
 78 later) and active_i^t indicates that agent i is not tagged out at step t . We additionally consider the
 79 *maximin* (Rawlsian) welfare criterion $\min_i R_i$, which optimizes the worst-off agent's return and
 80 serves as the second objective in our experiments.

81 2.2 Iterative LLM Policy Synthesis

82 Let Π denote the space of *code-based policies*: deterministic functions $\pi : \mathcal{S} \times [N] \rightarrow \mathcal{A}$ expressed
 83 as executable Python code. Each policy has access to the full environment state and a library of
 84 helpers (BFS pathfinding, beam targeting, coordinate transforms). This state access is a deliberate
 85 design choice: programmatic policies operate in algorithm space rather than the reactive observation-
 86 to-action space of neural policies, which lets a single LLM generation step encode rich coordination
 87 logic.

88 A frozen LLM $\mathcal{M}$ acts as the *policy synthesizer*. Given a system prompt p describing the environ-
 89 ment API and a feedback prompt q_k , it produces a new policy

$$\pi_{k+1} = \mathcal{M}(p, q(\pi_k, \mathcal{F}_k)),$$

90 where π_k is the previous policy (its source code) and $\mathcal{F}_k$ is the evaluation feedback. All N agents
 91 execute the same program π_k in self-play. We stress that this is *symmetric*, not *behaviorally ho-*
 92 *mogeneous*: since π_k takes `agent_id` as an argument, a single shared program can induce distinct
 93 per-agent behaviors (cleaner vs. gatherer assignment, time-rotated duty cycles, partitioned territo-
 94 ries; see the synthesized policies in Appendix D). What is shared is the source code; the sampled
 95 action distributions can differ across agents. Evaluation over a set of random seeds S yields the
 96 mean per-agent return $\bar{r}_k$ and the social metrics vector $\mathbf{m}_k = (U_k, E_k, S_k, P_k)$. Each generated
 97 policy passes an AST-based safety check (blocking `eval`, file I/O, network access) followed by a
 98 short smoke test; failures trigger regeneration (up to R attempts) with the error message appended
 99 to the prompt.

100 **Feedback.** We package the previous policy's source code together with all available evaluation
 101 signals:

$$\mathcal{F}_k = (\text{code}(\pi_k), \bar{r}_k, \mathbf{m}_k, \mathbf{d}), \quad (5)$$

102 where $\mathbf{d}$ contains natural-language definitions of each social metric. The LLM consumes $\mathcal{F}_k$ to
 103 revise and improve the policy. This is a starting point; the choice of feedback content is a single
 104 design decision within a much larger pipeline configuration space: our two-level framework (Sec-
 105 tion 3) opens the full space to automated search.

106 3 Two-Level Framework

107 We introduce a two-level system where a *researcher agent* $\mathcal{R}$ autonomously discovers configurations
 108 that optimize the output of an inner-loop system. While we instantiate this for multi-agent policy
 109 synthesis, the architecture is general: any pipeline where an LLM generates artifacts, evaluates
 110 them, and iterates can serve as the inner loop. The fundamental insight is that the entire inner-loop
 111 codebase is a designable artifact that a code-based agent can search over. Figure 1 illustrates the
 112 architecture.




Figure 1: Two-level automated research framework (Algorithm 1).

Algorithm 1 Two-Level Automated Research

Require: Game $\mathcal{G}$, policy synthesizer $\mathcal{M}$, researcher $\mathcal{R}$, system prompt $p_{\mathcal{R}}$, initial config c_0 , outer iterations J , ground-truth objective Φ , held-out seeds S_{ho}

Ensure: Best configuration c^* , best policy π^*

```

1:  $\pi_0^* \leftarrow \text{INNERLOOP}(\mathcal{M}, \mathcal{G}, c_0)$ 
2:  $\Phi_0 \leftarrow \text{EVAL}(\pi_0^*; \mathcal{G}, S_{\text{ho}})$ 
3:  $\text{history} \leftarrow \{(c_0, \Phi_0, \mathbf{m}_0, \emptyset)\}$ 
4: for  $j = 1, \dots, J$  do
5:    $c_j \leftarrow \mathcal{R}(p_{\mathcal{R}}, \text{CODE}(c_{j-1}), \text{history})$ 
6:   if  $\neg \text{VALIDATECONFIG}(c_j)$  then  $\text{retry} \leq R$ 
7:      $\pi_j^* \leftarrow \text{INNERLOOP}(\mathcal{M}, \mathcal{G}, c_j)$ 
8:      $\Phi_j \leftarrow \text{EVAL}(\pi_j^*; \mathcal{G}, S_{\text{ho}})$ 
9:      $\Delta_j \leftarrow \text{DIFF}(c_{j-1}, c_j)$  // code diff
10:     $\text{history} \leftarrow \text{history} \cup \{(c_j, \Phi_j, \mathbf{m}_j, \Delta_j)\}$ 
11:  end for
12:  $j^* \leftarrow \arg \max_j \Phi_j$ 
13: return  $c_{j^*}, \pi_{j^*}^*$ 

```

Table 1: Configuration components modifiable by the researcher, with concrete edits that $\mathcal{R}$ made in our runs. The environment simulator, ground-truth evaluation, and policy LLM weights are frozen.

	Scope	Concrete edits by $\mathcal{R}$ (see appendix)
p	System prompt	Multi-section strategic briefing with explicit duty-rotation template (agent_id + step//T) % n (App. C.3, Listing 4)
ϕ	Feedback fn	Thresholded diagnostics: “FAIRNESS ALERT” fires when $\min_i R_i < 0$; “DO NOT REGRESS” guard fires when $U \geq 2.5$ (App. C.4, Listing 5)
$\mathcal{H}$	Helper library	BFS-Voronoi territories with respawn-timer awareness; band-based apple zoning by agent_id (App. C.2, Listings 3, 2)
ι	Iteration logic	Per-condition $K \in \{2, 3\}$; $ S $ walked 5→8→12 for variance control; thinking budget 10–32k (App. C.5, Table 6)

3.1 Configuration Space

Let $\mathcal{C}$ denote the space of *pipeline configurations*. Each configuration $c \in \mathcal{C}$ specifies the full inner-loop setup:

$$c = (p, \phi, \mathcal{H}, \iota) \quad (6)$$

where p is the system prompt, ϕ is the feedback construction function (which metrics and diagnostics to include, how to frame it, whether to inject adaptive hints and thresholds, etc), $\mathcal{H}$ is the helper function library (auxiliary functions for pathfinding, getting aggregates of useful environment quantities, etc.), and ι specifies the iteration logic (number of inner iterations K , sampling strategy). Table 1 provides concrete examples. The validation pipeline is part of the frozen inner-loop infrastructure rather than a configurable component, the researcher cannot modify it in our experiments to prevent reward hacking.

The hand-designed feedback of [3] corresponds to a single fixed instantiation of ϕ . Our framework opens the full configuration space to automated search.

3.2 Inner Loop (Policy Synthesis)

Given a configuration c , the inner loop executes K iterations of LLM policy synthesis:

$$\pi_c^* = \text{INNERLOOP}(\mathcal{M}, \mathcal{G}, c) \quad (7)$$

Each iteration k proceeds in four stages, following [3]:

1. **Synthesize.** The policy LLM $\mathcal{M}$ receives the system prompt p , the previous policy’s source code π_{k-1} , and feedback $\mathcal{F}_{k-1}$ constructed by ϕ . It generates a new Python policy function π_k that has access to full environment state and the helper library $\mathcal{H}$.

- 131 2. **Validate.** The generated code undergoes AST-based safety checks (blocking dangerous operations such as file I/O and network access) followed by a short smoke test. Failures trigger
 132 re-generation (up to R retries), with the error message appended to the prompt.
 133 3. **Evaluate.** All N agents execute the same policy π_k in self-play over $|S|$ random seeds (note the
 134 policy is conditional on `agent_id`). The evaluation yields the mean per-agent reward $\bar{r}_k$ and the
 135 social metrics vector $\mathbf{m}_k = (U_k, E_k, S_k, P_k)$.
 136 4. **Feedback.** The feedback function ϕ constructs the prompt for the next iteration from
 137 $(\pi_k, \bar{r}_k, \mathbf{m}_k)$, packaging the previous policy’s source code together with the scalar reward, the
 138 social metrics vector, their natural-language definitions, and any adaptive diagnostics that ϕ in-
 139 jects.
 140

141 The inner loop output is evaluated on held-out seeds via a fixed ground-truth objective:

$$\Phi(c) = \text{EVAL}(\pi_c^*; \mathcal{G}, S_{\text{held-out}}) \quad (8)$$

142 where Φ maps from social outcomes to a scalar. We consider two alternative objectives to maximize:

$$\Phi_U = U = \frac{1}{H} \sum_{i=1}^N R_i \quad (\text{utilitarian efficiency}) \quad (9)$$

$$\Phi_{\min} = \min_i R_i \quad (\text{Rawlsian maximin}) \quad (10)$$

143 Φ_U rewards collective throughput and is indifferent to how reward is distributed across agents,
 144 whereas $\Phi_{\min}$ instead optimizes for the worst-off agent, pressuring the researcher toward configu-
 145 rations that distribute the cost of cooperation.

146 3.3 Outer Loop (Automated Research)

147 The researcher agent $\mathcal{R}$ iteratively modifies the pipeline configuration. Following the autoresearch
 148 paradigm [4], $\mathcal{R}$ operates on the inner-loop codebase as a modifiable artifact, proposing changes,
 149 observing outcomes, and refining. The procedure is formalized in Algorithm 1.

150 At each outer iteration j , the researcher $\mathcal{R}$ receives: i) the **full source code** of the current config-
 151 uration c_{j-1} (prompts, feedback construction, helpers, iteration logic); ii) the **experiment history**:
 152 for each prior iteration, the code diff Δ_i , ground-truth score Φ_i , and social metrics vector $\mathbf{m}_i$; iii)
 153 the **environment source code** (read-only), enabling the researcher to reason about game mechanics.
 154 The researcher proposes a new configuration c_j by generating code modifications. Concretely, $\mathcal{R}$
 155 is a coding agent (Claude Code CLI) that operates on a real software repository: it reads and edits
 156 Python source files, runs shell commands, inspects evaluation outputs, and commits changes to a
 157 dedicated git branch, following the same workflow a human researcher would follow.

158 3.4 Connection to Automated Mechanism Design

159 The two-level structure admits a mechanism design interpretation. The researcher $\mathcal{R}$ acts as a *mech-*
 160 *anism designer*: it controls the information structure (what metrics to reveal, how to frame them),
 161 the action space (what helper functions are available), and the incentive structure (how feedback is
 162 presented) under which the policy synthesizer $\mathcal{M}$ operates. The synthesizer acts as the *agent* within
 163 the designed mechanism.

164 This connects to the automated mechanism design literature [7], where a principal designs rules to
 165 induce desired behavior from self-interested agents. In our setting:

- 166 • The **principal** is the researcher $\mathcal{R}$, optimizing the ground-truth objective Φ .
- 167 • The **agent** is the synthesizer $\mathcal{M}$, optimizing per-agent reward as instructed.
- 168 • The **mechanism** is the configuration c : prompts, feedback, helpers, iteration logic.
- 169 • The **outcome** is the social welfare of the resulting multi-agent policy π_c^* .

170 A crucial distinction from classical mechanism design: $\mathcal{M}$ follows instructions but has *bounded*
 171 *rationality* in the sense that its ability to synthesize effective policies depends on the information
 172 and tools provided. The researcher’s task is thus closer to *information design* [8]: choosing what
 173 to reveal to help $\mathcal{M}$ navigate the cooperation–defection tension. Our experiments (Section 4) show
 174 that the researcher designs qualitatively different information structures depending on the welfare
 175 objective Φ , supporting this interpretation empirically.

176 4 Experiments

177 4.1 Setup

178 We conduct 12 autonomous researcher runs across a factorial design. The researcher agent $\mathcal{R}$ is
179 Claude Opus 4.6, invoked via the Claude Code CLI as a coding agent. Each run operates on a dedi-
180 cated git branch of a real Python codebase: $\mathcal{R}$ edits source files in `pipeline/`, executes evaluation
181 scripts, reads metric outputs, and iterates, without human intervention.

182 **Design.** For Cleanup: 2 policy LLMs $\times$ 2 objectives $\times$ 2 replications = 8 runs. For Gathering: 2
183 policy LLMs $\times$ 1 objective $\times$ 2 replications = 4 runs. Maximin runs are unnecessary for Gathering
184 because efficiency optimization alone achieves close to perfect equality.

185 **Models.** Policy synthesizer $\mathcal{M}$: Gemini 3.1 Pro (Google) or Claude Sonnet 4.6 (Anthropic), to
186 recent state-of-the-art LLMs. Both use extended thinking. We additionally test with Gemma 4 26B-
187 A4B-IT (Google), a smaller open-weight model, to probe the framework’s behavior when the policy
188 synthesizer has substantially lower capability (Appendix B.1).

189 **Baselines.** The hand-designed feedback configuration from prior work [3] serves as the initial
190 pipeline c_0 for all runs. On Cleanup ($N=10$), this baseline achieves $U=1.93/2.70$ (mean/max) with
191 Gemini and $U=0.86/1.56$ with Sonnet. We additionally compare against GEPA [9], an automated
192 prompt optimization method that iteratively refines the system prompt via LLM reflection. GEPA
193 optimizes only the system prompt p , whereas our researcher modifies the full pipeline $(p, \phi, \mathcal{H}, \iota)$.
194 We give GEPA a matched compute budget: the same number of optimization steps as outer iterations
195 used by our method, so all in all both methods use a comparable number of environment evaluations.

196 4.2 Results

197 Table 2 presents the main results, comparing our autoresearch framework against the hand-designed
198 baseline of [3] and GEPA [9].

199 **Finding 1: The researcher reliably improves over hand-designed baselines and outperforms**
200 **prompt-only optimization.** Every run improves substantially regardless of starting point (Fig-
201 ure 2). On Cleanup, autoresearch lifts both LLMs to $U \approx 3.1$ – 3.2 from baselines of 1.93 (Gemini)
202 and 0.86 (Sonnet), nearly closing the gap between them; on Gathering, all four runs converge to
203 $U \in [2.47, 2.52]$ from baselines spanning 0.03–2.42. Run-to-run spread is tight (gaps within 0.05
204 on Cleanup- Φ_U), suggesting the researcher reliably finds the performance ceiling of each policy
205 LLM via the pipeline modifications it discovers (helpers, prompts, and feedback; see appendices
206 C.2–C.4 for a selection of them). At matched environment queries, autoresearch beats GEPA by
207 2–3 $\times$ on Cleanup (both Φ_U and $\Phi_{\min}$) and 20% on Gathering, with the gap widening for the weaker
208 policy LLM: GEPA-Sonnet under $\Phi_{\min}$ can collapse to a pathological “everyone cleans, nobody
209 eats” regime ($U=-2.87$, $E=1.00$), while autoresearch-Sonnet reaches $\min_i R_i \approx 200$ reliably.
210 Modifying the full pipeline, not just the prompt, is what closes these gaps.

211 **Finding 2: No efficiency–fairness tradeoff in Cleanup (Gemini).** Maximin-optimized Gem-
212 ini pipelines sacrifice only 1% efficiency (U : 3.16 vs. 3.20) while achieving near-perfect equal-
213 ity ($\bar{E}$: 0.98 vs. 0.55) and transforming maximin from deeply negative baselines (−99 to −84)
214 to $\min_i R_i = 290$ (Figure 3). The researcher discovers *fair duty rotation* (Listing 7; primed by
215 the prompt rewrite of Listing 4)—time-based cycling using `agent_id` and `env._step_count`—
216 simultaneously improving worst-off welfare *and* collective output. Because cleaning is a public
217 good, distributing the cleaning cost fairly ensures enough cleaners to sustain apple production.

218 For Sonnet, there is a moderate tradeoff: efficiency drops from 3.12 to 2.57 under maximin opti-
219 mization. The gap reflects Sonnet’s harder time implementing complex coordination mechanisms
220 (role rotation, zone assignment) from strategic hints alone.

221 **Finding 3: Game structure determines whether fairness requires explicit optimization.** In
222 Cleanup, where cleaning costs are borne asymmetrically (cleaners pay −1, free-riders collect ap-
223 ples), baseline equality ranges from $E=0.04$ to 0.62, and maximin optimization is required to reach

Table 2: Results. Mean / max across the runs per condition. **Bold** marks the best mean per metric within each Policy LLM $\times$ Game sub-block. Baseline: unmodified, hand-designed pipeline from [3]. GEPA: automated prompt optimization [9] with matched compute budget.

Policy LLM	Target	U	E	$\min_i R_i$
<i>Cleanup — Baseline</i>				
Gemini 3.1 Pro	—	1.93/2.70	0.17/0.62	−159/−84
Sonnet 4.6	—	0.86/1.56	−0.02/0.25	−151/−59
<i>Cleanup — GEPA [9]</i>				
Gemini 3.1 Pro	Φ_U	1.34/1.37	−0.10/−0.05	−149/−126
	$\Phi_{\min}$	1.76/2.76	0.90/0.96	143/245
Sonnet 4.6	Φ_U	1.04/1.20	−0.59/−0.21	−164/−126
	$\Phi_{\min}$	−0.63/1.60	0.77/1.00	−147/6
<i>Cleanup — Two-level Autoresearch</i>				
Gemini 3.1 Pro	Φ_U	3.20 /3.25	0.55/0.61	−211/−182
	$\Phi_{\min}$	3.16/3.19	0.98 /0.98	290 /296
Sonnet 4.6	Φ_U	3.12 /3.14	0.66/0.70	−196/−146
	$\Phi_{\min}$	2.57/2.93	0.91 /0.97	179 /200
<i>Gathering — Baseline</i>				
Gemini 3.1 Pro	—	2.04/2.42	0.90/0.98	412/571
Sonnet 4.6	—	0.03/0.03	0.54/0.54	0/0
<i>Gathering — GEPA [9]</i>				
Gemini 3.1 Pro	Φ_U	2.08/2.35	0.94/0.96	436/518
Sonnet 4.6	Φ_U	1.20/1.23	0.63/0.71	86/123
<i>Gathering — Two-level Autoresearch</i>				
Gemini 3.1 Pro	Φ_U	2.49 /2.51	0.98 /0.98	582 /582
Sonnet 4.6	Φ_U	2.52 /2.52	0.96 /0.98	576 /597

224 $E > 0.9$. In Gathering, where all agents face a symmetric landscape, efficiency optimization alone
225 achieves $E > 0.94$ across all 4 runs: no separate maximin runs are needed. This generalizes: in
226 *provision* dilemmas with asymmetric costs, fairness requires designed mechanisms (role rotation,
227 duty sharing); in *restraint* dilemmas with symmetric costs, fairness emerges as a free byproduct of
228 efficient coordination. The researcher independently discovers this, it creates role differentiation
229 pipelines only for Cleanup, and pure spatial-coordination pipelines for Gathering.

230 **Finding 4: Convergent discovery of qualitatively different strategies per objective.** Despite
231 fully independent runs, the researcher converges on the same core strategies within each condition
232 (Table 3, Appendix B). In Cleanup, waste-counting helpers and spatial zone partitioning appear
233 across all runs. The qualitative dividing line is the presence of an explicit *fairness mechanism*,
234 which appears in 4/4 maximin runs but 0/4 efficiency runs. In 3/4 maximin runs (both Gemini runs
235 and one Sonnet run, Listings 7, 8) the researcher writes *time-based role rotation* into the synthe-
236 sizer prompt; in the remaining maximin run the researcher writes a structurally distinct “collective
237 threshold” mechanism in which all agents synchronously switch between cleaning and collecting
238 based on `waste_fraction(env)` (achieving comparable maximin without an agent-index phase).
239 Under efficiency optimization, the researcher instead writes *static* role assignment (some agents al-
240 ways clean), producing high collective output at the cost of equality (Listing 6). In Gathering, the
241 researcher discovers BFS-Voronoi territory partitioning and respawn-timer awareness (Listings 3, 9),
242 with no role differentiation: optimization is purely spatial (who collects which apples) and temporal
243 (respawn-aware positioning).

244 The convergence here is at the level of *which artifacts* $\mathcal{R}$ injects into $p, \phi, \mathcal{H}$. Once the researcher-
245 authored prompt contains a rotation template (e.g., Listing 4), the downstream synthesizer’s imple-
246 mentation is unsurprising. The non-trivial claim is that $\mathcal{R}$ writes such a template only under $\Phi_{\min}$,
247 never under Φ_U , despite the researcher system prompt $p_{\mathcal{R}}$ being identical across objectives and con-
248 taining no rotation formula, no objective-conditional guidance, and no link between any strategy
249 class and either welfare criterion (Listing 1, Appendix C.1).




Figure 2: Efficiency (U) across researcher iterations for all 12 runs. *Left*: Cleanup with 8 runs across 2 LLMs $\times$ 2 objectives (solid lines = kept experiments, open circles = discarded). Dashed horizontal line: hand-designed baseline from [3]. All runs converge to $U \approx 3.1$ – 3.2 despite diverse starting points. *Right*: Gathering with 4 runs (2 per LLM). All converge to $U \approx 2.5$ within 2–4 iterations despite baselines ranging from 0.03 to 2.42.

250 **Common failure modes.** Several patterns led to discarded autoresearch iterations across all runs:
 251 (1) *over-prescription*: adding too many strategic hints confuses the policy LLM, causing generation
 252 failures; (2) *iteration regression*: $K=4$ or $K=5$ inner iterations often cause catastrophic regression
 253 as the LLM “over-refines” a working policy; (3) *feedback overload*: verbose per-agent diagnostics
 254 cause over-correction rather than gradual improvement.

255 5 Related Work

256 **Sequential social dilemmas.** SSDs were introduced by Leibo et al. [1] (Gathering) and extended
 257 to public-goods settings by Hughes et al. [5] (Cleanup); Perolat et al. [6] formalized the social
 258 outcome metrics (U, E, S, P) used here. Subsequent work studies inequity aversion, intrinsic moti-
 259 vation, and reputational mechanisms for promoting cooperation in MARL agents. We complement
 260 this line by automating the search for cooperative *programs* rather than evolving neural policies, and
 261 by treating the welfare objective Φ as a designable parameter that the outer agent optimizes for, ob-
 262 taining qualitatively different cooperative behaviors (static role assignment vs. duty rotation) under
 263 different objectives without changing the environment.

264 **LLMs for policy and program synthesis.** FunSearch [10] evolves programs for combinato-
 265 rial discovery; Eureka [11] synthesizes reward functions from environment source code; Voy-
 266 ager [12] and Code as Policies [13] generate executable skill code for single-agent embodied control;
 267 ReEvo [14] evolves heuristics with reflective feedback. These works target single-agent settings or
 268 non-strategic optimization. Gallego [3], on which our inner loop is based, extended LLM program
 269 synthesis to the multi-agent SSD setting, where one program must coordinate N self-play copies.
 270 Our contribution is one level above: rather than tuning the synthesizer for one task, we let an outer
 271 agent rewrite the synthesis pipeline itself.

272 **LLM reflection and prompt optimization.** Reflexion [15], Self-Refine [16], OPRO [17], and
 273 GEPA [9] demonstrate that structured verbal feedback loops improve LLM outputs; ERL [18] in-
 274 ternalizes such reflection via self-distillation. These methods optimize the prompt or the model’s
 275 reasoning trajectory in isolation. Our framework instead modifies the entire surrounding pipeline
 276 (system prompt, feedback construction, helper library, iteration logic) making prompt optimization
 277 one component of a strictly larger search space. The empirical gap is large: Section 4 shows full-
 278 pipeline autoresearch achieving $3\times$ higher efficiency in Cleanup and 37% higher in Gathering than
 279 GEPA at matched compute, with run-to-run spread an order of magnitude tighter on the harder
 280 fairness objective.

281 **Automated AI research.** A small but growing body of work delegates the design of ML pipelines
 282 to LLM coding agents. Karpathy’s autoresearch [4] runs a coding agent that modifies `train.py`
 283 for nanoGPT pretraining and is rewarded for validation loss. Our system shares this autoresearch
 284 architecture (coding agent + frozen evaluation harness + diff-based history) but targets a different in-
 285 ner loop (multi-agent policy synthesis instead of single-model pretraining) and a different objective
 286 (multi-agent social welfare instead of validation bits-per-byte). To our knowledge this is the first
 287 instantiation of the autoresearch paradigm in a multi-agent decision-making domain.

288 **Automated mechanism and information design.** Classical automated mechanism design [7]
 289 computes optimal allocation rules for self-interested agents under explicit incentive constraints,
 290 while information design [8] studies how a principal commits to a signaling policy that shapes a
 291 receiver’s behavior. Our outer agent solves a related but distinct problem: $\mathcal{M}$ is not strategically
 292 deceptive (it follows instructions) but is *boundedly rational*, so the researcher must decide what
 293 information, helpers, and structure to expose so that $\mathcal{M}$ writes policies achieving the principal’s
 294 welfare goal. Section 4 shows that the pipelines $\mathcal{R}$ produces are genuinely a function of the welfare
 295 objective, supporting this information-design framing empirically.

296 6 Discussion and Conclusion

297 We have presented a two-level framework where a coding agent autonomously discovers pipeline
 298 configurations that improve the output of an inner-loop LLM system. Applied to multi-agent policy
 299 synthesis in social dilemmas, the researcher agent reliably exceeds hand-designed baselines across
 300 multiple independent runs, and converges on qualitatively similar strategies within each game-
 301 objective condition. The framework requires no task-specific scaffolding beyond a standard CLI
 302 and git: the agent operates on a standard software repository using the same tools (file editing,
 303 shell commands, git) available to a human researcher. The inner-loop validation pipeline, helper-
 304 library skeleton, and orchestrator API are themselves deliberately scoped artifacts, however (see
 305 Appendix A).

306 **Mechanism design in action.** Our results empirically support the mechanism design interpreta-
 307 tion of Section 3.4. The researcher acts as an information designer: under Φ_U , it reveals efficiency-
 308 oriented information (waste counts, zone assignments) that guides $\mathcal{M}$ toward productive but un-
 309 equal role allocation (Listing 6). Under $\Phi_{\min}$, it additionally reveals fairness-oriented structure such
 310 as rotation schedules and per-agent equity feedback (Listings 4, 5), guiding $\mathcal{M}$ toward egalitarian
 311 coordination.

312 **Two types of dilemma.** The Cleanup–Gathering contrast reveals a structural insight about social
 313 dilemmas: *provision* dilemmas (asymmetric costs, one agent’s contribution benefits all) require ex-
 314 plicit fairness mechanisms, while *restraint* dilemmas (symmetric costs, each agent faces the same
 315 temptation) achieve fairness naturally when coordination improves. This distinction, well-known
 316 in the common-pool resource literature, is rediscovered here by an autonomous AI agent, which
 317 suggests it is a robust structural feature of these environments.

318 **Human oversight and audit.** Our system is fully autonomous by design, but its architecture of-
 319 fers natural affordances for human-in-the-loop oversight. Because the researcher $\mathcal{R}$ operates on a
 320 standard git repository, a domain expert can audit the full history of code diffs Δ_j alongside their
 321 metric outcomes $(\Phi_j, \mathbf{m}_j)$. The keep/discard decision (Algorithm 1, line 11) is a lightweight in-
 322 tervention point: a human could override the accept threshold, veto modifications that introduce
 323 undesirable mechanisms (e.g., exploitative role assignments), or steer the researcher toward under-
 324 explored regions of the configuration space by editing the experiment history. More broadly, the
 325 separation between the researcher $\mathcal{R}$ and the frozen evaluation Φ creates a *delegation boundary*: the
 326 human defines *what* to optimize (the welfare objective), while the agent decides *how*. This mirrors
 327 the expert-in-the-loop design patterns identified in deployment-oriented work on AI agents for dis-
 328 covery, where trust calibration depends on the legibility of the agent’s decision trail rather than on
 329 real-time supervision.

330 **Future work.** Several directions emerge. First, we are interested in applying the two-level frame-
 331 work to other LLM-driven pipelines (code optimization, scientific experiment design, infrastructure

tuning), where the same architecture applies with a different inner loop and objective Φ ; next, adversarial objectives can be intriguing, to test whether the researcher discovers exploitative pipeline configurations, exposing reward hacking risks; and asymmetric programs: extend to settings where agents run *different* source code (the present setup is symmetric in code but already supports heterogeneous behavior via `agent_id`; relaxing the shared-program constraint enables richer per-agent specialization and partial-information settings).

References

- [1] Joel Z Leibo, Vinicius Zambaldi, Marc Lanctot, Janusz Marecki, and Thore Graepel. Multi-agent reinforcement learning in sequential social dilemmas. In *Proc. 16th Conference on Autonomous Agents and MultiAgent Systems*, pages 464–473, 2017.
- [2] Lucian Buşoniu, Robert Babuška, and Bart De Schutter. A comprehensive survey of multiagent reinforcement learning. *IEEE Transactions on Systems, Man, and Cybernetics, Part C (Applications and Reviews)*, 38(2):156–172, 2008.
- [3] Víctor Gallego. Cooperation and exploitation in llm policy synthesis for sequential social dilemmas. *arXiv preprint arXiv:2603.19453*, 2026.
- [4] Andrej Karpathy. autoresearch: AI agents running research on single-GPU nanochat training automatically. <https://github.com/karpathy/autoresearch>, 2026.
- [5] Edward Hughes, Joel Z Leibo, Matthew Phillips, Karl Tuyls, Edgar A Dueñez-Guzmán, Antonio García Castañeda, Iain Dunning, Tina Zhu, Kevin R McKee, R. Koster, Heather Roff, and Thore Graepel. Inequity aversion improves cooperation in intertemporal social dilemmas. In *Neural Information Processing Systems*, 2018.
- [6] Julien Perolat, Joel Z Leibo, Vinicius Zambaldi, Charles Beattie, Karl Tuyls, and Thore Graepel. A multi-agent reinforcement learning model of common-pool resource appropriation. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- [7] Vincent Conitzer and Tuomas Sandholm. Complexity of mechanism design. In *Proc. 18th Conference on Uncertainty in Artificial Intelligence*, pages 103–110, 2002.
- [8] Emir Kamenica and Matthew Gentzkow. Bayesian persuasion. *American Economic Review*, 101(6):2590–2615, 2011.
- [9] Lakshya A Agrawal et al. GEPA: Reflective prompt evolution can outperform reinforcement learning. *arXiv preprint arXiv:2507.19457*, 2026.
- [10] Bernardino Romera-Paredes, Mohammadamin Barekatain, Alexander Novikov, Matej Balog, M Pawan Kumar, Emilien Dupont, Francisco JR Ruiz, Jordan S Ellenberg, Pengming Wang, Omar Fawzi, et al. Mathematical discoveries from program search with large language models. *Nature*, 625(7995):468–475, 2024.
- [11] Yecheng Jason Ma, William Liang, Guanzhi Wang, De-An Huang, Osbert Bastani, Dinesh Jayaraman, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Eureka: Human-level reward design via coding large language models. In *The Twelfth International Conference on Learning Representations*, 2024.
- [12] Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *Transactions on Machine Learning Research*, 2024.
- [13] Jacky Liang, Wenlong Huang, Fei Xia, Peng Xu, Karol Hausman, Brian Ichter, Pete Florence, and Andy Zeng. Code as policies: Language model programs for embodied control. In *2023 IEEE International Conference on Robotics and Automation (ICRA)*, pages 9493–9500, 2023.
- [14] Haoran Ye, Jiarui Wang, Zhiguang Cao, Federico Berto, Chuanbo Hua, Haeyeon Kim, Jinkyoo Park, and Guojie Song. ReEvo: Large language models as hyper-heuristics with reflective evolution. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.

- [15] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: language agents with verbal reinforcement learning. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.
- [16] Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. Self-refine: Iterative refinement with self-feedback. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.
- [17] Chengrun Yang, Xuezhi Wang, Yifeng Lu, Hanxiao Liu, Quoc V Le, Denny Zhou, and Xinyun Chen. Large language models as optimizers. In *The Twelfth International Conference on Learning Representations*, 2024.
- [18] Taiwei Shi, Sihao Chen, Bowen Jiang, Linxin Song, Longqi Yang, and Jieyu Zhao. Experiential reinforcement learning. *arXiv preprint arXiv:2602.13949*, 2026.

A Limitations

We flag three limitations of the present study, in roughly decreasing order of how much they bound the claims.

1. *Single researcher LLM*. All main-experiment runs (and the Gemma appendix) use the same researcher $\mathcal{R}$ (Claude Opus 4.6 via the Claude Code CLI). A researcher ablation across $\mathcal{R}$ -LLMs, in which the same fixed system prompt $p_{\mathcal{R}}$ (Appendix C.1) is run with several frontier coding agents, is the single most important follow-up.
2. *Inner-loop infrastructure is itself a designed artifact*. The line “ $\mathcal{R}$ uses no task-specific scaffolding” refers to the researcher’s tooling (CLI, git, file edits) and to the fact that the configuration $c=(p, \phi, \mathcal{H}, \iota)$ starts from a deliberately weak baseline. It does *not* mean the surrounding system is unscoped: the inner-loop validation pipeline (AST safety checks, sandboxed execution, multi-seed evaluation harness), the helper-library skeleton (`pipeline/helpers.py` starts non-empty), and the orchestrator API (`run_inner_loop.py`) are themselves engineered artifacts that bound what $\mathcal{R}$ can and cannot do. Generalization of the framework to settings without an equivalent harness is plausible but unevaluated.
3. *Gridworld scope*. The inner-loop SSDs are 2D gridworlds with fully-observed integer states, discrete action spaces, and short ($H=1000$ -step) episodes. The specific mechanisms the researcher discovers (BFS-Voronoi partitioning, time-rotated agent-id phase counters, waste-fraction threshold rules) exploit this structure and are gridworld-shaped; their transfer to higher-dimensional, continuous, or partially-observable multi-agent settings is an open question. The outer-loop *operating conditions* (noisy multi-seed evaluation, code-level edits to a real repository, bounded Φ queries) are realistic, but the conclusions about which strategies emerge are bounded by the inner-loop benchmark.

B Additional Results

Table 3: Strategies discovered by the researcher in Cleanup across 4 efficiency-optimized and 4 maximin-optimized independent runs. The grouped row records any explicit *fairness mechanism*; its two sub-rows decompose this into the dominant variant (time-based role rotation) and the alternative (synchronized whole-population clean/collect switching) found in one Sonnet maximin run.

Strategy	Φ_U (4 runs)	$\Phi_{\min}$ (4 runs)
Waste-counting helpers	4/4	4/4
Zone/lane partitioning	3/4	4/4
Anti-regression feedback	3/4	2/4
Worked policy examples	2/4	3/4
Cleaning cost economics	2/4	3/4
Explicit fairness mechanism	0/4	4/4
<i>of which:</i> time-based role rotation	0/4	3/4
<i>of which:</i> synchronized clean/collect	0/4	1/4




Figure 3: Maximin ($\min_i R_i$) across researcher iterations for the 4 Cleanup maximin-optimized runs. All runs transform deeply negative baselines (worst-off agents losing reward) into substantially positive values. Gemini reaches ~ 290 while Sonnet reaches ~ 160 – 200 . The dashed line marks $\min_i R_i = 0$ (no agent loses reward overall).




Figure 4: Final metrics across all conditions. (a) Efficiency: all conditions converge to $U \approx 2.5$ – 3.2 . (b) Equality: maximin-optimized Cleanup runs achieve $E \approx 1.0$, while efficiency-optimized runs show $E \approx 0.5$ – 0.7 ; Gathering achieves high equality regardless. (c) Maximin: the sharpest contrast—efficiency optimization leaves the worst-off agent deeply negative ($\min_i R_i \approx -200$), while maximin optimization transforms it to $+290$ (Gemini) or $+179$ (Sonnet). Gathering achieves high maximin (~ 580) under efficiency optimization alone. Error bars show s.d. across runs.

416 **Inner-loop averages confirm a broad-based improvement.** Figures 2 and 3 report the metric of
 417 the *kept* inner-loop output (π_c^*) at each outer iteration. A natural concern is that this could overstate
 418 pipeline quality if the researcher is benefitting from variance: with K inner iterations per outer
 419 step, a single lucky generation could carry the curve while the rest of the inner trajectory is noise.
 420 Figures 5 and 6 replot the same runs with the metric averaged across *all* inner iterations of each outer
 421 step. The trajectories ramp at essentially the same rate, with only mild attenuation: Cleanup mean
 422 efficiency still climbs to $\bar{U} \approx 2.5$ – 3.0 (vs. 3.1 – 3.2 for the kept policy), Gathering still saturates at
 423 $\bar{U} \approx 2.5$, and Cleanup mean maximin still reaches $\min_i R_i \approx 200$ – 275 (vs. 179 – 290). The whole
 424 inner-loop output distribution improves, not just its tail, consistent with the interpretation that the
 425 researcher is shaping the synthesizer’s behavior rather than amplifying lucky draws.




Figure 5: Mean efficiency $\bar{U}$ averaged across all inner iterations of each outer step (compare Figure 2, which shows the kept policy only). *Left*: Cleanup ($N=10$); *right*: Gathering ($N=4$). The trajectories track Figure 2 closely, indicating that the researcher’s gains come from improving the entire inner-loop output distribution, not just the best of K samples.




Figure 6: Mean maximin $\overline{\min_i R_i}$ averaged across all inner iterations of each outer step for the 4 Cleanup $\Phi_{\min}$ runs (compare Figure 3). The deeply negative baselines lift to ~ 200 – 275 mean maximin, only mildly below the kept-policy values of Figure 3. The horizontal dashed line marks $\min_i R_i = 0$.

426 B.1 Smaller Open-Weight Model: Gemma 4 26B

427 We test the framework with Gemma 4 26B-A4B-IT (Google), a 26B-parameter open-weight model
 428 substantially smaller than the frontier models in the main experiments. We run three experiment runs:
 429 two on Cleanup optimizing efficiency and maximin respectively, and one on Gathering optimizing
 430 efficiency. As with the other synthesizer models, all with Opus 4.6 as the researcher $\mathcal{R}$.

431 **Setup.** In all three experiments, Gemma starts from complete failure: the baseline pipeline pro-
 432 duces $U = -10.0$ on Cleanup (agents CLEAN-spam without collecting apples) and $U = 0.0$ on Gath-
 433 ering (broken BFS calls), compared to $U = 1.93/0.86$ (Gemini/Sonnet on Cleanup) and $U = 2.04/0.03$
 434 (Gemini/Sonnet on Gathering). The researcher ran $J = 10$ – 18 outer iterations per condition.

435 **Results.** Table 4 presents the best configurations discovered. Under efficiency optimization, the re-
 436 searcher achieves $\bar{U} = 0.87$, a dramatic recovery from the broken baseline but far below frontier mod-

437 els ($U \approx 3.2$). The researcher compensates for the model’s limited capability by reducing inner-loop
 438 iterations to $K=1$ (avoiding the regression that plagued $K \geq 2$ runs) and providing highly structured
 439 worked examples with explicit cleaning-role assignment.

Table 4: Autoresearch with Gemma 4 26B-A4B-IT. Single run per condition. Baseline: pipeline from [3].

Game	Target	U	E	$\min_i R_i$	Keep rate
Cleanup	Baseline	-10.0	1.00 [†]	-1000	—
	Φ_U	0.87	-1.11	-371	6/19
	$\Phi_{\min}$	1.71	0.94	137	9/17
Gathering	Baseline	0.0	1.00 [†]	0	—
	Φ_U	2.44	0.98	580	4/11

[†]Equality is trivially 1.0 because all agents receive near-zero reward.

440 **Maximin optimization rescues efficiency.** Strikingly, under maximin optimization the researcher
 441 achieves *higher* efficiency ($U=1.71$) than under direct efficiency optimization ($U=0.87$), along-
 442 side strong equality ($E=0.94$) and positive worst-off welfare ($\min_i R_i=137$). This reversal, absent
 443 in frontier models where both objectives yield $U \approx 3.1$ – 3.2 , occurs because the maximin objective
 444 forces the researcher toward coordination mechanisms (rotating cleaning duties, index-based apple
 445 assignment) that simultaneously improve collective output. Under efficiency-only optimization, the
 446 researcher converges on a local optimum that a 26B model can implement but that caps performance
 447 well below the frontier.

448 This suggests that *for models below a capability threshold, fairness objectives may serve as better*
 449 *optimization targets for overall social welfare than direct efficiency maximization.* The structured
 450 coordination enforced by the maximin objective provides a scaffold that compensates for the weaker
 451 model’s difficulty implementing complex strategies from strategic hints alone.

452 **Gathering: full recovery on a simpler game.** On Gathering ($N=4$), the researcher brings
 453 Gemma from $U=0.0$ to $U=2.44$ with $E=0.98$ and $\min_i R_i=580$, after 10 outer iterations (4 kept).
 454 These values nearly match frontier models (Gemini: $U=2.49$, Sonnet: $U=2.52$; Table 2). The re-
 455 searcher discovers the same Voronoi partitioning and respawn-aware camping strategies found in the
 456 main experiments, and increases inner-loop iterations to $K=5$ to allow sufficient refinement. The
 457 contrast with Cleanup suggests that the researcher can fully compensate for model weakness on sim-
 458 pler coordination tasks (4 agents, symmetric costs) but not on harder provision dilemmas (10 agents,
 459 asymmetric costs).

460 B.2 Compute Requirements

461 Table 5 reports wall-clock time for all 12 autoresearch runs. *Inner loop* measures total time spent
 462 on policy LLM generation and environment simulation across all evaluations; *total* (estimated from
 463 run-directory timestamps) additionally includes the researcher agent’s analysis, code editing, and
 464 decision-making between evaluations.

465 **Cost breakdown.** Each inner loop evaluation involves $K=2$ – 3 policy LLM generation calls (each
 466 ~ 2 k input tokens, ~ 1.5 k output tokens) followed by multi-seed simulation (5–12 seeds, ~ 15 – 30 s
 467 total). Policy LLM generation dominates inner loop time (86–97%), with Sonnet evaluations taking
 468 $\sim 2\times$ longer than Gemini due to extended thinking. The researcher agent (Claude Opus 4.6, running
 469 via Claude Code CLI) accounts for 41% of total wall-clock time (25.6h of the 62.2h total across
 470 all 12 runs), spent reading results, editing pipeline source files, and planning modifications. In
 471 monetary terms, the researcher agent is the dominant cost: each outer iteration consumes ~ 50 – 100 k
 472 context tokens in the Opus session, whereas each inner loop evaluation uses only ~ 5 – 10 k policy
 473 LLM tokens.

Table 5: Wall-clock time per autonomous run. Evals: total inner loop evaluations including initial baseline. Per eval: mean wall-clock per single evaluation. Total: end-to-end including researcher overhead.

Policy LLM	Φ	Evals	Inner loop (h)	Per eval (min)	Total (h)
<i>Cleanup</i> ($N=10$)					
Gemini	Φ_U	11	2.6	14	3.7
Gemini	Φ_U	18	4.1	14	6.4
Sonnet	Φ_U	4	2.1	32	6.1
Sonnet	Φ_U	7	5.0	43	6.1
Gemini	$\Phi_{\min}$	17	3.0	11	4.3
Gemini	$\Phi_{\min}$	11	3.6	20	5.0
Sonnet	$\Phi_{\min}$	8	4.5	33	8.8
Sonnet	$\Phi_{\min}$	9	5.3	36	13.2
<i>Gathering</i> ($N=4$)					
Sonnet	Φ_U	3	1.7	33	2.2
Gemini	Φ_U	5	1.4	17	1.9
Gemini	Φ_U	3	0.9	19	1.1
Sonnet	Φ_U	4	2.3	35	3.4
All 12 runs		100	36.6	22	62.2

C Researcher-Authored Pipeline Artifacts

The previous appendix shows the policies π_c^* that the inner loop *outputs*; this one shows the configuration $c = (p, \phi, \mathcal{H}, \iota)$ (Section 4, Table 1) that the researcher $\mathcal{R}$ *authored* to produce them. Excerpts are taken from the final commit on the dedicated git branch of the corresponding run; we reproduce the prose verbatim, with author commentary in [brackets] and elisions marked We organize by artifact type so each subsection makes a within-type contrast (e.g., Φ_U vs. $\Phi_{\min}$, or Cleanup vs. Gathering). Subsection C.1 comes first because it shows the upstream input that the rest of the appendix is derived from: the prompt $p_{\mathcal{R}}$ given to the researcher itself.

C.1 Researcher system prompt $p_{\mathcal{R}}$

The researcher $\mathcal{R}$ is a coding agent (Claude Opus 4.6 via the Claude Code CLI) instantiated with a single fixed system prompt that is identical across all experiment runs (it differs only in the game name and the primary-metric flag, which is one of `efficiency` or `maximin`). This is the input from which every artifact in the rest of this appendix is derived. Because much of the paper’s claim hinges on the asymmetry between *what $p_{\mathcal{R}}$ tells $\mathcal{R}$* and *what $\mathcal{R}$ in turn writes for $\mathcal{M}$* , we reproduce the strategically relevant portions of $p_{\mathcal{R}}$ verbatim below; the full file is in the released repository under `autoresearch/program.md`.

Listing 1: Excerpt from $p_{\mathcal{R}}$. The remaining parts of $p_{\mathcal{R}}$ describe the file layout of the inner-loop pipeline, the evaluation script invocation, the keep/discard rule, and the logging format; none of those add task-specific strategic guidance.

```

490  ## The metrics
491
492  The primary metric is specified at launch via '--metric' (default: 'efficiency').
493  The two options are:
494
495  Efficiency (U): collective apple collection rate across all agents per
496  timestep. Higher is better. ...
497
498  Maximin (Rawlsian welfare): minimum total per-agent return across all
499  agents. Higher is better. Inspired by Rawls' difference principle -- a just
500  policy maximizes the welfare of the worst-off agent. ...
501
502  ## Strategy space
503
504  Here are categories of modifications to explore:
505
506  ### Prompt engineering (p)
507  - Add strategic hints about the Cleanup dilemma (e.g., "cleaning is a public
508  good -- someone must do it")
509  - Add worked examples of sophisticated policies
510

```

```

511 - Restructure the API documentation for clarity
512 - Add game-theoretic reasoning frameworks
513 - Mention optimal strategies from the literature (Voronoi partitioning,
514   role assignment)
515
516 ### Feedback engineering (l, phi)
517 - Show per-agent reward breakdown (not just average)
518 - Add derived metrics (e.g., cleaning rate, waste level trends, apple growth)
519 - Frame feedback to emphasize cooperation
520 - Add temporal analysis ...
521 - Show metric trends across iterations ...
522 - Provide diagnostic hints based on metrics ...
523
524 ### Helper functions (H)
525 - count_waste(env), waste_fraction(env), bfs_to_waste(env, agent_id),
526   should_clean(env), assign_role(env, agent_id),
527   find_cleaning_position(env, agent_id)
528
529 ### Iteration logic (iota)
530 - Change K (more iterations = more refinement but more cost)
531 - Change eval seeds, retry budget, thinking budget

```

533 **What $p_{\mathcal{R}}$ does *not* say.** For the convergent-discovery claim of Finding 4 to be informative, $p_{\mathcal{R}}$
534 must not itself encode the specific mechanisms that $\mathcal{R}$ later writes into $p, \phi, \mathcal{H}$ under $\Phi_{\min}$. Inspect-
535 ing Listing 1, three absences matter:

- 536 • *No rotation formula.* $p_{\mathcal{R}}$ never mentions `(agent_id + env._step_count // T) % n`,
537 `env._step_count`, or any other time-based cycling pattern. The closest item – “Mention op-
538 timal strategies from the literature (Voronoi partitioning, role assignment)” – names *static* role
539 assignment, not time-rotated duty.
- 540 • *No objective-conditional guidance.* The strategy-space listing is identical regardless of whether
541 $\mathcal{R}$ is launched with `-metric efficiency` or `-metric maximin`. $p_{\mathcal{R}}$ does not tell $\mathcal{R}$ to behave
542 differently under the two objectives; the only objective-sensitive input is the (scalar) score Φ_j
543 returned after each inner-loop run.
- 544 • *No “rotation is for fairness.”* $p_{\mathcal{R}}$ does not link any strategy to either welfare criterion. The
545 connection “rotation = fairness = high maximin” is constructed by $\mathcal{R}$ from Φ_j observations across
546 iterations on the branch.

547 The researcher-authored prompt p shown in Listing 4 and its Sonnet counterpart, both written under
548 $\Phi_{\min}$, contain all three of these elements; the corresponding Φ_U prompts do not. The role of $p_{\mathcal{R}}$ in
549 the experiment is therefore to define the configuration space and the evaluation harness, not to script
550 the discoveries themselves.

551 C.2 Helper library $\mathcal{H}$: coordination primitives

552 The researcher adds primitives that the policy LLM can call as black boxes, sparing it from re-
553 implementing tricky logic on every iteration. Two patterns dominate: (i) *state inspection* (waste
554 counts, fractions, beam-yield scoring) and (ii) *coordination primitives* (zone assignment, role rota-
555 tion). We show one of each.

Listing 2: Cleanup, $\Phi_{\min}$ – band-based apple zoning helper added by the researcher in exp5. This is
the primitive that lets the policy in Listing 7 keep collectors within their own row band, preventing
all 7 gatherers from racing to the same apple.

```

556 1 def get_my_apples(env, agent_id):
557 2     """Alive apples in this agent's horizontal row band.
558 3     Divides the orchard into n_agents bands; falls back to all apples if empty."""
559 4     aid = int(agent_id)
560 5     n = int(env.n_agents)
561 6     band_h = env.height / n
562 7     band_lo, band_hi = aid * band_h, (aid + 1) * band_h
563 8
564 9     my_apples, all_apples = set(), set()
565 10    for i in range(env.n_apples):
566 11        if env.apple_alive[i]:
567 12            ar = int(env._apple_pos[i, 0]); ac = int(env._apple_pos[i, 1])
568 13            all_apples.add((ar, ac))
569 14            if band_lo <= ar < band_hi:
570 15                my_apples.add((ar, ac))
571 16    return my_apples if my_apples else all_apples
572

```


Listing 3: Gathering – BFS-Voronoi territory + respawn-aware waiting helpers added by the researcher in gather-exp3. These are the primitives that Listing 9 composes into a one-line policy: “go to my_zone_apples, otherwise nearest_respawning_apple.”

```

574 1 def voronoi_zones(env):
575 2     """Multi-source BFS Voronoi over walkable cells.
576 3     Returns (row, col) -> agent_id; ties broken by lower agent_id."""
577 4     queue, visited = deque(), {}
578 5     for a_id in range(env.n_agents):
579 6         if int(env.agent_timeout[a_id]) == 0:
580 7             ar = int(env.agent_pos[a_id][0]); ac = int(env.agent_pos[a_id][1])
581 8             queue.append((ar, ac, a_id))
582 9             visited[(ar, ac)] = a_id
583 10    while queue:
584 11        r, c, a_id = queue.popleft()
585 12        for dr, dc in [(-1,0),(1,0),(0,-1),(0,1)]:
586 13            nr, nc = r + dr, c + dc
587 14            if 0 <= nr < env.height and 0 <= nc < env.width \
588 15                and not env.walls[nr, nc] and (nr, nc) not in visited:
589 16                visited[(nr, nc)] = a_id
590 17                queue.append((nr, nc, a_id))
591 18    return visited
592 19
593 20 def nearest_respawning_apple(env, agent_id, zones=None, max_timer=10):
594 21     """Nearest dead apple in MY zone respawning within max_timer steps.
595 22     Returns (row, col) of best wait spot, or None."""
596 23     if zones is None: zones = voronoi_zones(env)
597 24     best_pos, best_t, best_d = None, max_timer + 1, float('inf')
598 25     ar = int(env.agent_pos[agent_id][0]); ac = int(env.agent_pos[agent_id][1])
599 26     for i in range(env.n_apples):
600 27         if not env.apple_alive[i]:
601 28             pos = (int(env._apple_pos[i][0]), int(env._apple_pos[i][1]))
602 29             if zones.get(pos) == agent_id:
603 30                 t = int(env.apple_timer[i])
604 31                 if t <= max_timer:
605 32                     d = abs(pos[0] - ar) + abs(pos[1] - ac)
606 33                     if t < best_t or (t == best_t and d < best_d):
607 34                         best_pos, best_t, best_d = pos, t, d
608 35     return best_pos
609

```

The choice of helper depends on the dilemma’s geometry: row-bands suffice when fairness is the main concern (Cleanup maximin) but full BFS-Voronoi is needed when wall-aware territory ownership matters (Gathering). The researcher does not fix this in advance, it picks the simpler primitive whenever it works.

C.3 System prompt p : from neutral framing to strategic briefing

The unmodified system prompt is a neutral API description: “Write a policy that maximizes per-agent reward.” Under maximin optimization, the researcher rewrites the prompt into a multi-section strategic briefing. The excerpt below shows the pieces it added to the Cleanup prompt in exp5; the full prompt grew from 165 to 325 lines.

Listing 4: Cleanup, $\Phi_{\min}$ – excerpts from the researcher-rewritten system prompt (exp5). The unmodified baseline contained only API documentation; everything below was added by $\mathcal{R}$.

```

620 ## CRITICAL OBJECTIVE: Rawlsian Fairness (Maximin)
621
622 Your goal is to maximize the minimum per-agent total return across all
623 agents. This is the "maximin" or Rawlsian welfare criterion.
624
625 This means: it is NOT enough for the *average* agent to do well. The
626 worst-off agent must do as well as possible. A policy where 8 agents
627 each earn +200 but 2 agents each earn -100 is TERRIBLE (maximin = -100).
628
629 Key implication: cleaning costs (-1 per CLEAN action) must be shared
630 equitably among ALL agents. If you assign fixed "cleaner" roles, those
631 agents will accumulate large negative rewards and destroy your maximin score.
632
633 ## Strategy for Maximin: ROLE ROTATION + APPLE ZONING
634
635 The optimal strategy for maximin involves:
636 1. Shared cleaning duty: ALL agents take turns cleaning using a rotation
637    schedule based on 'agent_id' and 'env_step_count'. For example:
638    'is_my_cleaning_turn = (agent_id + env_step_count // SHIFT) % env.n_agents < NUM_CLEANERS'
639    where SHIFT is ~50 steps and NUM_CLEANERS is 2-3.
640 2. When NOT your turn: collect apples from YOUR ZONE using
641

```

```

642     'get_my_apples(env, agent_id)' -- prevents all gatherers competing
643     for the same nearest apple.
644 3. NEVER use BEAM (action 6): it costs -1 and causes -50 to the target.
645 4. Keep waste density low: aim for ~2-3 active cleaners at any time.
646
647 [... API documentation, helper documentation, working example ...]
648
649 IMPORTANT:
650 - NEVER assign permanent cleaning roles to specific agents -- this kills maximin.
651 - Use env._step_count for time-based role rotation so all agents share cleaning duty.
652 - NEVER use BEAM (action 6) -- it destroys both agents' rewards.

```

Two things stand out. First, the researcher writes the formula it expects $\mathcal{M}$ to use almost verbatim (`(agent_id + env._step_count // SHIFT) % env.n_agents < NUM_CLEANERS`) and the policy in Listing 7 uses essentially this template. Second, the researcher *teaches by counter-example*: the explicit “8 agents earn +200, 2 earn -100, maximin = -100” makes the failure mode of Φ_U -style static roles (Listing 6) concretely visible to $\mathcal{M}$. This is the information-design move predicted by the mechanism-design framing in Section 3.4.

The Gathering prompt evolves along a different axis (no maximin, no rotation). Its researcher-added “Key Strategic Insights” section instead emphasizes (i) *self-play implies never beam* and (ii) *Voronoi partition each step*, exactly matching what the policy in Listing 9 implements.

C.4 Feedback ϕ : adaptive diagnostics

The baseline feedback function from [3] shows reward + four social metrics with definitions and stops there. The researcher turns it into a state machine: it inspects the latest metrics and conditionally injects different hints. Two different runs produced two qualitatively different diagnostics: the same metric ($\bar{r}_k$ or $\min_i R_i$) is repurposed to either *stabilize* a working policy or *redirect* a failing one.

Listing 5: Adaptive feedback diagnostics. *Top*: efficiency-optimized run (exp1) – a stability guard prevents iter-3 regression once U is high. *Bottom*: maximin-optimized run (exp5) – two fairness diagnostics that fire when the worst-off agent loses reward.

```

669 1 # --- Cleanup, Phi_U feedback (exp1, pipeline/feedback.py final) ---
670 2 last_eff = history[-1]["metrics"]["efficiency"]
671 3 if last_eff >= 2.5:
672 4     parts.append(
673 5         """CRITICAL -- DO NOT REGRESS!: The current policy achieves high "
674 6         "efficiency (>=2.5). You MUST output a policy that is nearly identical "
675 7         "to the current one. Copy the current policy code and make AT MOST "
676 8         "one small targeted improvement (e.g., adjust a single numeric "
677 9         "threshold). If you are not confident a change will help, output the "
678 10        "current policy UNCHANGED. A regression here means the run fails."
679 11    )
680 12 # --- Cleanup, Phi_min feedback (exp5, pipeline/feedback.py final) ---
681 13 last_maximin = history[-1]["metrics"].get("maximin", 0)
682 14 last_avg = history[-1]["reward_avg"]
683 15 if last_maximin < 0:
684 16     parts.append(
685 17         f"""FAIRNESS ALERT!: maximin={last_maximin:.1f} is NEGATIVE. "
686 18         f"The worst-off agent lost reward overall while average was {last_avg:.1f}. "
687 19         f"This means cleaning duties are NOT shared equitably. "
688 20         f"Use time-based role rotation (env._step_count) so ALL agents share "
689 21         f"cleaning costs equally. NEVER assign permanent cleaner roles."
690 22     )
691 23 elif last_maximin < last_avg * 0.5:
692 24     parts.append(
693 25         f"""FAIRNESS WARNING!: maximin={last_maximin:.1f} is much lower than "
694 26         f"average={last_avg:.1f}. The gap suggests unequal cleaning burden. "
695 27         f"Ensure ALL agents rotate through cleaning duty."

```

The two diagnostics are written by independent researcher runs but converge on a common pattern: *thresholded interventions on a primary metric*. They are also a common cause of the low run-to-run variance reported in Finding 1: when a working policy lands in the high- U basin, the stability guard prevents the LLM from over-refining and tipping out of it (the $K=3$ regressions visible in Section 4’s “Common failure modes” (4)), and when an unfair policy lands in the negative-maximin region, the alert injects exactly the rotation hint that recovers it.

C.5 Iteration logic ι : per-condition hyperparameters

Although ι has the smallest source surface, the researcher’s edits to it explain a meaningful fraction of the variance reduction noted in Finding 1. Table 6 summarizes the values shipped in the final commit of each best-performing run.

Table 6: Iteration-logic (ι) values in the final commit of selected runs. K = inner-loop iterations, $|S|$ = evaluation seeds, “thinking” = extended-thinking token budget passed to $\mathcal{M}$. Baseline ([3]) is $K=3$, $|S|=5$, thinking 16k.

Run	Game / objective	K	$ S $	thinking	Researcher’s stated rationale
exp1 (Gemini)	Cleanup, Φ_U	3	5	16k	default; stability comes from feedback hint
exp4 (Sonnet)	Cleanup, Φ_U	3	5	32k	“Sonnet was over-refining at 16k thinking”
exp5 (Gemini)	Cleanup, $\Phi_{\min}$	2	12	16k	“ $K=2$ avoids iter-3 regression; 12 seeds for variance control”
exp7 (Sonnet)	Cleanup, $\Phi_{\min}$	3	5	10k	“reduced thinking from 16k – Sonnet was generating runaway code at 16k”
gather-exp3 (Gemini)	Gathering, Φ_U	3	5	16k	default; Gathering converges in $K \leq 3$

The maximin Gemini run is the most informative case: the researcher discovered that with $K=3$ and $|S|=5$, identical configurations could yield $\min_i R_i=295$ on one set of seeds and $\min_i R_i=100$ on another (this is the “variance exposure” failure mode of Section 4’s common failures). It then walked $|S|$ from $5 \rightarrow 8 \rightarrow 12$ over three outer iterations until the seed-averaged signal was stable enough that the policy LLM stopped chasing noise. The policy in Listing 7 is sampled at $|S|=12$.

C.6 Cross-artifact observations

- *Helpers do the algebra; prompts pick the strategy class.* The researcher uses helpers (2, 3) to put non-trivial primitives at $\mathcal{M}$ ’s fingertips, and the prompt (4) to tell $\mathcal{M}$ which primitive to compose. Prompts without supporting helpers produced policies that “mention rotation” but failed to implement it; helpers without prompt updates often went unused.
- *Feedback is the only adaptive component.* p , $\mathcal{H}$, and ι are static within a run; ϕ is the only place where the researcher gets to change what $\mathcal{M}$ sees during the inner loop, and it uses thresholded diagnostics (Listing 5) to do so. This is what mostly drives the run-to-run variance reduction.
- *The same artifact has different content under each objective.* The Cleanup prompt under Φ_U omits “rotation” and “maximin” entirely; the same prompt under $\Phi_{\min}$ devotes its entire opening section to it. The configuration c is genuinely a function of the welfare objective Φ , not a fixed pipeline parameterized by it. This is consistent with the information-design framing of Section 3.4.
- *The researcher’s edits are short and targeted.* The full diff between the baseline pipeline and the best maximin Gemini configuration totals ~ 300 added lines spread over four files; ~ 80 of them appear in the listings above. The remaining content is API documentation and worked examples, included for completeness but not strategically novel.

D Selected Generated Policies

The code excerpts below are verbatim outputs of the policy synthesizer $\mathcal{M}$, taken from the final inner-loop iteration of the best-performing run in each condition. To fit the page, we elide repeated boilerplate (BFS scaffolding, fallback branches) with $\dots$ and add commentary in [brackets]; nothing else has been edited. The listings illustrate the qualitative differences highlighted in Findings 2–4 of Section 4: *static* vs. *rotating* cleaning roles in Cleanup, and the *Voronoi* + *respawn-timer* structure in Gathering. All four policies were independently generated by the LLM under different seeds, models, and outer-loop runs.

D.1 Cleanup, Φ_U (Gemini): static interleaved roles + dynamic threshold

Best run: $U=3.25$, $E=0.61$, $\min_i R_i=-182$. The synthesizer assigns *permanent* cleaning roles via $\text{agent_id} \bmod n < \tau$, where the threshold τ depends smoothly on the current waste fraction. Higher pollution recruits more cleaners, but the *identity* of each cleaner is fixed: agents $0-\tau-1$ always pay the cost. This drives U to its ceiling but leaves equality at 0.61 and pushes the worst-off agent to $\min_i R_i=-182$.

Listing 6: Cleanup, Φ_U – Gemini 3.1 Pro. Static role assignment with waste-tiered threshold (verbatim, abridged).

```

742 1 def policy(env, agent_id) -> int:
743 2     if int(env.agent_timeout[agent_id]) > 0:
744 3         return 7 # STAND while removed
745 4
746 5     waste_frac = waste_fraction(env)
747 6     n_agents = env.n_agents
748 7
749 8     # [Smooth dynamic threshold: more pollution -> more cleaners]
750 9     if waste_frac > 0.35: threshold_cleaners = n_agents
751 10    elif waste_frac > 0.25: threshold_cleaners = int(n_agents * 0.7)
752 11    elif waste_frac > 0.15: threshold_cleaners = int(n_agents * 0.5)
753 12    elif waste_frac > 0.05: threshold_cleaners = int(n_agents * 0.3)
754 13    else: threshold_cleaners = max(1, int(n_agents * 0.2))
755 14
756 15    # [STATIC role: cleaners are ALWAYS the lowest-id agents -- no rotation]
757 16    is_cleaner = (agent_id % n_agents) < threshold_cleaners
758 17
759 18    if is_cleaner:
760 19        best_orient, n_waste = find_best_clean_orientation(env, agent_id)
761 20        if n_waste >= 1:
762 21            cur_orient = int(env.agent_orient[agent_id])
763 22            if best_orient == cur_orient: return 8 # CLEAN
764 23            if (cur_orient + 1) % 4 == best_orient: return 5 # ROTATE_RIGHT
765 24            elif (cur_orient - 1) % 4 == best_orient: return 4 # ROTATE_LEFT
766 25            else: return 4
767 26            # [Competitor-aware BFS toward waste in own row band]
768 27            c_idx = agent_id % n_agents
769 28            target_r = int(((c_idx + 0.5) / threshold_cleaners) * env.height)
770 29            # ... pick waste cell minimizing dist + |row - target_r| + 50 *
771 30            # closer_cleaners
772 31            # ... bfs_toward and direction_to_action
773 32        else:
774 33            # GATHERER: row-banded sweep, competitor-aware
775 34            # assigned_row = ((g_idx + 0.5) / n_gatherers) * env.height
776 35            # ... pick alive apple minimizing dist + |row - assigned_row| + 50 *
777 36            # closer_gatherers
778 37            ...
779 38

```

781 D.2 Cleanup, $\Phi_{\min}$ (Gemini): time-rotated, geographically interleaved roles

782 Best run: $U=3.19$, $E=0.98$, $\min_i R_i=296$. Switching the outer objective from U to $\min_i R_i$ flips
783 the role-assignment scheme entirely: `role_idx` now depends on `agent_id + step/50`, so every
784 agent rotates through cleaner and gatherer duty in a 50-step cycle. The cleaner indices $\{1, 4, 8\}$ are
785 spatially *interleaved* along the river (rather than contiguous) so that whoever is on duty has a short
786 walk to their assigned river slice. This single change accounts for the $E=0.55 \rightarrow 0.98$ jump and the
787 $\min_i R_i=-182 \rightarrow 296$ swing reported in Table 2, with negligible efficiency loss.

Listing 7: Cleanup, $\Phi_{\min}$ – Gemini 3.1 Pro. Time-rotated roles with interleaved geography (verbatim, abridged).

```

788 1 def policy(env, agent_id) -> int:
789 2     if int(env.agent_timeout[agent_id]) > 0:
790 3         return 7 # STAND while removed
791 4
792 5     aid = int(agent_id)
793 6     step = int(env._step_count)
794 7     n = int(env.n_agents)
795 8
796 9     # [TIME ROTATION: every 50 steps each agent's role index advances]
797 10    if n == 10:
798 11        role_idx = (aid + step // 50) % 10
799 12        # [Cleaner indices {1,4,8} are SPATIALLY INTERLEAVED, not contiguous,
800 13        # so each on-duty cleaner has a short walk to its river slice]
801 14        if role_idx == 1: role_type, zone_idx = 'C', 0
802 15        elif role_idx == 4: role_type, zone_idx = 'C', 1
803 16        elif role_idx == 8: role_type, zone_idx = 'C', 2
804 17        else:
805 18            role_type = 'G'
806 19            g_map = {0:0, 2:1, 3:2, 5:3, 6:4, 7:5, 9:6}
807 20            zone_idx = g_map[role_idx]
808 21            num_cleaners, num_gatherers = 3, 7
809 22
810 23    if role_type == 'C':
811 24        # [Each on-duty cleaner owns one horizontal slice of the river]
812 25

```

```

813 25 row_min = int( zone_idx * env.height / num_cleaners)
814 26 row_max = int((zone_idx + 1) * env.height / num_cleaners) - 1
815 27
816 28 # [Score every (row, col, orientation) in the slice by beam yield;
817 29 # fire if facing >=2 waste, else move to a >=2-yield position]
818 30 best_positions, good_positions = set(), set()
819 31 for r in range(row_min, row_max + 1):
820 32     for c in range(10):
821 33         if env.walls[r, c] or env.waste[r, c]: continue
822 34         y_max = max(get_clean_yield(env, r, c, o) for o in range(4))
823 35         if y_max >= 2: best_positions.add((r, c))
824 36         if y_max >= 1: good_positions.add((r, c))
825 37     # ... rotate / CLEAN / bfs_to_target_set as appropriate
826 38 else:
827 39     # GATHERER: collect apples ONLY inside the rotating row band
828 40     # row_min/row_max for num_gatherers slices ...
838 41 ...

```

831 D.3 Cleanup, $\Phi_{\min}$ (Sonnet): independent rediscovery of duty rotation

832 This run was launched on a separate dedicated git branch from the Gemini maximin runs of List-
833 ing 7 and from $p_{\mathcal{R}}$ identical to the one shown in Appendix C.1, with no shared state. The researcher-
834 authored synthesizer prompt p on this branch ended up containing a rotation hint structurally similar
835 to Listing 4 but with different phrasing and a different recommended period (phase_length = 100
836 in a worked example, vs. SHIFT ≈ 50 in Listing 4); $\mathcal{M}$ then chose its own period of 50. The conver-
837 gence is therefore at the level of which artifacts $\mathcal{R}$ injects into p , not at the level of $\mathcal{M}$ inventing
838 rotation from a neutral prompt.

839 Best run: $U=2.93$, $E=0.83$, $\min_i R_i=154$. The Sonnet synthesizer arrives at the *same* structural
840 insight as the Gemini maximin run: a phase counter ($\text{agent_id} + \text{step}/50$) mod n rotates which
841 2 of the 10 agents clean at any given time, and a *separate* 200-step zone counter rotates which 5-
842 row band each collector sweeps. The two rotation periods (50 and 200 steps) ensure every agent
843 visits every cleaning slot and every apple zone within one episode, resulting in a structural fairness
844 invariant.

845 **An alternative Sonnet maximin run uses a different mechanism.** The second Sonnet $\Phi_{\min}$
846 run took a structurally distinct route. There, $\mathcal{R}$ never wrote a rotation template into p ; instead it
847 authored a “collective threshold” worked example in which *all* agents synchronously enter a cleaning
848 mode whenever $\text{waste_fraction}(\text{env}) > 0.22$ and a collecting mode when it drops below 0.08,
849 with no agent_id phase shift. The resulting synthesizer policy follows this pattern and reaches
850 $\min_i R_i=200$, slightly above the rotation policy of Listing 8, confirming that the convergence in
851 Table 3 is on the broader *class* (explicit fairness mechanism, 4/4 maximin) rather than on rotation
852 specifically.

Listing 8: Cleanup, $\Phi_{\min}$ – Sonnet 4.6 (rotation variant; see Appendix D.3 for the alternative syn-
synchronized clean/collect mechanism in the second Sonnet maximin run). Phase-rotated cleaning +
zone-rotated collection (verbatim, abridged).

```

853 1 def policy(env, agent_id) -> int:
854 2     if int(env.agent_timeout[agent_id]) > 0:
855 3         return 7 # STAND while removed
856 4
857 5     n = env.n_agents
858 6     step = env._step_count
859 7     wf = waste_fraction(env)
860 8
861 9     # --- ROLE ASSIGNMENT: fair 2-cleaner rotation ---
862 10    # [50-step phases x 20 phases / episode -> each agent cleans 4 phases (~20%)]
863 11    phase = step // 50
864 12    cleaner_rank = (agent_id + phase) % n # [cycles 0..9 fairly]
865 13    is_cleaner = cleaner_rank < 2 # [exactly 2 cleaners per phase]
866 14
867 15    # [Emergency override: all agents clean when waste crosses the spawn cliff]
868 16    in_emergency = wf >= 0.40
869 17    if in_emergency: is_cleaner = True
870 18
871 19    # [River split: cleaner_slot 0 -> top half, 1 -> bottom half]
872 20    cleaner_slot = (agent_id % 2) if in_emergency else cleaner_rank
873 21
874 22    if is_cleaner and wf > 0.04:
875 23        # [Scan 4 orientations from current position; if best beam shot

```

```

877 24 # covers >=3 waste cells, fire; else step to a strictly better neighbor]
878 25 # ... (best_o, best_cnt) = max over 4 directions
879 26 # ... if best_cnt >= 3: rotate to best_o, then CLEAN
880 27 # ... else: scan 4 adjacent cells x 4 orientations for higher yield
881 28 # ... if no waste in beam range: bfs to waste_set in OWN river half,
882 29 # fall back to whole river
883 30
884 31 # --- COLLECTOR LOGIC: 5-zone rotating apple collection ---
885 32 # [5 zones x 200 steps = 1000 steps -> each agent visits every zone once;
886 33 # 2 agents share each zone at any time, reducing competition 10 -> 2]
887 34 bounds = zone_boundaries(env, 5)
888 35 zone_phase = step // 200
889 36 zone = (agent_id + zone_phase) % 5
890 37 z_start, z_end = bounds[zone], bounds[zone + 1]
891 38
892 39 zone_apples = {(int(env._apple_pos[i][0]), int(env._apple_pos[i][1]))
893 40               for i in range(env.n_apples) if env.apple_alive[i]
894 41               and z_start <= int(env._apple_pos[i][0]) < z_end}
895 42 # ... bfs_to_target_set(zone_apples), else bfs_nearest_apple as fallback
896 43 return 7

```

898 D.4 Gathering (Gemini): wall-aware Voronoi + spatiotemporal targeting

899 Best run: $U=2.47$, $E=0.98$, $\min_i R_i=580$. In Gathering, where costs are symmetric, no role differentiation is needed: the entire optimization is spatial and temporal. The synthesizer (i) partitions cells into BFS-Voronoi territories owned by individual agents, (ii) runs a single centralized BFS from its own position to score every reachable cell with exact wall-aware distances, and (iii) for both alive and respawning apples in its territory, computes the *earliest collection time* $\max(\text{walk_dist}, \text{respawn_timer})$ and targets the minimum. When the agent must wait on a dead apple, it steps onto a *non-spawn* adjacent cell rather than blocking a respawn point. This single sort key ($\max(\text{distance}, \text{timer})$) subsumes both “go to nearest apple” and “camp respawn” as special cases.

Listing 9: Gathering – Gemini 3.1 Pro. Voronoi territories + spatiotemporal priority targeting (verbatim, abridged).

```

908 1 def policy(env, agent_id) -> int:
909 2     if int(env.agent_timeout[agent_id]) > 0:
910 3         return 7
911 4
912 5     r, c = int(env.agent_pos[agent_id][0]), int(env.agent_pos[agent_id][1])
913 6     orient = int(env.agent_orient[agent_id])
914 7
915 8     # [Wall-aware Voronoi over the gridworld; spawn_points = all apple cells]
916 9     zones = voronoi_zones(env)
917 10    spawn_points = {(int(p[0]), int(p[1])) for p in env._apple_pos}
918 11
919 12    # [Single centralized BFS: exact distance + first move to every reachable cell]
920 13    distances, first_moves = {}, {}
921 14    q = deque([(r, c, 0, None)])
922 15    visited = {(r, c)}
923 16    while q:
924 17        cr, cc, d, fm = q.popleft()
925 18        distances[(cr, cc)] = d
926 19        first_moves[(cr, cc)] = fm
927 20        for dr, dc in [(-1,0),(1,0),(0,-1),(0,1)]:
928 21            nr, nc = cr + dr, cc + dc
929 22            if 0 <= nr < env.height and 0 <= nc < env.width and not env.walls[nr,
930 23                nc]:
931 24                if (nr, nc) not in visited:
932 25                    visited.add((nr, nc))
933 26                    q.append((nr, nc, d + 1, fm if fm is not None else (dr, dc)))
934 27
935 28    # [Both alive AND respawning apples in MY zone, with their respawn timer]
936 29    my_zone_spawns = []
937 30    for i in range(env.n_apples):
938 31        pr, pc = int(env._apple_pos[i][0]), int(env._apple_pos[i][1])
939 32        if zones.get((pr, pc)) == agent_id and (pr, pc) in distances:
940 33            my_zone_spawns.append((pr, pc, int(env.apple_timer[i]), distances[(pr,
941 34                pc)]))
942 35
943 36    # [SPATIOTEMPORAL PRIORITY: earliest collection time first.
944 37    # score = max(walk_distance, respawn_timer); ties broken by sooner timer]
945 38    my_zone_spawns.sort(key=lambda x: (max(x[3], x[2]), x[2], x[3]))
946 39
947 40    for pr, pc, timer, dist in my_zone_spawns:

```

```

949 39         if timer == 0:                                # [Alive apple: walk straight to it]
950 40             if dist == 0: return 7
951 41             fm = first_moves[(pr, pc)]
952 42             return direction_to_action(fm[0], fm[1], orient)
953 43         else:                                          # [Dead apple: camp on safe adjacent cell]
954 44             # ... pick nearest neighbor (nr,nc) that is NOT in spawn_points,
955 45             #         navigate there and STAND while waiting for respawn
956 46             ...
957 47         # [Global poaching fallback if zone is empty: nearest alive apple anywhere]
958 48         ...

```

960 D.5 Cross-condition observations

961 Several patterns are visible across these four listings (and confirmed by inspecting the remaining 8
962 runs not shown):

- 963 • *Role assignment encodes the welfare objective.* Static $\text{agent_id} < \tau$ (Listing 6) maximizes
964 U but harms $\min_i R_i$. Time-rotated $(\text{agent_id} + \text{step} // T) \% n$ (Listings 7, 8) maximizes
965 $\min_i R_i$ at $\leq 1\%$ efficiency cost (Gemini).
- 966 • *Convergent rediscovery across LLMs.* Gemini and Sonnet, on independent runs with separate
967 dedicated git branches, both arrive at the $(\text{id} + \text{phase}) \% n$ duty-rotation idiom with similar
968 phase lengths ($T \approx 50$ steps) in 3/4 maximin runs, and both omit any explicit fairness mechanism
969 under efficiency. The remaining 1/4 maximin run (Sonnet) converges on a structurally distinct
970 synchronized clean/collect mechanism with comparable maximin (Appendix D.3). The point of
971 convergence is the *class* of mechanism ($\mathcal{R}$ writes an explicit fairness mechanism into p under
972 $\Phi_{\min}$, never under Φ_U), not necessarily the rotation idiom itself.
- 973 • *Game structure dictates strategy class.* Cleanup policies are dominated by role-assignment logic
974 (cleaner vs. gatherer); Gathering policies are dominated by territory partitioning and respawn-
975 timer reasoning, with no role differentiation at all. The researcher does not need to be told which
976 class of strategy to write.
- 977 • *Earliest-collection-time as a unifying score.* The Gathering policy’s $\text{max}(\text{walk}, \text{timer})$ key (List-
978 ing 9) is structurally identical across the 4 Gathering runs (both Gemini and Sonnet), differing only
979 in how the Voronoi diagram is computed, e.g. by Manhattan approximation, fully BFS-based, or
980 cached helper.