
Discovering Cooperative Pipelines: Autoresearch for Sequential Social Dilemmas

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 We study *two-level autoresearch*: an outer-loop AI agent that autonomously re-
2 designs the inner-loop pipeline of an LLM policy-synthesis system for multi-agent
3 Sequential Social Dilemmas (SSDs). A *researcher agent* $\mathcal{R}$ (Claude Opus 4.6,
4 run as a coding agent) reads the inner-loop source code, edits system prompts,
5 feedback functions, helper libraries, and iteration logic, runs evaluations, and de-
6 cides what to keep, following the *autoresearch* paradigm. Across 12 autonomous
7 runs spanning two games (Cleanup and Gathering), two policy-synthesizer LLMs
8 (Claude Sonnet 4.6, Gemini 3.1 Pro), and two welfare objectives (utilitarian ef-
9 ficiency and Rawlsian maximin), the researcher reliably exceeds hand-designed
10 baselines, sharply tightens run-to-run variance, and substantially outperforms
11 prompt-only optimization (GEPA), which cannot overcome the stochasticity of
12 LLM code generation. The discovered pipelines are genuinely a function of the
13 welfare objective: under maximin the researcher independently rediscovers *duty*
14 *rotation*—never found under efficiency optimization—supporting an *information-*
15 *design* interpretation in which the researcher chooses what to expose to the bound-
16 edly rational synthesizer. The Cleanup–Gathering contrast further reveals that
17 *game structure* dictates whether fairness needs explicit optimization: in *provi-*
18 *sion* dilemmas it requires designed mechanisms, while in *restraint* dilemmas
19 it emerges as a byproduct of coordination. Code at [https://github.com/](https://github.com/anon590/llm-policies-social-dilemmas)
20 [anon590/llm-policies-social-dilemmas](https://github.com/anon590/llm-policies-social-dilemmas).

21 1 Introduction

22 Sequential Social Dilemmas (SSDs) [1] are the multi-agent analogue of the prisoner’s dilemma
23 in temporally rich Markov games: individually rational play leads to collectively suboptimal out-
24 comes through pollution, over-harvesting, or open conflict. Standard multi-agent reinforcement
25 learning (MARL) struggles in this regime due to credit assignment, non-stationarity, and large
26 joint action spaces [2]. A complementary approach, recently introduced by Gallego [3], replaces
27 parameter-space optimization with *algorithm-space synthesis*: a frozen LLM writes a Python policy
28 function, evaluates it in self-play, and iteratively refines it from performance feedback. A single
29 generation step can produce coordination logic—territory partitioning, role assignment, conditional
30 cooperation—that would take millions of RL episodes to discover.

31 This shifts where the design problem lives, rather than removing it. The inner-loop pipeline that
32 drives the synthesizer has many free parameters: which system prompt, which feedback variables,
33 which helper functions, how many refinement steps. Each materially affects the resulting policies,
34 and prior work tuned them by hand. A natural question follows: *can an AI agent design the pipeline?*

35 We answer affirmatively with a two-level autoresearch framework. An *outer-loop researcher agent*
36 (Claude Opus 4.6, run as a coding agent via the Claude Code CLI) edits the source files of an *inner-*

loop policy synthesizer (Gemini 3.1 Pro or Claude Sonnet 4.6), runs evaluations on held-out seeds, and keeps modifications that improve a fixed welfare objective Φ . The outer agent operates on an ordinary git repository—reading code, writing diffs, running shell commands—without task-specific scaffolding, mirroring the autoresearch paradigm of Karpathy [4] for single-GPU LLM pretraining. Although the inner-loop SSDs are gridworld benchmarks rather than physical systems, the outer-loop discovery process itself runs under conditions a deployed discovery agent faces: noisy multi-seed evaluations, stochastic code generation, an LLM-evaluation budget that bounds how often Φ can be queried, and a heterogeneous Python repository the agent must navigate end-to-end on its own.

Our contributions are:

- A general two-level framework that delegates the design of an LLM synthesis pipeline to a coding agent operating on a real software repository (Section 3).
- The first instantiation of the autoresearch paradigm in a multi-agent decision-making domain, with 12 autonomous runs across two SSDs (Cleanup and Gathering), two policy LLMs, and two welfare objectives (utilitarian efficiency U and Rawlsian maximin $\min_i R_i$). The researcher reliably exceeds hand-designed baselines and dominates prompt-only optimization (GEPA) [5] (Section 4).
- A mechanism-design interpretation supported by the qualitatively different pipelines the agent produces under different welfare objectives—including the autonomous discovery of *time-based duty rotation* as a fairness mechanism, never found under efficiency optimization.

2 Background

We build on the iterative LLM policy synthesis framework of Gallego [3], which serves as the (frozen) inner loop of our two-level system. This section recalls the SSD formalism, the social outcome metrics, and the base synthesis loop.

2.1 Sequential Social Dilemmas

A Sequential Social Dilemma is a partially observable Markov game $\mathcal{G} = \langle N, \mathcal{S}, \{\mathcal{A}_i\}_{i=1}^N, T, \{R_i\}_{i=1}^N, H \rangle$ with N agents, state space $\mathcal{S}$ (the gridworld configuration), per-agent action spaces $\mathcal{A}_i$, transition function T , reward functions R_i , and episode horizon H [1]. Beyond the dilemma’s matrix-game structure, SSDs add temporal richness: agents must learn *when* and *where* to cooperate, not just *whether* to.

We study two canonical SSDs that capture complementary dilemma types.

Cleanup [6] is a *public goods provision* game ($N=10$). The map has two regions: a river that accumulates waste, and an orchard where apples grow. Apples regrow only when river pollution is below threshold. Each agent can fire a cleaning beam (cost -1) that removes waste, collect apples ($+1$ each), or fire a tagging beam (cost -1 , inflicting -50 on the target and removing it for 25 steps). The dilemma: cleaning is costly but its benefits are public, so purely self-interested agents free-ride.

Gathering [1, 7] is a *common pool resource* game ($N=4$). Agents collect shared apples on a fixed respawn timer and may fire tagging beams to temporarily remove rivals. The dilemma: agents can coexist and share resources, or aggress to monopolize them; aggression wastes time and reduces total welfare.

Both games use 8–9 discrete actions (movement, rotation, beam, stand, optionally clean) and episodes of $H=1000$ steps. The two dilemmas differ in cost structure—asymmetric provision (cleaners pay, all benefit) vs. symmetric restraint (every agent faces the same temptation)—a distinction that drives our key findings.

80 **Social metrics.** Following Perolat et al. [7], let $R_i = \sum_{t=0}^{H-1} r_i^t$ denote agent i 's episode return.
 81 We evaluate four social outcomes:

$$U = \frac{1}{H} \sum_{i=1}^N R_i \quad (\text{efficiency}) \quad (1)$$

$$E = 1 - \frac{\sum_{i,j} |R_i - R_j|}{2N \sum_i R_i} \quad (\text{equality}) \quad (2)$$

$$S = \frac{1}{N} \sum_{i=1}^N \bar{t}_i \quad (\text{sustainability}) \quad (3)$$

$$P = \frac{1}{H} \sum_{t=0}^{H-1} |\{i : \text{active}_i^t\}| \quad (\text{peace}) \quad (4)$$

82 where $\bar{t}_i$ is the mean timestep at which agent i collects positive reward (higher $\Rightarrow$ resources preserved
 83 later) and active_i^t indicates that agent i is not tagged out at step t . We additionally consider the
 84 *maximin* (Rawlsian) welfare criterion $\min_i R_i$, which optimizes the worst-off agent's return and
 85 serves as the second objective in our experiments.

86 2.2 Iterative LLM Policy Synthesis

87 Let Π denote the space of *programmable policies*: deterministic functions $\pi : \mathcal{S} \times [N] \rightarrow \mathcal{A}$ ex-
 88 pressed as executable Python code. Each policy has access to the full environment state and a
 89 library of helpers (BFS pathfinding, beam targeting, coordinate transforms). This state access is a
 90 deliberate design choice: programmatic policies operate in algorithm space rather than the reactive
 91 observation-to-action space of neural policies, which lets a single LLM generation step encode rich
 92 coordination logic.

93 A frozen LLM $\mathcal{M}$ acts as the *policy synthesizer*. Given a system prompt p describing the environ-
 94 ment API and a feedback prompt q_k , it produces a new policy

$$\pi_{k+1} = \mathcal{M}(p, q(\pi_k, \mathcal{F}_k)),$$

95 where π_k is the previous policy (its source code) and $\mathcal{F}_k$ is the evaluation feedback. All N agents ex-
 96 ecute the same policy in homogeneous self-play, and evaluation over a set of random seeds S yields
 97 the mean per-agent return $\bar{r}_k$ and the social metrics vector $\mathbf{m}_k = (U_k, E_k, S_k, P_k)$. Each generated
 98 policy passes an AST-based safety check (blocking eval, file I/O, network access) followed by a
 99 short smoke test; failures trigger regeneration (up to R attempts) with the error message appended
 100 to the prompt.

101 **Feedback.** The hand-designed feedback used by Gallego [3] as our starting point packages the
 102 previous policy's source code together with all available evaluation signals:

$$\mathcal{F}_k = (\text{code}(\pi_k), \bar{r}_k, \mathbf{m}_k, \mathbf{d}), \quad (5)$$

103 where $\mathbf{d}$ contains natural-language definitions of each social metric. Exposing the full social-metrics
 104 vector (rather than the scalar reward alone) disambiguates distinct failure modes that share the same
 105 average return—e.g., under- vs. over-cleaning in Cleanup—and consistently improves cooperative
 106 outcomes. On Cleanup, this baseline feedback reaches $U=1.93/2.70$ (mean/max) with Gemini and
 107 $U=0.86/1.56$ with Sonnet across 4 runs each. The choice of feedback content is a single design
 108 decision within a much larger pipeline configuration space; our two-level framework (Section 3)
 109 opens the full space to automated search.

110 3 Two-Level Framework

111 We introduce a two-level system where a *researcher agent* $\mathcal{R}$ autonomously discovers configurations
 112 that optimize the output of an inner-loop system. While we instantiate this for multi-agent policy
 113 synthesis, the architecture is general: any pipeline where an LLM generates artifacts, evaluates them,
 114 and iterates can serve as the inner loop. The key insight is that the entire inner-loop codebase—not
 115 just the feedback level—is a designable artifact that an AI agent can search over. Figure 1 illustrates
 116 the architecture.




Figure 1: Two-level automated research framework (Algorithm 1).

Algorithm 1 Two-Level Automated Research

Require: Game $\mathcal{G}$, policy synthesizer $\mathcal{M}$, researcher $\mathcal{R}$, system prompt $p_{\mathcal{R}}$, initial config c_0 , outer iterations J , ground-truth objective Φ , held-out seeds S_{ho}

Ensure: Best configuration c^* , best policy π^*

```

1:  $\pi_0^* \leftarrow \text{INNERLOOP}(\mathcal{M}, \mathcal{G}, c_0)$ 
2:  $\Phi_0 \leftarrow \text{EVAL}(\pi_0^*; \mathcal{G}, S_{\text{ho}})$ 
3:  $\text{history} \leftarrow \{(c_0, \Phi_0, \mathbf{m}_0, \emptyset)\}$ 
4: for  $j = 1, \dots, J$  do
5:    $c_j \leftarrow \mathcal{R}(p_{\mathcal{R}}, \text{CODE}(c_{j-1}), \text{history})$ 
6:   if  $\neg \text{VALIDATECONFIG}(c_j)$  then  $\text{retry} \leq R$ 
7:      $\pi_j^* \leftarrow \text{INNERLOOP}(\mathcal{M}, \mathcal{G}, c_j)$ 
8:      $\Phi_j \leftarrow \text{EVAL}(\pi_j^*; \mathcal{G}, S_{\text{ho}})$ 
9:      $\Delta_j \leftarrow \text{DIFF}(c_{j-1}, c_j)$  // code diff
10:     $\text{history} \leftarrow \text{history} \cup \{(c_j, \Phi_j, \mathbf{m}_j, \Delta_j)\}$ 
11:  end for
12:  $j^* \leftarrow \arg \max_j \Phi_j$ 
13: return  $c_{j^*}, \pi_{j^*}^*$ 

```

Table 1: Configuration components modifiable by the researcher. The environment simulator, ground-truth evaluation, and policy LLM weights are frozen.

	Scope	Examples
p	System prompt	Hints, worked examples
ϕ	Feedback fn	Metric subset, framing, derived metrics, trends
$\mathcal{H}$	Helper library	Pathfinding, zones
ι	Iteration logic	K , eval seeds, best-of- n
v	Validation	Safety checks, smoke tests

3.1 Configuration Space

Let $\mathcal{C}$ denote the space of *pipeline configurations*. Each configuration $c \in \mathcal{C}$ specifies the full inner-loop setup:

$$c = (p, \phi, \mathcal{H}, \iota, v) \quad (6)$$

where p is the system prompt, ϕ is the feedback construction function (which content to include, how to frame it, whether to inject adaptive hints), $\mathcal{H}$ is the helper function library, ι specifies the iteration logic (number of iterations K , sampling strategy, temperature schedule), and v is the validation pipeline. Table 1 provides concrete examples.

The hand-designed feedback of [3] corresponds to a single fixed instantiation of ϕ . Our framework opens the full configuration space to automated search.

3.2 Inner Loop (Policy Synthesis)

Given a configuration c , the inner loop executes K iterations of LLM policy synthesis:

$$\pi_c^* = \text{INNERLOOP}(\mathcal{M}, \mathcal{G}, c) \quad (7)$$

Each iteration k proceeds in four stages, following [3]:

- Synthesize.** The policy LLM $\mathcal{M}$ receives the system prompt p , the previous policy’s source code π_{k-1} , and feedback $\mathcal{F}_{k-1}$ constructed by ϕ . It generates a new Python policy function π_k that has access to full environment state and the helper library $\mathcal{H}$.
- Validate.** The generated code undergoes AST-based safety checks (blocking dangerous operations such as file I/O and network access) followed by a short smoke test. Failures trigger re-generation (up to R retries), with the error message appended to the prompt.
- Evaluate.** All N agents execute the same policy π_k in self-play over $|S|$ random seeds. The evaluation yields the mean per-agent reward $\bar{r}_k$ and the social metrics vector $\mathbf{m}_k = (U_k, E_k, S_k, P_k)$.

137 **4. Feedback.** The feedback function ϕ constructs the prompt for the next iteration from
 138 $(\pi_k, \bar{r}_k, \mathbf{m}_k)$, packaging the previous policy’s source code together with the scalar reward, the
 139 social metrics vector, their natural-language definitions, and any adaptive diagnostics that ϕ in-
 140 jects.

141 The inner loop output is evaluated on held-out seeds via a fixed ground-truth objective:

$$\Phi(c) = \text{EVAL}(\pi_c^*; \mathcal{G}, S_{\text{held-out}}) \quad (8)$$

142 where Φ maps from social outcomes to a scalar. We consider two objectives:

$$\Phi_U = U \quad (\text{utilitarian efficiency}) \quad (9)$$

$$\Phi_{\min} = \min_i R_i \quad (\text{Rawlsian maximin}) \quad (10)$$

143 3.3 Outer Loop (Automated Research)

144 The researcher agent $\mathcal{R}$ iteratively modifies the pipeline configuration. Following the autoresearch
 145 paradigm [4], $\mathcal{R}$ operates on the inner-loop codebase as a modifiable artifact, proposing changes,
 146 observing outcomes, and refining. The procedure is formalized in Algorithm 1.

147 At each outer iteration j , the researcher $\mathcal{R}$ receives:

- 148 • The **full source code** of the current configuration c_{j-1} (prompts, feedback construction, helpers,
 149 iteration logic).
- 150 • The **experiment history**: for each prior iteration, the code diff Δ_i , ground-truth score Φ_i , and
 151 social metrics vector $\mathbf{m}_i$.
- 152 • The **environment source code** (read-only), enabling the researcher to reason about game mechan-
 153 ics.

154 The researcher proposes a new configuration c_j by generating code modifications. Concretely, $\mathcal{R}$
 155 is a coding agent (Claude Code CLI) that operates on a real software repository: it reads and edits
 156 Python source files, runs shell commands, inspects evaluation outputs, and commits changes to
 157 a dedicated git branch—the same workflow a human researcher would follow. This mirrors the
 158 autoresearch setup, where an AI coding agent modifies `train.py` and observes validation loss;
 159 here, the researcher modifies the inner-loop codebase and observes social welfare.

160 3.4 Connection to Automated Mechanism Design

161 The two-level structure admits a mechanism design interpretation. The researcher $\mathcal{R}$ acts as a *mech-*
 162 *anism designer*: it controls the information structure (what metrics to reveal, how to frame them),
 163 the action space (what helper functions are available), and the incentive structure (how feedback is
 164 presented) under which the policy synthesizer $\mathcal{M}$ operates. The synthesizer acts as the *agent* within
 165 the designed mechanism.

166 This connects to the automated mechanism design literature [8], where a principal designs rules to
 167 induce desired behavior from self-interested agents. In our setting:

- 168 • The **principal** is the researcher $\mathcal{R}$, optimizing the ground-truth objective Φ .
- 169 • The **agent** is the synthesizer $\mathcal{M}$, optimizing per-agent reward as instructed.
- 170 • The **mechanism** is the configuration c : prompts, feedback, helpers, iteration logic.
- 171 • The **outcome** is the social welfare of the resulting multi-agent policy π_c^* .

172 A key distinction from classical mechanism design: $\mathcal{M}$ is not *strategically* self-interested—it fol-
 173 lows instructions—but it has *bounded rationality* in the sense that its ability to synthesize effective
 174 policies depends on the information and tools provided. The researcher’s task is thus closer to *infor-*
 175 *mation design* [9]: choosing what to reveal to help $\mathcal{M}$ navigate the cooperation–defection tension.
 176 Our experiments (Section 4) show that the researcher designs qualitatively different information
 177 structures depending on the welfare objective Φ , supporting this interpretation empirically.

4 Experiments

4.1 Setup

We conduct 12 autonomous researcher runs across a factorial design. The researcher agent $\mathcal{R}$ is Claude Opus 4.6, invoked via the Claude Code CLI as a coding agent. Each run operates on a dedicated git branch of a real Python codebase: $\mathcal{R}$ edits source files in `pipeline/`, executes evaluation scripts, reads metric outputs, and iterates—without human intervention.

Design. For Cleanup: 2 policy LLMs $\times$ 2 objectives $\times$ 2 replications = 8 runs. For Gathering: 2 policy LLMs $\times$ 1 objective $\times$ 2 replications = 4 runs. Maximin runs are unnecessary for Gathering because efficiency optimization alone achieves equality $E > 0.94$ (the game’s symmetric cost structure makes fairness a free byproduct of coordination).

Models. Policy synthesizer $\mathcal{M}$: Gemini 3.1 Pro (Google) or Claude Sonnet 4.6 (Anthropic). Both use extended thinking. These are the same models evaluated in [3]. We additionally test with Gemma 4 26B-A4B-IT (Google), a smaller open-weight model, to probe the framework’s behavior when the policy synthesizer has substantially lower capability (Appendix C).

Baselines. The hand-designed feedback configuration from [3] serves as the initial pipeline c_0 for all runs. On Cleanup ($N=10$), this baseline achieves $U=1.93/2.70$ (mean/max) with Gemini and $U=0.86/1.56$ with Sonnet across 4 runs each. We additionally compare against GEPA [5], an automated prompt optimization method that iteratively refines the system prompt via LLM reflection. GEPA optimizes only the system prompt p , whereas our researcher modifies the full pipeline $(p, \phi, \mathcal{H}, \iota, v)$. We give GEPA a matched compute budget: the same number of optimization steps as outer iterations used by our method (each GEPA step costs ~ 3 policy evaluations, comparable to our inner loop’s $K+1=4$).

Inner loop parameters. $K=3$ synthesis iterations (or $K=2$ if the researcher elects), evaluated over $|S|=5\text{--}12$ random seeds (researcher-configurable). The researcher ran $J=3\text{--}17$ outer iterations per run (mean: 8.3), with a keep rate of 18%–100% (mean: 63%).

4.2 Results

Table 2 presents the main results, comparing autoresearch against the hand-designed baseline of [3] and GEPA prompt optimization [5].

Finding 1: The researcher reliably improves over hand-designed baselines and outperforms prompt-only optimization. All 12 runs show substantial improvement regardless of starting point (Figure 2). In Cleanup, the researcher exceeds the unmodified pipeline baseline by 66% with Gemini ($U: 1.93 \rightarrow 3.20$) and 263% with Sonnet ($U: 0.86 \rightarrow 3.12$). Strikingly, the Sonnet improvement nearly closes the gap with Gemini ($U=3.12$ vs. 3.20), despite a $2\times$ baseline disadvantage (0.86 vs. 1.93). Final efficiency values cluster tightly around $U \approx 3.1\text{--}3.2$ despite baselines varying from 0.29 to 2.70 across individual runs, suggesting the researcher reliably discovers the performance ceiling of each policy LLM. Beyond improving the mean, the researcher yields remarkably consistent results: under efficiency optimization, Gemini’s runs cluster within $U=3.20\text{--}3.25$ (gap 0.05) despite baselines spanning 1.15–2.70, and Sonnet’s within 3.12–3.14 (gap 0.02) from baselines of 0.38–1.56. In Gathering, all four runs converge to $U=2.47\text{--}2.52$ from baselines ranging 0.03–2.42. The researcher discovers pipeline modifications—structured helpers (App. D.1), strategic hints (App. D.2), tailored feedback (App. D.3)—that constrain the policy LLM toward consistently good solutions.

Compared to GEPA (Table 2), autoresearch achieves $2.4\times$ higher efficiency under Φ_U in Cleanup ($U=3.20$ vs. 1.34), $2\times$ higher maximin under $\Phi_{\min}$ (290 vs. 143), and 20% higher efficiency in Gathering ($U=2.49$ vs. 2.08). Critically, GEPA shows far wider spread across runs on the harder fairness objective ($\Phi_{\min}$ maximin in Cleanup: 245/41 vs. autoresearch’s 296/284). The gap widens further with Sonnet: on Cleanup Φ_U , GEPA stalls at $U=1.04$ ($3\times$ below autoresearch’s 3.12), and on $\Phi_{\min}$ its two runs split between a modestly positive solution ($\min_i R_i=6$, $U=1.60$) and a pathological “everyone cleans, nobody eats” regime (perfect equality $E=1.00$ but negative efficiency $U=-2.87$).

Table 2: Results. Mean /max across 2 runs per condition (4 for Cleanup baselines). Baseline: unmodified pipeline from [3]. GEPA: automated prompt optimization [5] with matched compute budget.

Policy LLM	Target	U	E	$\min_i R_i$
<i>Cleanup (N=10) — Baseline</i>				
Gemini 3.1 Pro	—	1.93/2.70	0.17/0.62	−159/−84
Sonnet 4.6	—	0.86/1.56	−0.02/0.25	−151/−59
<i>Cleanup (N=10) — GEPA [5]</i>				
Gemini 3.1 Pro	Φ_U	1.34/1.37	−0.10/−0.05	−149/−126
	$\Phi_{\min}$	1.76/2.76	0.90/0.96	143/245
Sonnet 4.6	Φ_U	1.04/1.20	−0.59/−0.21	−164/−126
	$\Phi_{\min}$	−0.63/1.60	0.77/1.00	−147/6
<i>Cleanup (N=10) — Autoresearch</i>				
Gemini 3.1 Pro	Φ_U	3.20 /3.25	0.55/0.61	−211/−182
	$\Phi_{\min}$	3.16/3.19	0.98 /0.98	290 /296
Sonnet 4.6	Φ_U	3.12 /3.14	0.66/0.70	−196/−146
	$\Phi_{\min}$	2.57/2.93	0.91 /0.97	179 /200
<i>Gathering (N=4) — Baseline</i>				
Gemini 3.1 Pro	—	2.04/2.42	0.90/0.98	412/571
Sonnet 4.6	—	0.03/0.03	0.54/0.54	0/0
<i>Gathering (N=4) — GEPA [5]</i>				
Gemini 3.1 Pro	Φ_U	2.08/2.35	0.94/0.96	436/518
Sonnet 4.6	Φ_U	1.20/1.23	0.63/0.71	86/123
<i>Gathering (N=4) — Autoresearch</i>				
Gemini 3.1 Pro	Φ_U	2.49/2.51	0.98/0.98	582/582
Sonnet 4.6	Φ_U	2.52/2.52	0.96/0.98	576/597




Figure 2: Efficiency (U) across researcher iterations for all 12 runs. *Left*: Cleanup ($N=10$) with 8 runs across 2 LLMs $\times$ 2 objectives (solid lines = kept experiments, open circles = discarded). Dashed horizontal line: hand-designed baseline from [3]. All runs converge to $U \approx 3.1$ – 3.2 despite diverse starting points. *Right*: Gathering ($N=4$) with 4 runs (2 per LLM). All converge to $U \approx 2.5$ within 2–4 iterations despite baselines ranging from 0.03 to 2.42.

227 and $\min_i R_i = -299$, while autoresearch with the same model reaches $\min_i R_i = 179/200$ reliably.
 228 This highlights the advantage of modifying the full pipeline—helpers, feedback, and iteration logic
 229 provide structured scaffolding that reduces run-to-run variance, whereas prompt-only optimization
 230 cannot address the underlying stochasticity of code generation.

231 **Finding 2: No efficiency–fairness tradeoff in Cleanup (Gemini).** Maximin-optimized Gem-
 232 ini pipelines sacrifice only 1% efficiency (U : 3.16 vs. 3.20) while achieving near-perfect equal-
 233 ity (E : 0.98 vs. 0.55) and transforming maximin from deeply negative baselines (−99 to −84) to

234 $\min_i R_i = 290$. The researcher discovers *fair duty rotation* (Listing 6; primed by the prompt rewrite
 235 of Listing 3)—time-based cycling using `agent_id` and `env._step_count`—simultaneously im-
 236 proving worst-off welfare *and* collective output. Because cleaning is a public good, distributing the
 237 cleaning cost fairly ensures enough cleaners to sustain apple production.

238 For Sonnet, there is a moderate tradeoff: efficiency drops from 3.12 to 2.57 under maximin opti-
 239 mization. The gap reflects Sonnet’s harder time implementing complex coordination mechanisms
 240 (role rotation, zone assignment) from strategic hints alone.

241 **Finding 3: Game structure determines whether fairness requires explicit optimization.** In
 242 Cleanup, where cleaning costs are borne asymmetrically (cleaners pay -1 , free-riders collect ap-
 243 ples), baseline equality ranges from $E=0.04$ to 0.62 , and maximin optimization is required to reach
 244 $E > 0.9$. In Gathering, where all agents face a symmetric landscape, efficiency optimization alone
 245 achieves $E > 0.94$ across all 4 runs—no separate maximin runs are needed.

246 This generalizes: in *provision* dilemmas with asymmetric costs, fairness requires designed mecha-
 247 nisms (role rotation, duty sharing); in *restraint* dilemmas with symmetric costs, fairness emerges as
 248 a free byproduct of efficient coordination. The researcher independently discovers this—it creates
 249 role differentiation pipelines only for Cleanup, and pure spatial-coordination pipelines for Gather-
 250 ing.

251 **Finding 4: Convergent discovery of qualitatively different strategies per objective.** Despite
 252 fully independent runs, the researcher converges on the same core strategies within each condition
 253 (Table 3, Appendix A). In Cleanup, waste-counting helpers and spatial zone partitioning appear
 254 across all 8 runs. The key qualitative difference is *role rotation*, appearing in 4/4 maximin runs
 255 (Listings 6, 7) but 0/4 efficiency runs (Listing 5). Under efficiency optimization, the researcher
 256 assigns static cleaning roles (some agents always clean), achieving high collective output at the cost
 257 of equality. Under maximin optimization, it discovers rotation as a fairness mechanism.

258 In Gathering, the researcher discovers BFS-Voronoi territory partitioning and respawn-timer aware-
 259 ness (Listings 2, 8), with no role differentiation—optimization is purely spatial (who collects which
 260 apples) and temporal (respawn-aware positioning).

261 **Common failure modes.** Several patterns led to discarded experiments across all runs: (1) *over-*
 262 *prescription*—adding too many strategic hints confuses the policy LLM, causing generation failures;
 263 (2) *iteration regression*— $K=4$ or $K=5$ inner iterations often cause catastrophic regression as the
 264 LLM “over-refines” a working policy; (3) *feedback overload*—verbose per-agent diagnostics cause
 265 over-correction rather than gradual improvement; (4) *variance exposure*—identical configurations
 266 can yield $U=3.2$ then $U=0.09$ on different seeds, requiring multi-seed evaluation for stable signals
 267 (per-run $K, |S|$ in Table 6).

268 5 Related Work

269 **Sequential social dilemmas.** SSDs were introduced by Leibo et al. [1] (Gathering) and extended
 270 to public-goods settings by Hughes et al. [6] (Cleanup); Perolat et al. [7] formalized the social
 271 outcome metrics (U, E, S, P) used here. Subsequent work studies inequity aversion, intrinsic moti-
 272 vation, and reputational mechanisms for promoting cooperation in MARL agents. We complement
 273 this line by automating the search for cooperative *programs* rather than evolving neural policies,
 274 and by treating the welfare objective Φ as a designable parameter that the outer agent optimizes
 275 for—obtaining qualitatively different cooperative behaviors (static role assignment vs. duty rotation)
 276 under different objectives without changing the environment.

277 **LLMs for policy and program synthesis.** FunSearch [10] evolves programs for combinato-
 278 rial discovery; Eureka [11] synthesizes reward functions from environment source code; Voy-
 279 ager [12] and Code as Policies [13] generate executable skill code for single-agent embodied control;
 280 ReEvo [14] evolves heuristics with reflective feedback. These works target single-agent settings or
 281 non-strategic optimization. Gallego [3], on which our inner loop is based, extended LLM program
 282 synthesis to the multi-agent SSD setting, where one program must coordinate N self-play copies.
 283 Our contribution is one level above: rather than tuning the synthesizer for one task, we let an outer
 284 agent rewrite the synthesis pipeline itself.

LLM reflection and prompt optimization. Reflexion [15], Self-Refine [16], OPRO [17], and GEPA [5] demonstrate that structured verbal feedback loops improve LLM outputs; ERL [18] internalizes such reflection via self-distillation. These methods optimize the prompt or the model’s reasoning trajectory in isolation. Our framework instead modifies the entire surrounding pipeline—system prompt, feedback construction, helper library, iteration logic, validation—making prompt optimization (cf. GEPA) one component of a strictly larger search space. The empirical gap is large: Section 4 shows full-pipeline autoresearch achieving $3\times$ higher efficiency in Cleanup and 37% higher in Gathering than GEPA at matched compute, with run-to-run spread an order of magnitude tighter on the harder fairness objective.

Automated AI research. A small but growing body of work delegates the design of ML pipelines to LLM coding agents. Karpathy’s autoresearch [4] runs a coding agent that modifies `train.py` for nanoGPT pretraining and is rewarded for validation loss. Our system shares this autoresearch architecture—coding agent + frozen evaluation harness + diff-based history—but targets a different inner loop (multi-agent policy synthesis instead of single-model pretraining) and a different objective (multi-agent social welfare instead of validation bits-per-byte). To our knowledge this is the first instantiation of the autoresearch paradigm in a multi-agent decision-making domain.

Automated mechanism and information design. Classical automated mechanism design [8] computes optimal allocation rules for self-interested agents under explicit incentive constraints, while information design [9] studies how a principal commits to a signaling policy that shapes a receiver’s behavior. Our outer agent solves a related but distinct problem: $\mathcal{M}$ is not strategically deceptive (it follows instructions) but is *boundedly rational*, so the researcher must decide what information, helpers, and structure to expose so that $\mathcal{M}$ writes policies achieving the principal’s welfare goal. Section 4 shows that the pipelines $\mathcal{R}$ produces are genuinely a function of the welfare objective, supporting this information-design framing empirically.

6 Discussion and Conclusion

We have presented a two-level framework where an AI coding agent autonomously discovers pipeline configurations that improve the output of an inner-loop LLM system. Applied to multi-agent policy synthesis in social dilemmas, the researcher agent reliably exceeds hand-designed baselines across 12 independent runs, and converges on qualitatively similar strategies within each game-objective condition. The framework requires no domain-specific scaffolding: the agent operates on a standard software repository using the same tools (file editing, shell commands, git) available to a human researcher.

Mechanism design in action. Our results empirically support the mechanism design interpretation of Section 3.4. The researcher acts as an information designer: under Φ_U , it reveals efficiency-oriented information (waste counts, zone assignments) that guides $\mathcal{M}$ toward productive but unequal role allocation (Listing 5). Under $\Phi_{\min}$, it additionally reveals fairness-oriented structure—rotation schedules and per-agent equity feedback (Listings 3, 4)—guiding $\mathcal{M}$ toward egalitarian coordination. The striking absence of a fairness-efficiency tradeoff with Gemini suggests that in public goods provision, the mechanism design problem has a cooperative optimum where both objectives align—the researcher discovers this independently.

Two types of dilemma. The Cleanup–Gathering contrast reveals a structural insight about social dilemmas: *provision* dilemmas (asymmetric costs, one agent’s contribution benefits all) require explicit fairness mechanisms, while *restraint* dilemmas (symmetric costs, each agent faces the same temptation) achieve fairness naturally when coordination improves. This distinction, well-known in the common-pool resource literature, is rediscovered here by an autonomous AI agent—suggesting it is a robust structural feature of these environments.

Human oversight and audit. Our system is fully autonomous by design, but its architecture offers natural affordances for human-in-the-loop oversight. Because the researcher $\mathcal{R}$ operates on a standard git repository—committing each pipeline modification to a dedicated branch—a domain expert can audit the full history of code diffs Δ_j alongside their metric outcomes $(\Phi_j, \mathbf{m}_j)$, exactly as they would review a colleague’s pull requests. The keep/discard decision (Algorithm 1, line 11)

is a lightweight intervention point: a human could override the accept threshold, veto modifications that introduce undesirable mechanisms (e.g., exploitative role assignments), or steer the researcher toward under-explored regions of the configuration space by editing the experiment history. More broadly, the separation between the researcher $\mathcal{R}$ and the frozen evaluation Φ creates a *delegation boundary*: the human defines *what* to optimize (the welfare objective), while the agent decides *how*. This mirrors the expert-in-the-loop design patterns identified in deployment-oriented work on AI agents for discovery, where trust calibration depends on the legibility of the agent’s decision trail rather than on real-time supervision.

Limitations. Our experiments use only 2 replications per condition, yielding high uncertainty on some estimates (e.g., Sonnet $\Phi_{\min}$: $U=2.57/2.93$). The two SSDs are small gridworld environments; scaling to larger, more complex games remains open. The researcher runs are not fully controlled—the number of outer iterations varies from 3 to 17 (mean 8.3)—making direct comparison across conditions approximate. The GEPA comparison uses a single run per condition; additional replications would strengthen confidence in the comparison.

Future work. Five directions emerge: (1) *new domains*—apply the two-level framework to other LLM-driven pipelines (code optimization, scientific experiment design, infrastructure tuning), where the same architecture applies with a different inner loop and objective Φ ; (2) *ablation*—freeze all configuration components except one to isolate the relative importance of prompt, feedback, helpers, and iteration logic; (3) *adversarial objectives*—set $\Phi = R_0$ (agent 0’s reward only) to test whether the researcher discovers exploitative pipeline configurations, connecting to the reward hacking analysis of [3]; (4) *cross-game transfer*—run the best Cleanup configuration on Gathering and vice versa, testing game-generalizability; (5) *heterogeneous policies*—extend to settings where agents run different code, enabling richer role specialization.

References

- [1] Joel Z Leibo, Vinicius Zambaldi, Marc Lanctot, Janusz Janssen, and Thore Graepel. Multi-agent reinforcement learning in sequential social dilemmas. In *Proc. 16th Conference on Autonomous Agents and MultiAgent Systems*, pages 464–473, 2017.
- [2] Lucian Buşoniu, Robert Babuška, and Bart De Schutter. A comprehensive survey of multi-agent reinforcement learning. *IEEE Transactions on Systems, Man, and Cybernetics, Part C (Applications and Reviews)*, 38(2):156–172, 2008.
- [3] Víctor Gallego. Cooperation and exploitation in LLM policy synthesis for sequential social dilemmas. In *Proc. Adaptive and Learning Agents Workshop (ALA 2026)*, 2026.
- [4] Andrej Karpathy. autoresearch: AI agents running research on single-GPU nanochat training automatically. <https://github.com/karpathy/autoresearch>, 2026.
- [5] Lakshya A Agrawal et al. GEPA: Reflective prompt evolution can outperform reinforcement learning. *arXiv preprint arXiv:2507.19457*, 2026.
- [6] Edward Hughes, Joel Z Leibo, Matthew Phillips, Karl Tuyls, Edgar A Dueñez-Guzmán, Antonio García Castañeda, Iain Dunning, Tina Zhu, Kevin R McKee, R. Koster, Heather Roff, and Thore Graepel. Inequity aversion improves cooperation in intertemporal social dilemmas. In *Neural Information Processing Systems*, 2018.
- [7] Julien Perolat, Joel Z Leibo, Vinicius Zambaldi, Charles Beattie, Karl Tuyls, and Thore Graepel. A multi-agent reinforcement learning model of common-pool resource appropriation. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- [8] Vincent Conitzer and Tuomas Sandholm. Complexity of mechanism design. In *Proc. 18th Conference on Uncertainty in Artificial Intelligence*, pages 103–110, 2002.
- [9] Emir Kamenica and Matthew Gentzkow. Bayesian persuasion. *American Economic Review*, 101(6):2590–2615, 2011.

- [10] Bernardino Romera-Paredes, Mohammadamin Barekatain, Alexander Novikov, Matej Balog, M Pawan Kumar, Emilien Dupont, Francisco JR Ruiz, Jordan S Ellenberg, Pengming Wang, Omar Fawzi, et al. Mathematical discoveries from program search with large language models. *Nature*, 625(7995):468–475, 2024.
- [11] Yecheng Jason Ma, William Liang, Guanzhi Wang, De-An Huang, Osbert Bastani, Dinesh Jayaraman, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Eureka: Human-level reward design via coding large language models. In *The Twelfth International Conference on Learning Representations*, 2024.
- [12] Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *Transactions on Machine Learning Research*, 2024.
- [13] Jacky Liang, Wenlong Huang, Fei Xia, Peng Xu, Karol Hausman, Brian Ichter, Pete Florence, and Andy Zeng. Code as policies: Language model programs for embodied control. In *2023 IEEE International Conference on Robotics and Automation (ICRA)*, pages 9493–9500, 2023.
- [14] Haoran Ye, Jiarui Wang, Zhiguang Cao, Federico Berto, Chuanbo Hua, Haeyeon Kim, Jinkyoo Park, and Guojie Song. ReEvo: Large language models as hyper-heuristics with reflective evolution. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [15] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: language agents with verbal reinforcement learning. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.
- [16] Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. Self-refine: Iterative refinement with self-feedback. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.
- [17] Chengrun Yang, Xuezhi Wang, Yifeng Lu, Hanxiao Liu, Quoc V Le, Denny Zhou, and Xinyun Chen. Large language models as optimizers. In *The Twelfth International Conference on Learning Representations*, 2024.
- [18] Taiwei Shi, Sihao Chen, Bowen Jiang, Linxin Song, Longqi Yang, and Jieyu Zhao. Experiential reinforcement learning. *arXiv preprint arXiv:2602.13949*, 2026.

A Additional Results

Table 3: Strategies discovered by the researcher in Cleanup across 4 efficiency-optimized and 4 maximin-optimized independent runs.

Strategy	Φ_U (4 runs)	$\Phi_{\min}$ (4 runs)
Waste-counting helpers	4/4	4/4
Zone/lane partitioning	3/4	4/4
Anti-regression feedback	3/4	2/4
Worked policy examples	2/4	3/4
Cleaning cost economics	2/4	3/4
Time-based role rotation	0/4	4/4

B Compute Requirements

Table 4 reports wall-clock time for all 12 autonomous runs and the 6 GEPA comparison runs. *Inner loop* measures total time spent on policy LLM generation and environment simulation across all evaluations; *total* (estimated from run-directory timestamps) additionally includes the researcher agent’s analysis, code editing, and decision-making between evaluations.




Figure 3: Maximin ($\min_i R_i$) across researcher iterations for the 4 Cleanup maximin-optimized runs. All runs transform deeply negative baselines (worst-off agents losing reward) into substantially positive values. Gemini reaches ~ 290 while Sonnet reaches ~ 160 – 200 . The dashed line marks $\min_i R_i = 0$ (no agent loses reward overall).




Figure 4: Final metrics across all conditions. (a) Efficiency: all conditions converge to $U \approx 2.5$ – 3.2 . (b) Equality: maximin-optimized Cleanup runs achieve $E \approx 1.0$, while efficiency-optimized runs show $E \approx 0.5$ – 0.7 ; Gathering achieves high equality regardless. (c) Maximin: the sharpest contrast—efficiency optimization leaves the worst-off agent deeply negative ($\min_i R_i \approx -200$), while maximin optimization transforms it to $+290$ (Gemini) or $+179$ (Sonnet). Gathering achieves high maximin (~ 580) under efficiency optimization alone. Error bars show s.d. across runs.

419 **Cost breakdown.** Each inner loop evaluation involves $K=2$ – 3 policy LLM generation calls (each
420 ~ 2 k input tokens, ~ 1.5 k output tokens) followed by multi-seed simulation (5–12 seeds, ~ 15 – 30 s
421 total). Policy LLM generation dominates inner loop time (86–97%), with Sonnet evaluations taking
422 $\sim 2\times$ longer than Gemini due to extended thinking. The researcher agent (Claude Opus 4.6, running
423 via Claude Code CLI) accounts for 41% of total wall-clock time (25.6h of the 62.2h total across all
424 12 runs), spent reading results, editing pipeline source files, and planning modifications. In monetary
425 terms, the researcher agent is the dominant cost: each outer iteration consumes ~ 50 – 100 k context
426 tokens in the Opus session, whereas each inner loop evaluation uses only ~ 5 – 10 k policy LLM
427 tokens. The 6 GEPA comparison runs required 5.0h of compute with no researcher agent, operating
428 at lower cost per run but achieving substantially lower performance (Table 2).

Table 4: Wall-clock time per autonomous run. Evals: total inner loop evaluations including initial baseline. Per eval: mean wall-clock per single evaluation. Total: end-to-end including researcher overhead.

Policy LLM	Φ	Evals	Inner loop (h)	Per eval (min)	Total (h)
<i>Cleanup</i> ($N=10$)					
Gemini	Φ_U	11	2.6	14	3.7
Gemini	Φ_U	18	4.1	14	6.4
Sonnet	Φ_U	4	2.1	32	6.1
Sonnet	Φ_U	7	5.0	43	6.1
Gemini	$\Phi_{\min}$	17	3.0	11	4.3
Gemini	$\Phi_{\min}$	11	3.6	20	5.0
Sonnet	$\Phi_{\min}$	8	4.5	33	8.8
Sonnet	$\Phi_{\min}$	9	5.3	36	13.2
<i>Gathering</i> ($N=4$)					
Sonnet	Φ_U	3	1.7	33	2.2
Gemini	Φ_U	5	1.4	17	1.9
Gemini	Φ_U	3	0.9	19	1.1
Sonnet	Φ_U	4	2.3	35	3.4
All 12 runs		100	36.6	22	62.2
GEPA (6 runs)		6	5.0	50	—

429 C Smaller Open-Weight Model: Gemma 4 26B

430 We test the framework with Gemma 4 26B-A4B-IT (Google), a 26B-parameter open-weight model
 431 substantially smaller than the frontier models in the main experiments. We run three experiments—
 432 two on Cleanup ($N=10$) optimizing efficiency and maximin respectively, and one on Gathering
 433 ($N=4$) optimizing efficiency—all with Opus 4.6 as the researcher $\mathcal{R}$.

434 **Setup.** In all three experiments, Gemma starts from complete failure: the baseline pipeline pro-
 435 duces $U=-10.0$ on Cleanup (agents CLEAN-spam without collecting apples) and $U=0.0$ on Gath-
 436 ering (broken BFS calls), compared to $U=1.93/0.86$ (Gemini/Sonnet on Cleanup) and $U=2.04/0.03$
 437 (Gemini/Sonnet on Gathering). The researcher ran $J=10-18$ outer iterations per condition.

438 **Results.** Table 5 presents the best configurations discovered. Under efficiency optimization, the
 439 researcher achieves $U=0.87$ —a dramatic recovery from the broken baseline but far below fron-
 440 tier models ($U \approx 3.2$). The researcher compensates for the model’s limited capability by reducing
 441 inner-loop iterations to $K=1$ (avoiding the regression that plagued $K \geq 2$ runs) and providing highly
 442 structured worked examples with explicit cleaning-role assignment.

Table 5: Autoresearch with Gemma 4 26B-A4B-IT. Single run per condition. Baseline: pipeline from [3].

Game	Target	U	E	$\min_i R_i$	Keep rate
Cleanup	Baseline	-10.0	1.00 [†]	-1000	—
	Φ_U	0.87	-1.11	-371	6/19
	$\Phi_{\min}$	1.71	0.94	137	9/17
Gathering	Baseline	0.0	1.00 [†]	0	—
	Φ_U	2.44	0.98	580	4/11

[†]Equality is trivially 1.0 because all agents receive near-zero reward.

443 **Maximin optimization rescues efficiency.** Strikingly, under maximin optimization the researcher
 444 achieves *higher* efficiency ($U=1.71$) than under direct efficiency optimization ($U=0.87$), alongside
 445 strong equality ($E=0.94$) and positive worst-off welfare ($\min_i R_i=137$). This reversal—absent in
 446 frontier models, where both objectives yield $U \approx 3.1-3.2$ —occurs because the maximin objective
 447 forces the researcher toward coordination mechanisms (rotating cleaning duties, index-based apple

448 assignment) that simultaneously improve collective output. Under efficiency-only optimization, the
 449 researcher converges on a local optimum—static role assignment with 25% dedicated cleaners—that
 450 a 26B model can implement but that caps performance well below the frontier.

451 This suggests that *for models below a capability threshold, fairness objectives may serve as better*
 452 *optimization targets for overall social welfare than direct efficiency maximization.* The structured
 453 coordination enforced by the maximin objective provides a scaffold that compensates for the weaker
 454 model’s difficulty implementing complex strategies from strategic hints alone.

455 **Gathering: full recovery on a simpler game.** On Gathering ($N=4$), the researcher brings
 456 Gemma from $U=0.0$ to $U=2.44$ with $E=0.98$ and $\min_i R_i=580$, after 10 outer iterations (4 kept).
 457 These values nearly match frontier models (Gemini: $U=2.49$, Sonnet: $U=2.52$; Table 2). The
 458 researcher discovers the same Voronoi partitioning and respawn-aware camping strategies found
 459 in the main experiments, and increases inner-loop iterations to $K=5$ to allow sufficient refinement.
 460 The contrast with Cleanup—where Gemma peaks at $U=0.87/1.71$ versus frontier $U\approx 3.2$ —suggests
 461 that the researcher can fully compensate for model weakness on simpler coordination tasks (4 agents,
 462 symmetric costs) but not on harder provision dilemmas (10 agents, asymmetric costs).

463 D Researcher-Authored Pipeline Artifacts

464 The previous appendix shows the policies π_c^* that the inner loop *outputs*; this one shows the con-
 465 figuration $c = (p, \phi, \mathcal{H}, \iota, v)$ (Section 4, Table 1) that the researcher $\mathcal{R}$ *authored* to produce them.
 466 Excerpts are taken from the final commit on the dedicated git branch of the corresponding run; we
 467 reproduce the prose verbatim, with author commentary in [brackets] and elisions marked
 468 We organize by artifact type so each subsection makes a within-type contrast (e.g., Φ_U vs. $\Phi_{\min}$, or
 469 Cleanup vs. Gathering).

470 D.1 Helper library $\mathcal{H}$: coordination primitives

471 The researcher adds primitives that the policy LLM can call as black boxes, sparing it from re-
 472 implementing tricky logic on every iteration. Two patterns dominate: (i) *state inspection* (waste
 473 counts, fractions, beam-yield scoring) and (ii) *coordination primitives* (zone assignment, role rota-
 474 tion). We show one of each.

Listing 1: Cleanup, $\Phi_{\min}$ – band-based apple zoning helper added by the researcher in exp5. This is the primitive that lets the policy in Listing 6 keep collectors within their own row band, preventing all 7 gatherers from racing to the same apple.

```

475 1 def get_my_apples(env, agent_id):
476 2     """Alive apples in this agent's horizontal row band.
477 3     Divides the orchard into n_agents bands; falls back to all apples if empty."""
478 4     aid = int(agent_id)
479 5     n = int(env.n_agents)
480 6     band_h = env.height / n
481 7     band_lo, band_hi = aid * band_h, (aid + 1) * band_h
482 8
483 9     my_apples, all_apples = set(), set()
484 10    for i in range(env.n_apples):
485 11        if env.apple_alive[i]:
486 12            ar = int(env._apple_pos[i, 0]); ac = int(env._apple_pos[i, 1])
487 13            all_apples.add((ar, ac))
488 14            if band_lo <= ar < band_hi:
489 15                my_apples.add((ar, ac))
490 16    return my_apples if my_apples else all_apples
491

```

Listing 2: Gathering – BFS-Voronoi territory + respawn-aware waiting helpers added by the researcher in gather-exp3. These are the primitives that Listing 8 composes into a one-line policy: “go to my_zone_apples, otherwise nearest_respawning_apple.”

```

493 1 def voronoi_zones(env):
494 2     """Multi-source BFS Voronoi over walkable cells.
495 3     Returns (row, col) -> agent_id; ties broken by lower agent_id."""
496 4     queue, visited = deque(), {}
497 5     for a_id in range(env.n_agents):
498 6         if int(env.agent_timeout[a_id]) == 0:
499 7             ar = int(env.agent_pos[a_id][0]); ac = int(env.agent_pos[a_id][1])
500 8             queue.append((ar, ac, a_id))
501

```

```

502     visited[(ar, ac)] = a_id
503 while queue:
504     r, c, a_id = queue.popleft()
505     for dr, dc in [(-1,0),(1,0),(0,-1),(0,1)]:
506         nr, nc = r + dr, c + dc
507         if 0 <= nr < env.height and 0 <= nc < env.width \
508             and not env.walls[nr, nc] and (nr, nc) not in visited:
509             visited[(nr, nc)] = a_id
510             queue.append((nr, nc, a_id))
511 return visited
512
513 def nearest_respawning_apple(env, agent_id, zones=None, max_timer=10):
514     """Nearest dead apple in MY zone respawning within max_timer steps.
515     Returns (row, col) of best wait spot, or None."""
516     if zones is None: zones = voronoi_zones(env)
517     best_pos, best_t, best_d = None, max_timer + 1, float('inf')
518     ar = int(env.agent_pos[agent_id][0]); ac = int(env.agent_pos[agent_id][1])
519     for i in range(env.n_apples):
520         if not env.apple_alive[i]:
521             pos = (int(env._apple_pos[i][0]), int(env._apple_pos[i][1]))
522             if zones.get(pos) == agent_id:
523                 t = int(env.apple_timer[i])
524                 if t <= max_timer:
525                     d = abs(pos[0] - ar) + abs(pos[1] - ac)
526                     if t < best_t or (t == best_t and d < best_d):
527                         best_pos, best_t, best_d = pos, t, d
528 return best_pos

```

530 The choice of helper depends on the dilemma’s geometry: row-bands suffice when fairness is the
531 main concern (Cleanup maximin) but full BFS-Voronoi is needed when wall-aware territory owner-
532 ship matters (Gathering). The researcher does not fix this in advance — it picks the simpler primitive
533 whenever it works.

534 D.2 System prompt p : from neutral framing to strategic briefing

535 The unmodified system prompt is a neutral API description: “Write a policy that maximizes per-
536 agent reward.” Under maximin optimization, the researcher rewrites the prompt into a multi-section
537 strategic briefing. The excerpt below shows the pieces it added to the Cleanup prompt in exp5; the
538 full prompt grew from 165 to 325 lines.

Listing 3: Cleanup, $\Phi_{\min}$ – excerpts from the researcher-rewritten system prompt (exp5). The unmodified baseline contained only API documentation; everything below was added by $\mathcal{R}$.

```

539 ## CRITICAL OBJECTIVE: Rawlsian Fairness (Maximin)
540
541 Your goal is to maximize the minimum per-agent total return across all
542 agents. This is the "maximin" or Rawlsian welfare criterion.
543
544 This means: it is NOT enough for the *average* agent to do well. The
545 worst-off agent must do as well as possible. A policy where 8 agents
546 each earn +200 but 2 agents each earn -100 is TERRIBLE (maximin = -100).
547
548 Key implication: cleaning costs (-1 per CLEAN action) must be shared
549 equitably among ALL agents. If you assign fixed "cleaner" roles, those
550 agents will accumulate large negative rewards and destroy your maximin score.
551
552 ## Strategy for Maximin: ROLE ROTATION + APPLE ZONING
553
554 The optimal strategy for maximin involves:
555 1. Shared cleaning duty: ALL agents take turns cleaning using a rotation
556    schedule based on 'agent_id' and 'env.step_count'. For example:
557    'is_my_cleaning_turn = (agent_id + env.step_count // SHIFT) % env.n_agents < NUM_CLEANERS'
558    where SHIFT is ~50 steps and NUM_CLEANERS is 2-3.
559 2. When NOT your turn: collect apples from YOUR ZONE using
560    'get_my_apples(env, agent_id)' -- prevents all gatherers competing
561    for the same nearest apple.
562 3. NEVER use BEAM (action 6): it costs -1 and causes -50 to the target.
563 4. Keep waste density low: aim for ~2-3 active cleaners at any time.
564
565 [... API documentation, helper documentation, working example ...]
566
567 IMPORTANT:
568 - NEVER assign permanent cleaning roles to specific agents -- this kills maximin.
569 - Use env.step_count for time-based role rotation so all agents share cleaning duty.
570 - NEVER use BEAM (action 6) -- it destroys both agents' rewards.
571
572

```

Two things stand out. First, the researcher writes the formula it expects $\mathcal{M}$ to use almost verbatim — (agent_id + env._step_count // SHIFT) % env.n_agents < NUM_CLEANERS — and the policy in Listing 6 uses essentially this template. Second, the researcher *teaches by counter-example*: the explicit “8 agents earn +200, 2 earn -100, maximin = -100” makes the failure mode of Φ_U -style static roles (Listing 5) concretely visible to $\mathcal{M}$. This is the information-design move predicted by the mechanism-design framing in Section 3.4.

The Gathering prompt evolves along a different axis (no maximin, no rotation). Its researcher-added “Key Strategic Insights” section instead emphasizes (i) *self-play implies never beam* and (ii) *Voronoi partition each step*, exactly matching what the policy in Listing 8 implements.

D.3 Feedback ϕ : adaptive diagnostics

The baseline feedback function from [3] shows reward + four social metrics with definitions and stops there. The researcher turns it into a state machine: it inspects the latest metrics and conditionally injects different hints. Two different runs produced two qualitatively different diagnostics — the same metric ($\bar{r}_k$ or $\min_i R_i$) is repurposed to either *stabilize* a working policy or *redirect* a failing one.

Listing 4: Adaptive feedback diagnostics. *Top*: efficiency-optimized run (exp1) – a stability guard prevents iter-3 regression once U is high. *Bottom*: maximin-optimized run (exp5) – two fairness diagnostics that fire when the worst-off agent loses reward.

```

588 1 # --- Cleanup, Phi_U feedback (exp1, pipeline/feedback.py final) ---
589 2 last_eff = history[-1]["metrics"]["efficiency"]
590 3 if last_eff >= 2.5:
591 4     parts.append(
592 5         """CRITICAL -- DO NOT REGRESS*: The current policy achieves high "
593 6         "efficiency (>=2.5). You MUST output a policy that is nearly identical "
594 7         "to the current one. Copy the current policy code and make AT MOST "
595 8         "one small targeted improvement (e.g., adjust a single numeric "
596 9         "threshold). If you are not confident a change will help, output the "
597 10        "current policy UNCHANGED. A regression here means the run fails.")
598 11
599 12 # --- Cleanup, Phi_min feedback (exp5, pipeline/feedback.py final) ---
600 13 last_maximin = history[-1]["metrics"].get("maximin", 0)
601 14 last_avg = history[-1]["reward_avg"]
602 15 if last_maximin < 0:
603 16     parts.append(
604 17         f"""FAIRNESS ALERT*: maximin={last_maximin:.1f} is NEGATIVE. "
605 18         f"The worst-off agent lost reward overall while average was {last_avg:.1f}. "
606 19         f"This means cleaning duties are NOT shared equitably. "
607 20         f"Use time-based role rotation (env._step_count) so ALL agents share "
608 21         f"cleaning costs equally. NEVER assign permanent cleaner roles.")
609 22 elif last_maximin < last_avg * 0.5:
610 23     parts.append(
611 24         f"""FAIRNESS WARNING*: maximin={last_maximin:.1f} is much lower than "
612 25         f"average={last_avg:.1f}. The gap suggests unequal cleaning burden. "
613 26         f"Ensure ALL agents rotate through cleaning duty.")
614

```

The two diagnostics are written by independent researcher runs but converge on a common pattern: *thresholded interventions on a primary metric*. They are also a common cause of the low run-to-run variance reported in Finding 1: when a working policy lands in the high- U basin, the stability guard prevents the LLM from over-refining and tipping out of it (the $K=3$ regressions visible in Section 4’s “Common failure modes” (4)), and when an unfair policy lands in the negative-maximin region, the alert injects exactly the rotation hint that recovers it.

D.4 Iteration logic ι : per-condition hyperparameters

Although ι has the smallest source surface, the researcher’s edits to it explain a meaningful fraction of the variance reduction noted in Finding 1. Table 6 summarizes the values shipped in the final commit of each best-performing run.

The maximin Gemini run is the most informative case: the researcher discovered that with $K=3$ and $|S|=5$, identical configurations could yield $\min_i R_i=295$ on one set of seeds and $\min_i R_i=100$ on another (this is the “variance exposure” failure mode of Section 4’s common failures). It then walked $|S|$ from 5 $\rightarrow$ 8 $\rightarrow$ 12 over three outer iterations until the seed-averaged signal was stable enough that the policy LLM stopped chasing noise. The policy in Listing 6 is sampled at $|S|=12$.

Table 6: Iteration-logic (ι) values in the final commit of selected runs. K = inner-loop iterations, $|S|$ = evaluation seeds, “thinking” = extended-thinking token budget passed to $\mathcal{M}$. Baseline ([3]) is $K=3$, $|S|=5$, thinking 16k.

Run	Game / objective	K	$ S $	thinking	Researcher’s stated rationale
exp1 (Gemini)	Cleanup, Φ_U	3	5	16k	default; stability comes from feedback hint
exp4 (Sonnet)	Cleanup, Φ_U	3	5	32k	“Sonnet was over-refining at 16k thinking”
exp5 (Gemini)	Cleanup, $\Phi_{\min}$	2	12	16k	“ $K=2$ avoids iter-3 regression; 12 seeds for variance control”
exp7 (Sonnet)	Cleanup, $\Phi_{\min}$	3	5	10k	“reduced thinking from 16k – Sonnet was generating runaway code at 16k”
gather-exp3 (Gemini)	Gathering, Φ_U	3	5	16k	default; Gathering converges in $K \leq 3$

D.5 Cross-artifact observations

- *Helpers do the algebra; prompts pick the strategy class.* The researcher uses helpers (1, 2) to put non-trivial primitives at $\mathcal{M}$ ’s fingertips, and the prompt (3) to tell $\mathcal{M}$ *which* primitive to compose. Neither is sufficient alone — prompts without supporting helpers produced policies that “mention rotation” but failed to implement it; helpers without prompt updates often went unused.
- *Feedback is the only adaptive component.* p , $\mathcal{H}$, and ι are static within a run; ϕ is the only place where the researcher gets to change what $\mathcal{M}$ sees *during* the inner loop, and it uses thresholded diagnostics (Listing 4) to do so. This is what mostly drives the run-to-run variance reduction.
- *The same artifact has different content under each objective.* The Cleanup prompt under Φ_U omits “rotation” and “maximin” entirely; the same prompt under $\Phi_{\min}$ devotes its entire opening section to it. The configuration c is genuinely a function of the welfare objective Φ , not a fixed pipeline parameterized by it — consistent with the information-design framing of Section 3.4.
- *The researcher’s edits are short and targeted.* The full diff between the baseline pipeline and the best maximin Gemini configuration totals ~ 300 added lines spread over four files; ~ 80 of them appear in the listings above. The remaining content is API documentation and worked examples, included for completeness but not strategically novel.

E Selected Generated Policies

The code excerpts below are verbatim outputs of the policy synthesizer $\mathcal{M}$, taken from the final inner-loop iteration of the best-performing run in each condition. To fit the page, we elide repeated boilerplate (BFS scaffolding, fallback branches) with `...` and add commentary in [brackets]; nothing else has been edited. The listings illustrate the qualitative differences highlighted in Findings 2–4 of Section 4: *static* vs. *rotating* cleaning roles in Cleanup, and the *Voronoi* + *respawn-timer* structure in Gathering. All four policies were independently generated by the LLM under different seeds, models, and outer-loop runs.

E.1 Cleanup, Φ_U (Gemini): static interleaved roles + dynamic threshold

Best run: $U=3.25$, $E=0.61$, $\min_i R_i = -182$. The synthesizer assigns *permanent* cleaning roles via `agent_id mod n <  $\tau$` , where the threshold τ depends smoothly on the current waste fraction. Higher pollution recruits more cleaners, but the *identity* of each cleaner is fixed — agents $0-\tau-1$ always pay the cost. This drives U to its ceiling but leaves equality at 0.61 and pushes the worst-off agent to $\min_i R_i = -182$.

Listing 5: Cleanup, Φ_U – Gemini 3.1 Pro. Static role assignment with waste-tiered threshold (verbatim, abridged).

```

def policy(env, agent_id) -> int:
    if int(env.agent_timeout[agent_id]) > 0:
        return 7 # STAND while removed
    waste_frac = waste_fraction(env)
    n_agents = env.n_agents
    # [Smooth dynamic threshold: more pollution -> more cleaners]
    if waste_frac > 0.35: threshold_cleaners = n_agents
    elif waste_frac > 0.25: threshold_cleaners = int(n_agents * 0.7)
    elif waste_frac > 0.15: threshold_cleaners = int(n_agents * 0.5)
    elif waste_frac > 0.05: threshold_cleaners = int(n_agents * 0.3)

```

```

674 13     else:
675 14         threshold_cleaners = max(1, int(n_agents * 0.2))
676 15
677 16     # [STATIC role: cleaners are ALWAYS the lowest-id agents -- no rotation]
678 17     is_cleaner = (agent_id % n_agents) < threshold_cleaners
679 18
680 19     if is_cleaner:
681 20         best_orient, n_waste = find_best_clean_orientation(env, agent_id)
682 21         if n_waste >= 1:
683 22             cur_orient = int(env.agent_orient[agent_id])
684 23             if best_orient == cur_orient: return 8 # CLEAN
685 24             if (cur_orient + 1) % 4 == best_orient: return 5 # ROTATE_RIGHT
686 25             elif (cur_orient - 1) % 4 == best_orient: return 4 # ROTATE_LEFT
687 26             else: return 4
688 27         # [Competitor-aware BFS toward waste in own row band]
689 28         c_idx = agent_id % n_agents
690 29         target_r = int(((c_idx + 0.5) / threshold_cleaners) * env.height)
691 30         # ... pick waste cell minimizing dist + |row - target_r| + 50 *
692 31         # closer_cleaners
693 32         # ... bfs_toward and direction_to_action
694 33     else:
695 34         # GATHERER: row-banded sweep, competitor-aware
696 35         assigned_row = ((g_idx + 0.5) / n_gatherers) * env.height
697 36         # ... pick alive apple minimizing dist + |row - assigned_row| + 50 *
698 37         # closer_gatherers
699 38     ...

```

700 E.2 Cleanup, $\Phi_{\min}$ (Gemini): time-rotated, geographically interleaved roles

701 Best run: $U=3.19$, $E=0.98$, $\min_i R_i=296$. Switching the outer objective from U to $\min_i R_i$ flips
702 the role-assignment scheme entirely: `role_idx` now depends on `agent_id + step/50`, so every
703 agent rotates through cleaner and gatherer duty in a 50-step cycle. The cleaner indices $\{1, 4, 8\}$ are
704 spatially *interleaved* along the river (rather than contiguous) so that whoever is on duty has a short
705 walk to their assigned river slice. This single change accounts for the $E=0.55 \rightarrow 0.98$ jump and the
706 $\min_i R_i=-182 \rightarrow 296$ swing reported in Table 2, with negligible efficiency loss.

Listing 6: Cleanup, $\Phi_{\min}$ – Gemini 3.1 Pro. Time-rotated roles with interleaved geography (verbatim, abridged).

```

707 1 def policy(env, agent_id) -> int:
708 2     if int(env.agent_timeout[agent_id]) > 0:
709 3         return 7 # STAND while removed
710 4
711 5     aid = int(agent_id)
712 6     step = int(env._step_count)
713 7     n = int(env.n_agents)
714 8
715 9     # [TIME ROTATION: every 50 steps each agent's role index advances]
716 10    if n == 10:
717 11        role_idx = (aid + step // 50) % 10
718 12        # [Cleaner indices {1,4,8} are SPATIALLY INTERLEAVED, not contiguous,
719 13        # so each on-duty cleaner has a short walk to its river slice]
720 14        if role_idx == 1: role_type, zone_idx = 'C', 0
721 15        elif role_idx == 4: role_type, zone_idx = 'C', 1
722 16        elif role_idx == 8: role_type, zone_idx = 'C', 2
723 17        else:
724 18            role_type = 'G'
725 19            g_map = {0:0, 2:1, 3:2, 5:3, 6:4, 7:5, 9:6}
726 20            zone_idx = g_map[role_idx]
727 21            num_cleaners, num_gatherers = 3, 7
728 22
729 23    if role_type == 'C':
730 24        # [Each on-duty cleaner owns one horizontal slice of the river]
731 25        row_min = int(zone_idx * env.height / num_cleaners)
732 26        row_max = int((zone_idx + 1) * env.height / num_cleaners) - 1
733 27
734 28        # [Score every (row, col, orientation) in the slice by beam yield;
735 29        # fire if facing >=2 waste, else move to a >=2-yield position]
736 30        best_positions, good_positions = set(), set()
737 31        for r in range(row_min, row_max + 1):
738 32            for c in range(10):
739 33                if env.walls[r, c] or env.waste[r, c]: continue
740 34                y_max = max(get_clean_yield(env, r, c, o) for o in range(4))
741 35                if y_max >= 2: best_positions.add((r, c))
742 36                if y_max >= 1: good_positions.add((r, c))
743 37            # ... rotate / CLEAN / bfs_to_target_set as appropriate
744 38        else:
745 39            # GATHERER: collect apples ONLY inside the rotating row band

```

```

747 40 # row_min/row_max for num_gatherers slices ...
748 41 ...

```

750 E.3 Cleanup, $\Phi_{\min}$ (Sonnet): independent rediscovery of duty rotation

751 Best run: $U=2.93$, $E=0.83$, $\min_i R_i=154$. Run with a different policy LLM (Sonnet 4.6) and on
752 a different outer-loop trajectory, the synthesizer arrives at the *same* structural insight as the Gemini
753 maximin run: a phase counter ($\text{agent_id} + \text{step}/50$) mod n rotates which 2 of the 10 agents clean
754 at any given time, and a *separate* 200-step zone counter rotates which 5-row band each collector
755 sweeps. The two rotation periods (50 and 200 steps) ensure every agent visits every cleaning slot
756 and every apple zone within one episode — a structural fairness invariant the researcher discovered
757 without any explicit instruction to “rotate.”

Listing 7: Cleanup, $\Phi_{\min}$ – Sonnet 4.6. Phase-rotated cleaning + zone-rotated collection (verbatim, abridged).

```

758 1 def policy(env, agent_id) -> int:
759 2     if int(env.agent_timeout[agent_id]) > 0:
760 3         return 7 # STAND while removed
761 4
762 5     n = env.n_agents
763 6     step = env._step_count
764 7     wf = waste_fraction(env)
765 8
766 9     # --- ROLE ASSIGNMENT: fair 2-cleaner rotation ---
767 10    # [50-step phases x 20 phases / episode -> each agent cleans 4 phases (~20%)]
768 11    phase = step // 50
769 12    cleaner_rank = (agent_id + phase) % n # [cycles 0..9 fairly]
770 13    is_cleaner = cleaner_rank < 2 # [exactly 2 cleaners per phase]
771 14
772 15    # [Emergency override: all agents clean when waste crosses the spawn cliff]
773 16    in_emergency = wf >= 0.40
774 17    if in_emergency: is_cleaner = True
775 18
776 19    # [River split: cleaner_slot 0 -> top half, 1 -> bottom half]
777 20    cleaner_slot = (agent_id % 2) if in_emergency else cleaner_rank
778 21
779 22    if is_cleaner and wf > 0.04:
780 23        # [Scan 4 orientations from current position; if best beam shot
781 24        # covers >=3 waste cells, fire; else step to a strictly better neighbor]
782 25        # ... (best_o, best_cnt) = max over 4 directions
783 26        # ... if best_cnt >= 3: rotate to best_o, then CLEAN
784 27        # ... else: scan 4 adjacent cells x 4 orientations for higher yield
785 28        # ... if no waste in beam range: bfs to waste_set in OWN river half,
786 29        # fall back to whole river
787 30
788 31    # --- COLLECTOR LOGIC: 5-zone rotating apple collection ---
789 32    # [5 zones x 200 steps = 1000 steps -> each agent visits every zone once;
790 33    # 2 agents share each zone at any time, reducing competition 10 -> 2]
791 34    bounds = zone_boundaries(env, 5)
792 35    zone_phase = step // 200
793 36    zone = (agent_id + zone_phase) % 5
794 37    z_start, z_end = bounds[zone], bounds[zone + 1]
795 38
796 39    zone_apples = {(int(env._apple_pos[i][0]), int(env._apple_pos[i][1]))
797 40                  for i in range(env.n_apples) if env.apple_alive[i]
798 41                  and z_start <= int(env._apple_pos[i][0]) < z_end}
799 42    # ... bfs_to_target_set(zone_apples), else bfs_nearest_apple as fallback
800 43    return 7
801 44
802 45

```

803 E.4 Gathering (Gemini): wall-aware Voronoi + spatiotemporal targeting

804 Best run: $U=2.47$, $E=0.98$, $\min_i R_i=580$. In Gathering, where costs are symmetric, no role dif-
805 ferentiation is needed: the entire optimization is spatial and temporal. The synthesizer (i) parti-
806 tions cells into BFS-Voronoi territories owned by individual agents, (ii) runs a single centralized
807 BFS from its own position to score every reachable cell with exact wall-aware distances, and
808 (iii) for both alive and respawning apples in its territory, computes the *earliest collection time*
809 $\max(\text{walk_dist}, \text{respawn_timer})$ and targets the minimum. When the agent must wait on a
810 dead apple, it steps onto a *non-spawn* adjacent cell rather than blocking a respawn point. This single
811 sort key — $\max(\text{distance}, \text{timer})$ — subsumes both “go to nearest apple” and “camp respawn” as
812 special cases.

Listing 8: Gathering – Gemini 3.1 Pro. Voronoi territories + spatiotemporal priority targeting (verbatim, abridged).

```

813 1 def policy(env, agent_id) -> int:
814 2     if int(env.agent_timeout[agent_id]) > 0:
815 3         return 7
816 4
817 5     r, c = int(env.agent_pos[agent_id][0]), int(env.agent_pos[agent_id][1])
818 6     orient = int(env.agent_orient[agent_id])
819 7
820 8     # [Wall-aware Voronoi over the gridworld; spawn_points = all apple cells]
821 9     zones = voronoi_zones(env)
822 10    spawn_points = {(int(p[0]), int(p[1])) for p in env._apple_pos}
823 11
824 12    # [Single centralized BFS: exact distance + first move to every reachable cell]
825 13    distances, first_moves = {}, {}
826 14    q = deque([(r, c, 0, None)])
827 15    visited = {(r, c)}
828 16    while q:
829 17        cr, cc, d, fm = q.popleft()
830 18        distances[(cr, cc)] = d
831 19        first_moves[(cr, cc)] = fm
832 20        for dr, dc in [(-1,0),(1,0),(0,-1),(0,1)]:
833 21            nr, nc = cr + dr, cc + dc
834 22            if 0 <= nr < env.height and 0 <= nc < env.width and not env.walls[nr,
835 23                nc]:
836 24                if (nr, nc) not in visited:
837 25                    visited.add((nr, nc))
838 26                    q.append((nr, nc, d + 1, fm if fm is not None else (dr, dc)))
839 27
840 28    # [Both alive AND respawning apples in MY zone, with their respawn timer]
841 29    my_zone_spawns = []
842 30    for i in range(env.n_apples):
843 31        pr, pc = int(env._apple_pos[i][0]), int(env._apple_pos[i][1])
844 32        if zones.get((pr, pc)) == agent_id and (pr, pc) in distances:
845 33            my_zone_spawns.append((pr, pc, int(env.apple_timer[i]), distances[(pr,
846 34                pc)]))
847 35
848 36    # [SPATIOTEMPORAL PRIORITY: earliest collection time first.
849 37    # score = max(walk_distance, respawn_timer); ties broken by sooner timer]
850 38    my_zone_spawns.sort(key=lambda x: (max(x[3], x[2]), x[2], x[3]))
851 39
852 40    for pr, pc, timer, dist in my_zone_spawns:
853 41        if timer == 0: # [Alive apple: walk straight to it]
854 42            if dist == 0: return 7
855 43            fm = first_moves[(pr, pc)]
856 44            return direction_to_action(fm[0], fm[1], orient)
857 45        else: # [Dead apple: camp on safe adjacent cell]
858 46            # ... pick nearest neighbor (nr,nc) that is NOT in spawn_points,
859 47            # ... navigate there and STAND while waiting for respawn
860 48            # [Global poaching fallback if zone is empty: nearest alive apple anywhere]
861 49            ...

```

865 E.5 Cross-condition observations

866 Several patterns are visible across these four listings (and confirmed by inspecting the remaining 8
867 runs not shown):

- 868 • *Role assignment encodes the welfare objective.* Static $\text{agent_id} < \tau$ (Listing 5) maximizes
869 U but harms $\min_i R_i$. Time-rotated $(\text{agent_id} + \text{step} // T) \% n$ (Listings 6, 7) maximizes
870 $\min_i R_i$ at $\leq 1\%$ efficiency cost (Gemini).
- 871 • *Convergent rediscovery across LLMs.* Gemini and Sonnet, on independent runs, both arrive at
872 the $(\text{id} + \text{phase}) \% n$ duty-rotation idiom with similar phase lengths ($T \approx 50$ steps) under
873 maximin — and both omit it under efficiency.
- 874 • *Game structure dictates strategy class.* Cleanup policies are dominated by role-assignment logic
875 (cleaner vs. gatherer); Gathering policies are dominated by territory partitioning and respawn-
876 timer reasoning, with no role differentiation at all. The researcher does not need to be told which
877 class of strategy to write.
- 878 • *Earliest-collection-time as a unifying score.* The Gathering policy’s $\max(\text{walk}, \text{timer})$ key (List-
879 ing 8) is structurally identical across the 4 Gathering runs (both Gemini and Sonnet), differing
880 only in how the Voronoi diagram is computed — Manhattan approximation, fully BFS-based, or
881 cached helper.